



**Hewlett Packard  
Enterprise**

# Programming Environment and Modules

---

Comprehensive General LUMI Course

April 23–66, 2024

# Agenda

---

- The Cray Programming Environment
- Controlling the Environment with modules



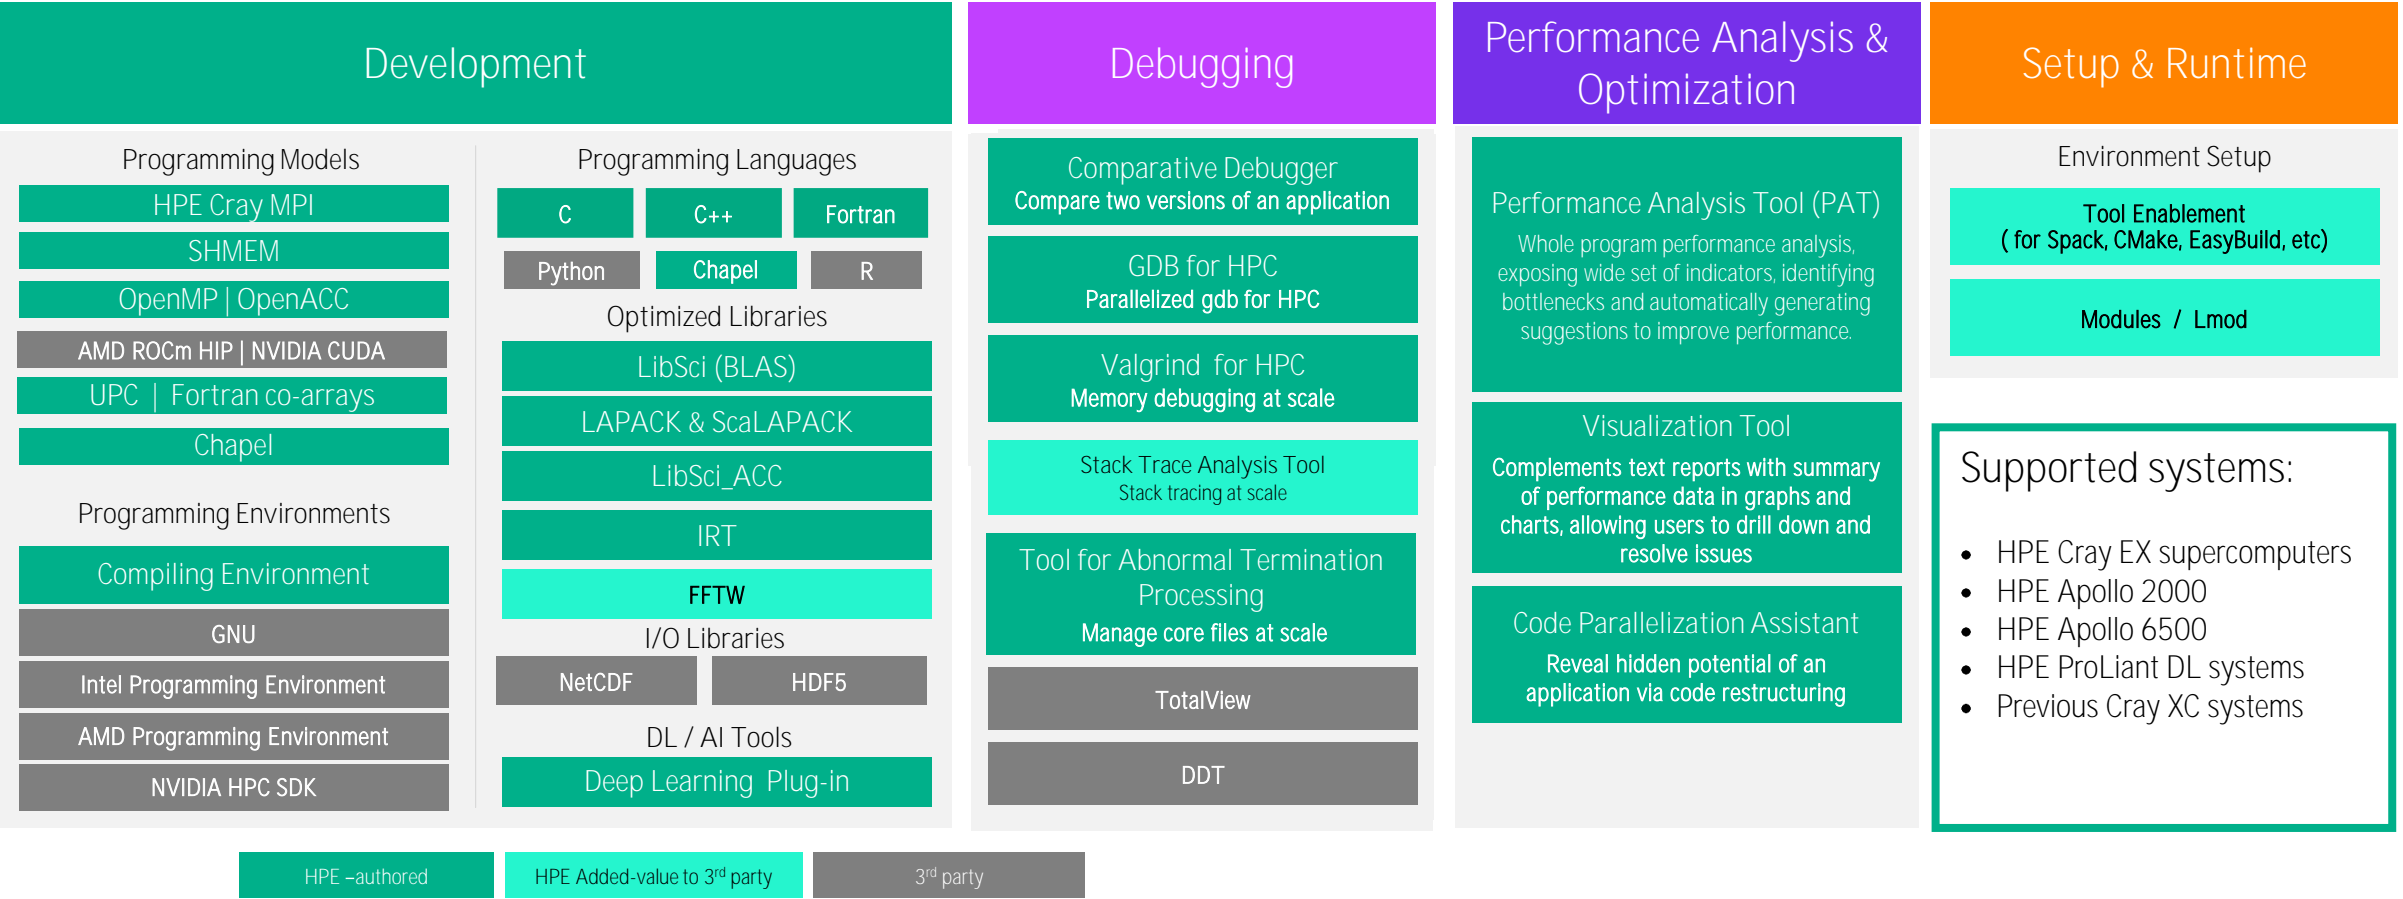
# Vision



- HPE Cray EX systems are designed to be High Productivity as well as High Performance Computers
- The Cray Programming Environment (CPE) provides a simple consistent interface to users and developers.
  - Focus on improving scalability and reducing complexity
- The default Programming Environment provides:
  - the highest levels of application performance
  - a rich variety of commonly used tools and libraries
  - a consistent interface to multiple compilers and libraries
  - an increased automation of routine tasks
- We continue to develop and refine the PE
  - Frequent communication and feedback to/from users
  - Strong collaborations with third-party developers



# Cray Developer Environment for EX



# Components of the Programming Environment

- Cray Compiling Environment (CCE)  
Optimizing compilers for Fortran, C, C++ and UPC
- Cray Scientific and Math Libraries (CSML)  
High Performance libraries for scientific applications (BLAS, LAPACK, ScaLAPCK, FFTW, NetCDF)
- Cray Message Passing Toolkit (CMPT)  
Provides the Message Passing Interface (MPI) and OpenSHMEM
- Cray Environment Setup and Compiling Support (CENV)  
Infrastructure to support the development environment.  
Includes compiler drivers, hugepages support and the PE packaging API
- Cray Performance Measurement and Analysis Tools (CPMAT)  
Provides tools to analyse performance and behaviour of programs and the PAPI performance API
- Cray Debugging Support Tools (CDST)  
Provides debugging tools including gdb4hpc and valgrind4hpc



# Cray Compiling Environment (CCE)

- The default compiler on EX systems
  - Specifically designed for HPC applications
  - Takes advantage of Cray's experience with automatic vectorization and shared memory parallelization
- Excellent standards support for multiple languages and programming models
  - Fortran, C, C++
    - Supports most of Fortran 2018 (ISO/IEC 1539:2018) with some exceptions
    - C/C++ compiler is based on Clang/LLVM
  - OpenMP
    - Full OpenMP 4.5 and partial OpenMP 5.0 and 5.1, see `man intro_omp`
  - OpenACC
    - Full OpenACC 2.0 and partial OpenACC 2.x/3.x, see `man intro_openacc`
  - HIP compilation
  - MPI Interfaces
- Full integrated and optimised support for PGAS languages
  - UPC 1.2 and Fortran 2008 coarray support
  - No preprocessor involved
  - Full debugger support (With ARM Forge)
- OpenMP and automatic multithreading fully integrated
  - Share the same runtime and resource pool
  - Aggressive loop restructuring and scalar optimization done in the presence of OpenMP
  - Consistent interface for managing OpenMP and automatic multithreading



# Other Compilers and Programming Language Support

---

- Compilers
  - GNU
  - AOCC (AMD CPU Compiler)
  - AMD LLVM (AMD GPU support)
- Python: 3.10.10
- R: Includes R 4.2.1



# Status of Clang into CCE (introduced in CCE9)

## Cray Enhancements to Clang:

- Loopmark (-fsave-loopmark)
- Decompile (-fsave-decompile)
- UPC support (-hupc or automatic with the .hupc file extension)
- Enhanced assembly listings (line numbers, comments)
- Support for Cray Performance Measurement and Analysis Tools
- Cray-specific performance improvements
- **Cray OpenMP runtime** supports CCE Classic/Clang interoperability for codes that mix Fortran and C/C++
- Support for OpenMP 4.5 offload to NVIDIA and AMD GPUs





# AMD Optimizing C/C++ Compiler (AOCC)

- The CPE enables this but does not bundle it
- Supported as a programming environment via PrgEnv-aocc
  - Based on LLVM 13
  - Provides C/C++ and Fortran (based “classical” FLANG)
- Loading the module will provide the compiler via the standard (ftn, cc , CC) wrappers and will link appropriate libraries
- More details from:  
<https://developer.amd.com/wp-content/resources/AOCC-3.0-Install-Guide.pdf>  
<https://developer.amd.com/wp-content/resources/AOCC-3.0-User-Guide.pdf>



# Cray MPI and SHMEM

## Cray MPI

- Implementation based on MPICH3 source from ANL
- Includes many improved algorithms and tweaks for Cray hardware
  - Improved algorithms for many collectives
  - Asynchronous progress engine allows overlap of computation and comms
  - Customizable collective buffering when using MPI-IO
  - Optimized Remote Memory Access (one-sided) fully supported including passive RMA
- Full MPI-3.1 support with the exception of
  - MPI\_LONG\_DOUBLE and MPI\_C\_LONG\_DOUBLE\_COMPLEX for CCE
- Includes support for Fortran 2008 bindings (since CCE 8.3.3)

## Cray SHMEM

- Fully optimized Cray SHMEM library supported
- Fully compliant with OpenSHMEM v1.5
- Cray XC implementation close to the T3E model



# Cray Scientific and Maths Libraries (CSML)

The programming environment provides integrated tuned versions of popular scientific libraries

- LibSci
  - BLAS (Basic Linear Algebra Subroutines)
  - BLACS (Basic Linear Algebra Communication Subprograms)
  - CBLAS (wrappers providing a C interface to the FORTRAN BLAS library)
  - IRT (Iterative Refinement Toolkit)
  - LAPACK (Linear Algebra Routines)
  - LAPACKe (C interfaces to LAPACK Routines)
  - ScaLAPACK (Scalable LAPACK)
- LibSci\_ACC
  - Subset of GPU-optimized GPU routines from LibSci
- FFTW3
  - Fastest Fourier Transforms in the West, release 3
- Data libraries
  - NetCDF and HDF5



# Packages where only Build Instructions Provided

- TPSL - Cray's collection of third-party scientific libraries
    - HYPRE, METIS, ParMETIS, MUMPS, SUNDIALS, SuperLU, SuperLU\_DIST, Scotch, GLM, Matio
  - Boost
  - PETSc
  - SLEPc
  - Trilinos
  - ADIOS
  - CMake
- 
- Build instructions can be found at <https://github.com/Cray/pe-scripts>
  - Customer sites often have good documentation on building applications/packages



# Cray DSMML

- Distributed Symmetric Memory Management Library (DSMML)
- Standalone proprietary library for managing distributed shared symmetric memory heaps
- Useful for implementing PGAS languages
- See  
intro\_dsmml(3)  
Content will appear here in due course <https://pe-cray.github.io/cray-dsmml/>



# Python and R Modules

- R, cray-R (version 4.2.1)
- Python 3
- cray-python/3.10.10 (from CPE 23.09) contains:

python 3.10.10  
numpy 1.23.5  
pandas 1.4.2  
SciPy 1.10.0  
mpi4py 3.1.4  
dask 2022.7.0



# Cray PE Deep Learning plugin

- Provides a highly-optimized communication layer for Deep Learning training
- This provides mechanisms for the following which can be added to a single-process application:
  - Process identification and job size
  - Weight and bias broadcasting
  - Gradient averaging
- Accessible via Python and C APIs
- Python API supports:  
TensorFlow, PyTorch, Keras and NumPy
- Modules
  - craype-dl-plugin-py3** (standard module)
  - craype-dl-plugin-ftr** (with fault tolerant support)
- See **man intro\_dl\_plugin**
- (note that you need to have gcc loaded to see above modules with module avail)



# Debugging Support Tools (CDST)

- gdb4hpc  
Command-line parallel debugger
- valgrid4hpc  
Parallel-debugging tool for detection of memory leaks parallel application errors.  
(partially implemented on Shasta)
- Stack Trace Analysis Tool (STAT)  
Merged-stack backtrace analysis tool
- Abnormal Termination Processing (ATP)  
Scalable core file generator and analysis tool
- Cray Comparative Debugging (CCDB)  
**Side by side debugging tool for two 'versions' of an application.**





# Cray Performance Measurement and Analysis Tools

- Perftools lite modules for one-shot build/run/analysis
  - perftools-lite, perftools-lite-events, perftools-lite-loops, perftools-lite-hbm
- Full featured perftools module for multi-step collection and analysis
  - pat\_build (program instrumentation)
  - pat\_report (report generation)
  - Craypat runtime library
- pat\_run Launches a dynamically-linked program for analysis
- Also
  - PAPI, Apprentice2, stat\_view and reveal



# Controlling the Environment with Modules

---



# Modules

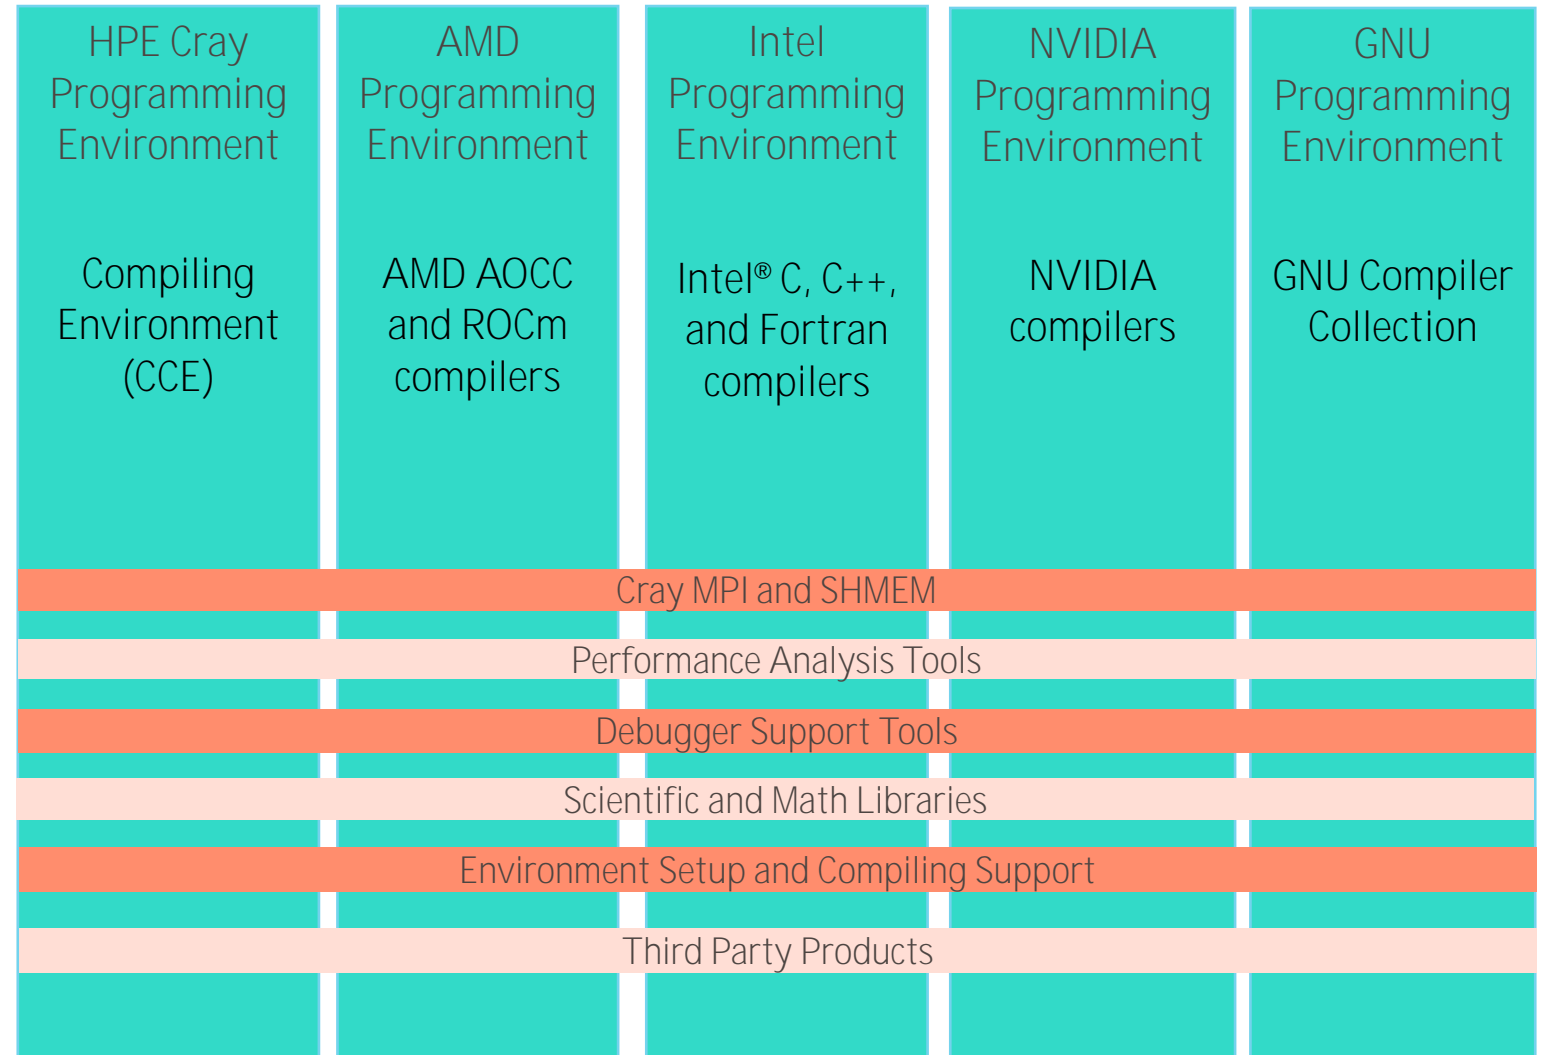
---

- The Cray Programming Environment uses a “modules” framework to support multiple software versions and to create integrated software packages
- Either Environment Modules or Lmod will be set as the default system wide
- As new versions of the supported software and associated man pages become available, they are installed and added to the Programming Environment as a new module version, while earlier versions are retained to support legacy applications
- System administrators will set the default software versions, or you can choose another version by using modules system commands
- Users can create their own modules, or administrators can install site specific modules available to many users
- Modules are used both to set high-level context (for example choose a compiler toolchain) and to select individual tool and library components and versions



# Compiler choice when using the PE

- Use **modules** to select compiling environment
- Automatically uses our math, scientific, and communication libraries with chosen compiler
- Can use debug and profiling tools with chosen compiler



# Viewing the Current Module State

- Each login session has its own module state which can be modified by loading, swapping or unloading the available *modules*
- This state affects the functioning of the compiler wrappers and in some cases runtime of applications
- A standard, default set of modules is always loaded at login for all users
- Current state can be viewed by running:

**\$> module list**



# Default Modules Example: PE modules based on version 23.09 (*year.month*)

```
% module list
```

```
Currently Loaded Modules:
```

- |                                          |                               |     |
|------------------------------------------|-------------------------------|-----|
| 1) <b>craype-x86-rome</b>                | 8) cray-dsmml/0.2.2           |     |
| 2) libfabric/1.15.2.0                    | 9) cray-mpich/8.1.27          |     |
| 3) craype-network-ofi                    | 10) cray-libsci/23.09.1.1     |     |
| 4) perftools-base/23.09.0                | 11) <b>PrgEnv-cray</b> /8.4.0 |     |
| 5) xpmem/2.5.2-2.4_3.20__gd0f7936.shasta | 12) ModuleLabel/label         | (S) |
| 6) cce/16.0.1                            | 13) lumi-tools/23.11          | (S) |
| 7) craype/2.7.23                         | 14) init-lumi/0.2             | (S) |

```
Where:
```

```
S: Module is Sticky, requires --force to unload or purge
```



# Additional CPE versions

```
% module avail cpe
```

```
----- HPE-Cray PE modules -----  
cpe/22.08    cpe/22.12    cpe/23.03    cpe/23.09 (D)
```

Where:

D: Default Module

```
% module load cpe/23.03
```

The following have been reloaded with a version change:

- 1) PrgEnv-cray/8.4.0 => PrgEnv-cray/8.3.3
- 2) cce/16.0.1 => cce/15.0.1
- 3) cray-libsci/23.09.1.1 => cray-libsci/23.02.1.1
- 4) craype/2.7.23 => craype/2.7.20
- 5) perftools-base/23.09.0 => perftools-base/23.03.0



# Viewing Available Modules

- There may be many hundreds of possible modules available to users
  - Beyond the pre-loaded defaults there are many additional packages provided by Cray
  - Sites may choose to install their own versions
- Users can see all the modules that can be loaded using the command:
  - **module avail**
- Cray PE modules will have a location in /opt, customer specific modules may also be installed.
- Searches can be narrowed by passing the first few characters of the desired module, e.g.

```
% module avail cce/
```

```
----- HPE-Cray PE modules -----  
cce/14.0.2    cce/15.0.0    cce/15.0.1    cce/16.0.1 (L,D)
```

Where:

L: Module is loaded

D: Default Module



# Further Refining Available Modules List

- **avail [switches] [sub-command [sub-command-args]]**
  - List all available modulefiles in the current **MODULEPATH**
- Useful options for filtering
  - **-d, --default**
    - List only default versions of modulefiles with multiple available versions
  - **--terse, t**
    - List in terse output format
  - **% module avail <string>**
    - List modules that include <string>



# Module spider

- Lmod includes the *spider* modules command which is very useful in searching and providing information on how to use modules
- If you use it with a specific module name and version, it will show exactly which hierarchy of modules need to be loaded first

```
% module spider cray-netcdf/4.9.0.7
```

```
-----  
cray-netcdf: cray-netcdf/4.9.0.7  
-----
```

You will need to load all module(s) on any one of the lines below before the "cray-netcdf/4.9.0.7" module is available to load.

```
LUMI/22.08  partition/C  amd/5.2.3  cray-hdf5/1.12.2.7  
LUMI/22.08  partition/C  aocc/3.2.0  cray-hdf5/1.12.2.7  
LUMI/22.08  partition/C  gcc/12.2.0  cray-hdf5/1.12.2.7  
cray-hdf5/1.12.2.7
```

```
...
```

Help:

Release info: /opt/cray/pe/netcdf/4.8.1.5/release\_info

# Module whatis info

- To find out what a module does (show whatis info)  
module whatis
- To search for a string in module name and the whatis info:  
module key

```
% module whatis cray-mpich
cray-mpich/8.1.27 : cray-mpich - Cray MPICH Message Passing Interface
```

```
% module key processor
```

```
-----
The following modules match your search criteria: "processor"
-----
```

```
...
  craype-x86-rome: craype-x86-rome
...
  craype-x86-trento: craype-x86-trento
...
-----
```

# Modules Command Output

- Be aware that module commands output to standard error
- This makes it tricky to search the (voluminous) module avail output
- csh/tcsh

```
module avail >& mavail.txt ; grep netcdf mavail.txt
( module avail >/dev/null ) |& grep netcdf
module --redirect avail | grep netcdf
```
- bash/ksh

```
module avail 2> mavail.txt ; grep netcdf mavail.txt
module avail 2>&1 | grep netcdf
module --redirect avail | grep netcdf
```
- The modules commands paginate with less(1)



# Modifying the Default Environment

Loading, swapping or unloading modules:

- The default version of any individual module can be loaded by name
  - e.g.: module load perftools-base
- A specific version can be specified after the forward slash
  - e.g.: module load perftools-base/23.09.0
- Modules can be swapped out in place
  - e.g.: module swap cce cce/15.0.1
- Or removed entirely
  - e.g.: module unload perftools

Modules will automatically change values of variables like PATH, MANPATH, LM\_LICENSE\_FILE... etc

- Modules also provide a simple mechanism for updating certain environment variables, such as PATH, MANPATH, and LD\_LIBRARY\_PATH
- In general, you should make use of the modules system rather than embedding specific directory paths into your startup files, makefiles, and scripts



# Tips for Module Usage

- Put module list in job scripts
- If you want to test which programming environment is loaded (say in a Makefile)
  - Test the PE\_ENV environment variable  
but be aware it is not supported
- If you really need to directly access something from the installed location of software (for example some build really wants to know where software is located)
  - Load (or show) the module  
module show cray-fftw
  - Look for envvars that point to the location of interest



# Summary of Useful Module Commands

Which modules are available?

- **module avail**, **module avail cce/**

Which modules are currently loaded?

- **module list**

Load software

- **module load perftools**

Change software version

- **module swap cce/16.0.1 cce/15.0.1**

Unload module

- **module unload cce**

Display module [release notes](#)

- **module help cce**

Show summary of module environment changes

- **module show cce** (or **module display cce**)



# Compiler Driver Wrappers

- All applications *that will run in parallel* on the Cray EX should be compiled with the standard language wrappers
- The compiler drivers for each language are:
  - **cc** – wrapper for the C compiler
  - **CC** – wrapper for the C++ compiler
  - **ftn** – wrapper for the Fortran compiler
- These wrappers will choose the required compiler version, target architecture options, scientific libraries and required include files automatically from the module environment.
  - They enable MPI compilation by default
- Use them exactly like you would the original compiler, e.g. to compile **prog1.f90** run  
**ftn -c prog1.f90**
- It can be necessary to set the env variables CC, CXX, FC for building tools, e.g.  
**CC=cc CXX=CC FC=ftn ./configure <flags>**





# About the `-I`, `-L` and `-l` Flags

- For libraries and include files covered by module files, you should not add anything to your Makefile
  - No additional MPI flags are needed (included by wrappers)
  - You do not need to add any `-I`, `-l` or `-L` flags for the Cray provided libraries
- If your Makefile needs an input for `-L` to work correctly, try using ‘
- If you really need a specific path, try checking ‘module show X’ for some environment variables
  - You will want to avoid a reference to a specific version as this fixes you build to that version or may create a conflict if the module environment changes



# Choosing a Programming Environment

- The wrappers choose which compiler to use from the **PrgEnv** collection loaded

PrgEnv	Description	Real Compilers
PrgEnv-cray	Cray Compilation Environment	crayftn, craycc, crayCC
PrgEnv-gnu	GNU Compiler Collection	gfortran, gcc, g++
PrgEnv-aocc	AMD Optimizing Compilers (AOCC, CPU only support)	flang, clang, clang++
PrgEnv-amd	AMD LLVM Compilers (GPU support)	amdflang, amdclang, amdclang++
PrgEnv-cray-amd	AMD Clang C/C++ compiler and the Cray Compiling Environment (CCE) Fortran compiler	
PrgEnv-gnu-amd	AMD Clang C/C++ compiler and the GNU compiler suite Fortran compiler	



# Choosing a Programming Environment

- **PrgEnv-cray** is loaded by default at login
  - use **module list** to check what is currently loaded
- List the PrgEnv- meta modules

```
> module avail PrgEnv
```

- Switch to a new programming environment

```
> module swap PrgEnv-cray PrgEnv-gnu
```

Lmod is automatically replacing "cce/16.0.1" with "gcc/12.2.0".

Due to MODULEPATH changes, the following have been reloaded:

1) **cray-mpich/8.1.27**

- The Cray MPI module is loaded by default (**cray-mpich**) and reloaded accordingly upon PrgEnv change
- Compiler wrappers will link to the proper compiler version of the modules
- Check the compiler version to know which backend you are using (e.g. **cc -v**)



# Backend Compiler Version

- Check the compiler version to know which backend you are using in the compiler wrappers
  - **--version** flag
- E.g.
  - PrgEnv-cray

```
> cc --version
Cray clang version 16.0.1
> ftn --version
Cray Fortran : Version 16.0.1
```
  - PrgEnv-amd

```
> cc --version
AMD clang version 14.0.0 (https://github.com/RadeonOpenCompute/llvm-project roc-5.2.3
22324 d6c88e5a78066d5d7a1e8db6c5e3e9884c6ad10e)
> ftn --version
AMD flang-new version 14.0.0 (https://github.com/RadeonOpenCompute/llvm-project roc-
5.2.3 22324 d6c88e5a78066d5d7a1e8db6c5e3e9884c6ad10e)
```



## PE modules: non-default modules

- The default GCC for PrgEnv-gnu is version 12.2.0
  - Wrong libs (GCC 11.2.0) are taken in `/opt/cray/pe/gcc-libs`, suggested to use GCC 11.2.0

```
module load gcc/11.2.0
```

- New PE rocm module based on version 5.2.3
  - Newer modules are available provided by LUST (check via `module spider rocm`)
- If you want to use non-default modules, set LD\_LIBRARY\_PATH before running the application

```
export LD_LIBRARY_PATH=${CRAY_LD_LIBRARY_PATH}:${LD_LIBRARY_PATH}
```

- `${CRAY_LD_LIBRARY_PATH}` is set by the PE modules
- use the `ldd` command to check the linked libraries of you application executable



# Compiler Versions

There are usually multiple versions of each compiler available to users.

- The most recent version is usually the default and will be loaded when swapping PrgEnvs.
- To change the version of the compiler in use, swap the Compiler Module.  
eg `module swap cce/16.0.1 cce/15.0.1`

PrgEnv	Compiler Module
PrgEnv-cray	cce
PrgEnv-gnu	gcc
PrgEnv-aocc	aocc
PrgEnv-amd	amd

- Note that you may not be able to swap vastly different versions without also swapping other modules
  - Example with PrgEnv-gnu loaded  
> `module swap gcc cce`

Lmod is automatically replacing "PrgEnv-gnu/8.4.0" with "PrgEnv-cray/8.4.0".

Due to MODULEPATH changes, the following have been reloaded:

1) `cray-mpich/8.1.27`



# OpenMP

- OpenMP is supported by all of the **PrgEnvs**

PrgEnv	Enable OpenMP
PrgEnv-cray (Fortran)	-homp / -fopenmp
PrgEnv-cray (C/C++)	-fopenmp
PrgEnv-gnu	-fopenmp
PrgEnv-aocc	-fopenmp
PrgEnv-amd	-fopenmp



# Compiler Man Pages and Information

- For more information on individual compilers

PrgEnv	C	C++	Fortran
PrgEnv-cray	man craycc	man crayCC	man crayftn
PrgEnv-gnu	man gcc	man g++	man gfortran
PrgEnv-aocc	clang --help	clang++ --help	flang --help
PrgEnv-amd	amdclang --help	amdclang++ --help	amdflang --help
Wrappers	man cc	man CC	man ftn

- All compilers provide --help output
  - E.g. for PrgEnv-cray: **crayftn --help**
- To verify that you are using the correct version of a compiler, use:
  - -V option on a cc, CC, or ftn for cray Cray
  - -dumpversion option on a cc, CC, or ftn command with GNU, AOCC, AMD





# Recap

---

- The HPE Cray Programming Environment provides a consistent interface into a host of user and developer software
- Various toolchains are supported: Cray compilers, AMD, GNU, etc.
- Different architectures of CPU and GPU are supported
- The environment is used via a combination of
  - Modules
  - Compiler wrappers





# Questions?