



**Hewlett Packard
Enterprise**

Advanced Placement

Comprehensive General LUMI Course
April 23–26, 2024

Goal: Speed Up Your Application Without Changing or Recompiling Code (1/3)

*“Never underestimate the bandwidth of a station wagon
full of tapes hurtling down the highway”*

— Andrew S. Tanenbaum

Extending the metaphor, this talk is about configuring a “GPS” to:

1. Use the fastest routes from point A to point B (reduce latency)
2. Avoid traffic jams (hitting hardware limitation due to link saturation)



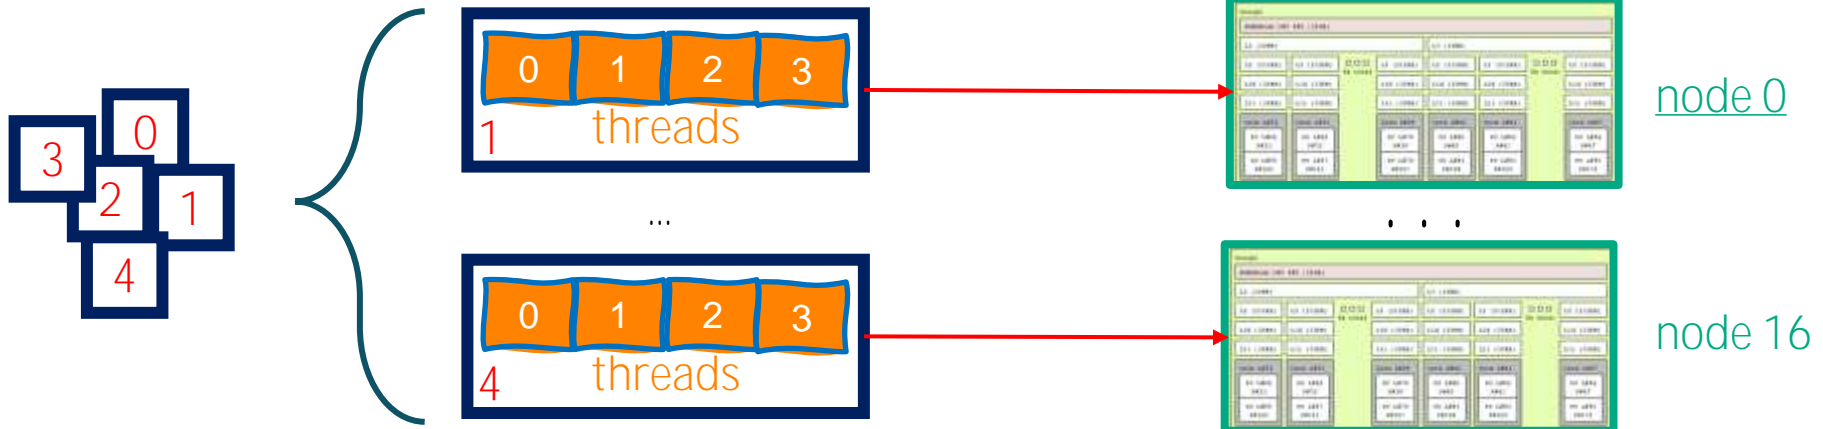
Goal: Speed Up Your Application Without Changing or Recompiling Code (2/3)

- Task **distribution (and ordering)**: distribution of MPI tasks among the nodes.
- Task **affinity**: assign each MPI task a set of CPU cores.
- Thread **binding**: pin/attach each thread in each MPI task to one or many cores.

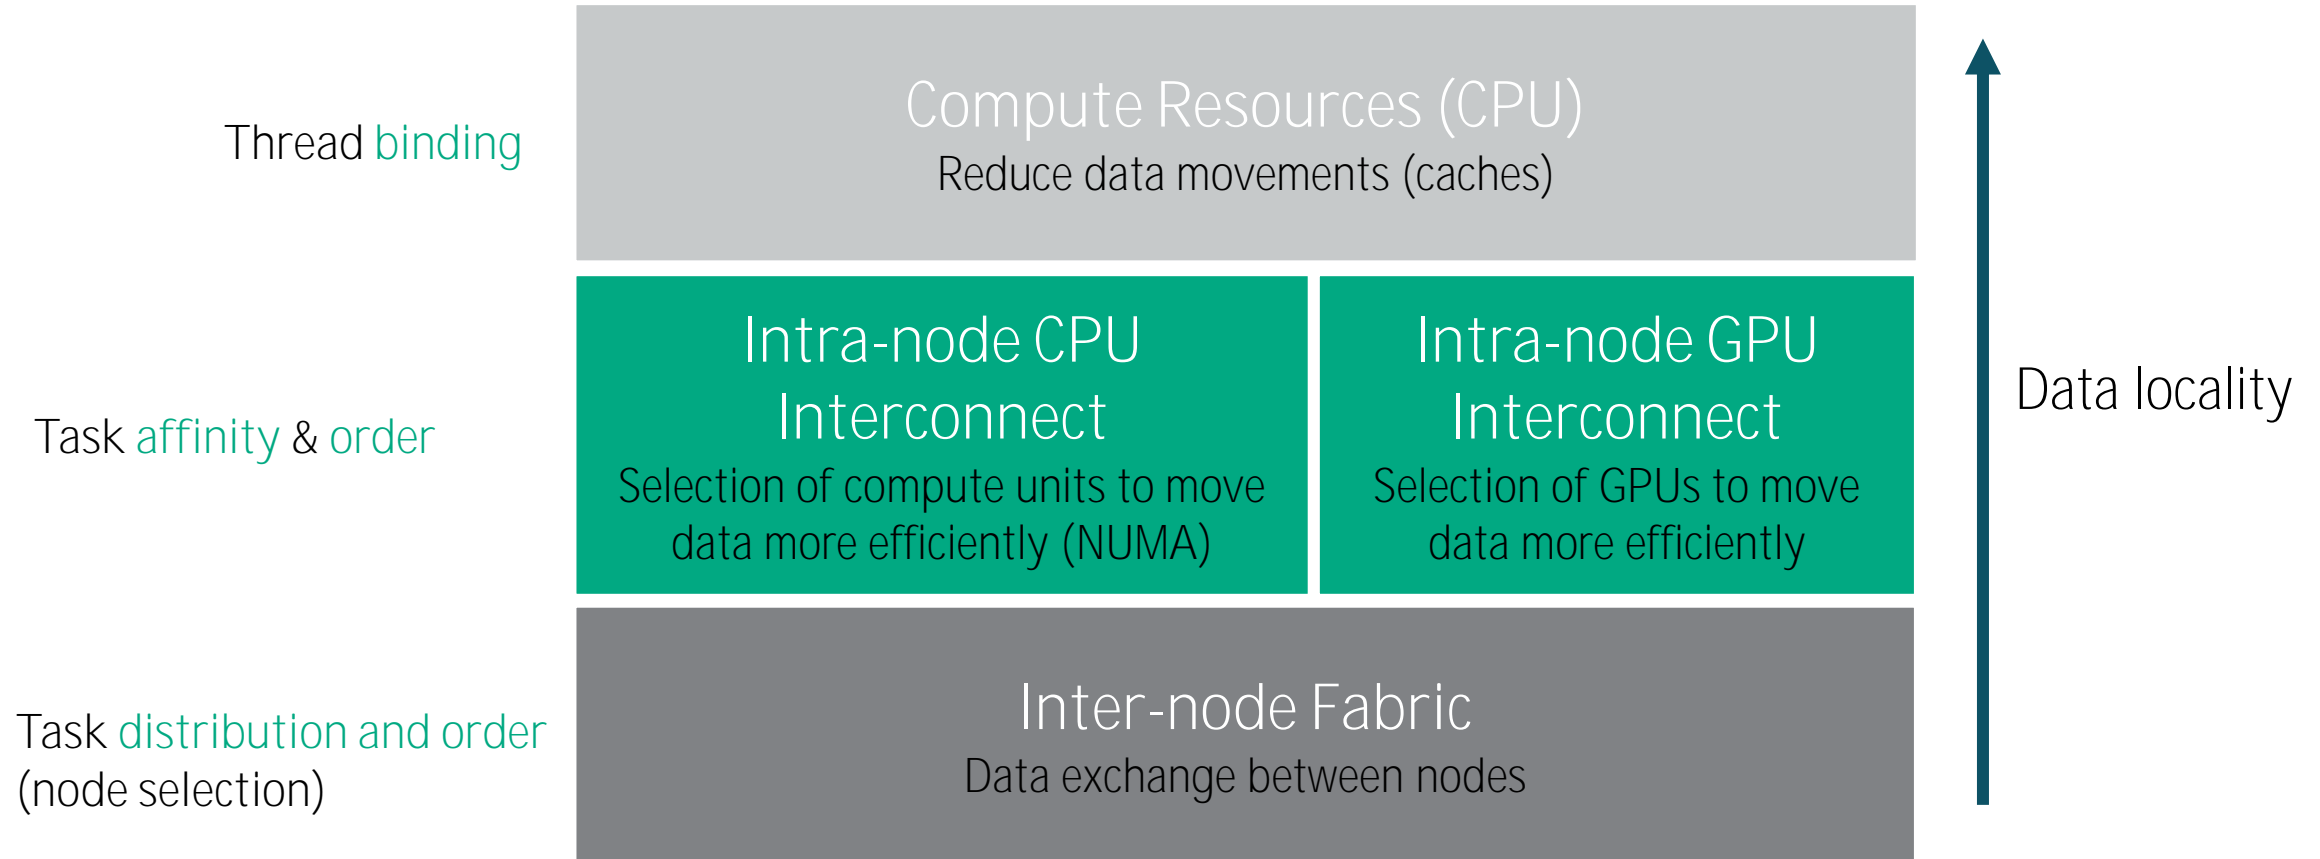
1 JOB =
N MPI TASKS

1 MPI TASK =
N THREADS

1 CLUSTER =
N NODES



Goal: Speed Up Your Application Without Changing or Recompiling Code (3/3)



Agenda

- Placement Check tools
- Inter-node Placement (fabric)
- Intra-node Placement (memory affinity)
- Compute Resources Affinity (OpenMP thread binding)
- Intra-node Placement for GPUs





PLACEMENT CHECK TOOLS

Tools and options to check distributions and bindings



PLACEMENT CHECK TOOLS – hybrid_check

```
module load LUMI/23.09
module load partition/C
module load lumi-CPEtools
srun -c2 --nodes=1 --tasks-per-node=8 hybrid_check
```

Running 8 MPI ranks with 2 threads each (total number of threads: 16).

```
++ hybrid_check: MPI rank    0/8    OpenMP thread    0/2    on cpu    0/256 of nid002052
++ hybrid_check: MPI rank    0/8    OpenMP thread    1/2    on cpu    1/256 of nid002052
++ hybrid_check: MPI rank    1/8    OpenMP thread    0/2    on cpu    2/256 of nid002052
++ hybrid_check: MPI rank    1/8    OpenMP thread    1/2    on cpu    3/256 of nid002052
++ hybrid_check: MPI rank    2/8    OpenMP thread    0/2    on cpu    4/256 of nid002052
++ hybrid_check: MPI rank    2/8    OpenMP thread    1/2    on cpu    5/256 of nid002052
++ hybrid_check: MPI rank    3/8    OpenMP thread    0/2    on cpu    6/256 of nid002052
++ hybrid_check: MPI rank    3/8    OpenMP thread    1/2    on cpu    7/256 of nid002052
++ hybrid_check: MPI rank    4/8    OpenMP thread    0/2    on cpu    64/256 of nid002052
++ hybrid_check: MPI rank    4/8    OpenMP thread    1/2    on cpu    65/256 of nid002052
++ hybrid_check: MPI rank    5/8    OpenMP thread    0/2    on cpu    67/256 of nid002052
++ hybrid_check: MPI rank    5/8    OpenMP thread    1/2    on cpu    66/256 of nid002052
++ hybrid_check: MPI rank    6/8    OpenMP thread    0/2    on cpu    68/256 of nid002052
++ hybrid_check: MPI rank    6/8    OpenMP thread    1/2    on cpu    69/256 of nid002052
++ hybrid_check: MPI rank    7/8    OpenMP thread    0/2    on cpu    70/256 of nid002052
++ hybrid_check: MPI rank    7/8    OpenMP thread    1/2    on cpu    71/256 of nid002052
```


PLACEMENT CHECK TOOLS – gpu_check

```
module load LUMI/23.09
module load partition/G
module load lumi-CPEtools
srun -n 8 -c 7 gpu_check
```

```
MPI 000 - OMP 000 - HWT 001 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 000 - OMP 001 - HWT 002 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 000 - OMP 002 - HWT 003 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 000 - OMP 003 - HWT 004 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 000 - OMP 004 - HWT 005 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 000 - OMP 005 - HWT 006 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 000 - OMP 006 - HWT 007 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc

MPI 001 - OMP 000 - HWT 009 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 001 - OMP 001 - HWT 010 - Node nid005014 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
```

...



PLACEMENT CHECK TOOLS – Summary and Other Options

- You should be aware **WHAT is doing** the binding for you!
- **Various software** components may try do this for you
 - WLM (SLURM, ALPS, ...)
 - MPI (MVAPICH, ...)
 - Compiler (CCE,GNU,...)
 - OpenMP
 - Tools (numactl, taskset)

Cheat sheet

| | |
|--|--|
| <code>m1 lumi-CPEtools; srun ... hybrid_check</code> | Show hybrid MPI/OpenMP placement |
| <code>m1 lumi-CPEtools; srun ... gpu_check</code> | Show hybrid MPI/OpenMP placement + GPU affinities |
| <code>srun --cpu-bind=verbose,<mode></code> | Report task affinity set by Slurm |
| <code>export OMP_DISPLAY_AFFINITY=TRUE</code> | Show OpenMP affinities |
| <code>export MPICH_CPUMASK_DISPLAY=1</code> | Report MPI task affinities |
| <code>export MPICH_RANK_REORDER_DISPLAY=1</code> | Show MPI Rank reordering |
| <code>export MPICH_OFI_NIC_VERBOSE=2</code> | Show information about network interface selection |





INTER-NODE PLACEMENT

Distributing tasks between nodes to better harness the fabric



INTER-NODE PLACEMENT – What can be done?

Selection of nodes (almost no control)

Slurm allocates sets of nodes which are roughly consecutive (likely to be in the same cabinet with a few nodes).

Task placement on allocated nodes

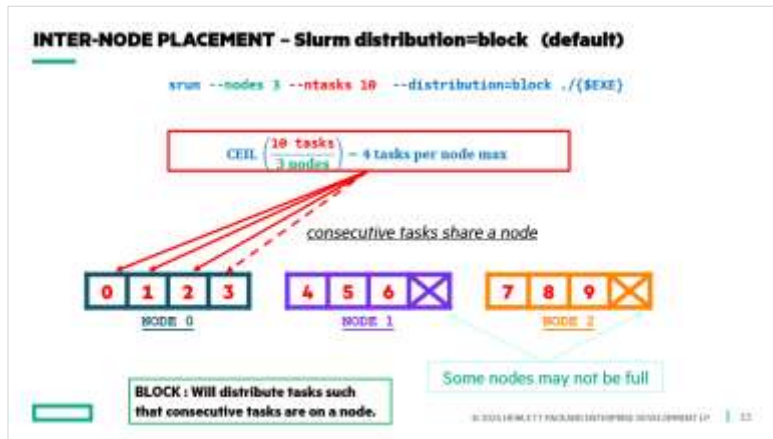
If possible, place ranks exchanging highest volume of data on the same compute. Two complementary options:

- Distribution pattern: Coarse grain control on placement with Slurm (srun **--distribution** argument).
- Reordering of MPI tasks (***MPICH_RANK_REORDER_METHOD***): More details during the ***MPI rank reordering*** talk and hints provided by the Perftools.

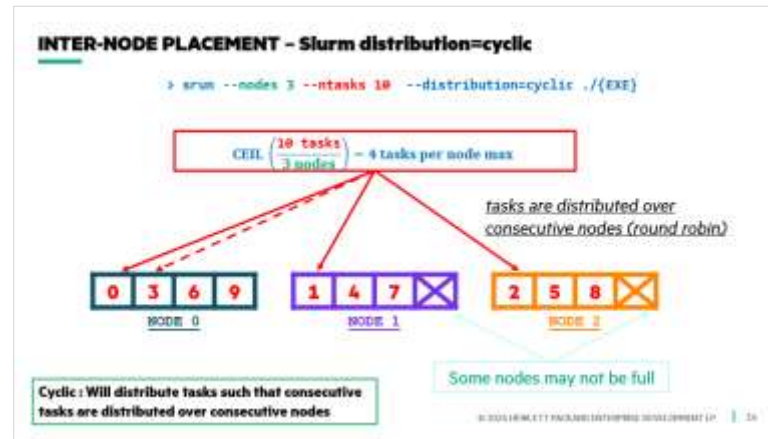


INTER-NODE PLACEMENT- Distributing tasks between nodes (coarse grain control)

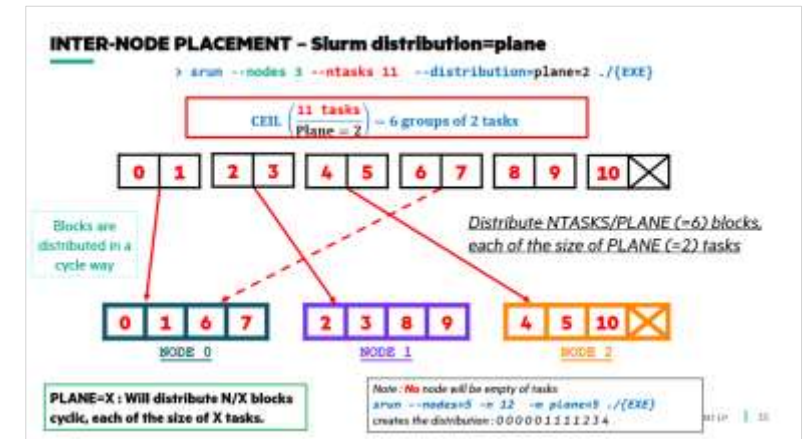
- To control the distribution of the MPI ranks (tasks) across nodes
- In *srun*: use the '**--distribution/-m**' with either { block | cyclic | plane=<size> }



--distribution=block
(default)



--distribution=cyclic




--distribution=plane=X

INTER-NODE PLACEMENT – Slurm distribution=block (default)

```
srun --nodes 3 --ntasks 10 --distribution=block ./{$EXE}
```

$$\text{CEIL} \left(\frac{10 \text{ tasks}}{3 \text{ nodes}} \right) = 4 \text{ tasks per node max}$$




BLOCK : Will distribute tasks such that consecutive tasks are on a node.

Some nodes may not be full

INTER-NODE PLACEMENT – Slurm distribution=cyclic

```
> srun --nodes 3 --ntasks 10 --distribution=cyclic ./{EXE}
```

$$\text{CEIL} \left(\frac{10 \text{ tasks}}{3 \text{ nodes}} \right) = 4 \text{ tasks per node max}$$




Cyclic : Will distribute tasks such that consecutive tasks are distributed over consecutive nodes

Some nodes may not be full

INTER-NODE PLACEMENT – Slurm distribution=plane

```
> srun --nodes 3 --ntasks 11 --distribution=plane=2 ./{EXE}
```

$$\text{CEIL} \left(\frac{11 \text{ tasks}}{\text{Plane} = 2} \right) = 6 \text{ groups of 2 tasks}$$




Blocks are distributed in a cycle way


Distribute $NTASKS/PLANE (=6)$ blocks, each of the size of $PLANE (=2)$ tasks



NODE 0



NODE 1



NODE 2

PLANE=X : Will distribute N/X blocks cyclic, each of the size of X tasks.

Note : No node will be empty of tasks
`srun --nodes=5 -n 12 -m plane=5 ./{EXE}`
creates the distribution : 0 0 0 0 0 1 1 1 1 2 3 4

INTER-NODE PLACEMENT – Takeaways

- Keep in mind variations in network performance might be observed. Slurm may allocate nodes which are not in the same cabinet.
- Default Slurm distribution across nodes (block) is well suited in most cases.
- MPI rank reordering might be useful to reduce pressure on the fabric (perftools may provide hints)

Cheat sheet

| | |
|---|---|
| <code>salloc/sbatch -x <nodelist></code> | Exclude specific nodes |
| <code>srun --distribution=block cyclic plane=<X></code> | Modify task distribution pattern across nodes |
| <code>export MPICH_RANK_REORDER_METHOD=3</code> <code>export MPICH_RANK_REORDER_FILE=<file></code> | Manually defining rank reordering (more details in another presentation) |





INTRA-NODE PLACEMENT (NUMA NODES)

Distributing tasks to better harness shared memory communications (node interconnect)



INTRA-NODE PLACEMENT – Lumi-G and Lumi-C Nodes


Accelerator blade



Module GPU - HPE Cray EX235a

- Bard Peak Module
 - 2 nodes per module
- Each Bard Peak Node
 - 1 processor AMD Trento 64 cores
 - 4 GPUs AMD MI250x
 - 4 network interfaces Slingshot-11 (200Gbps each)

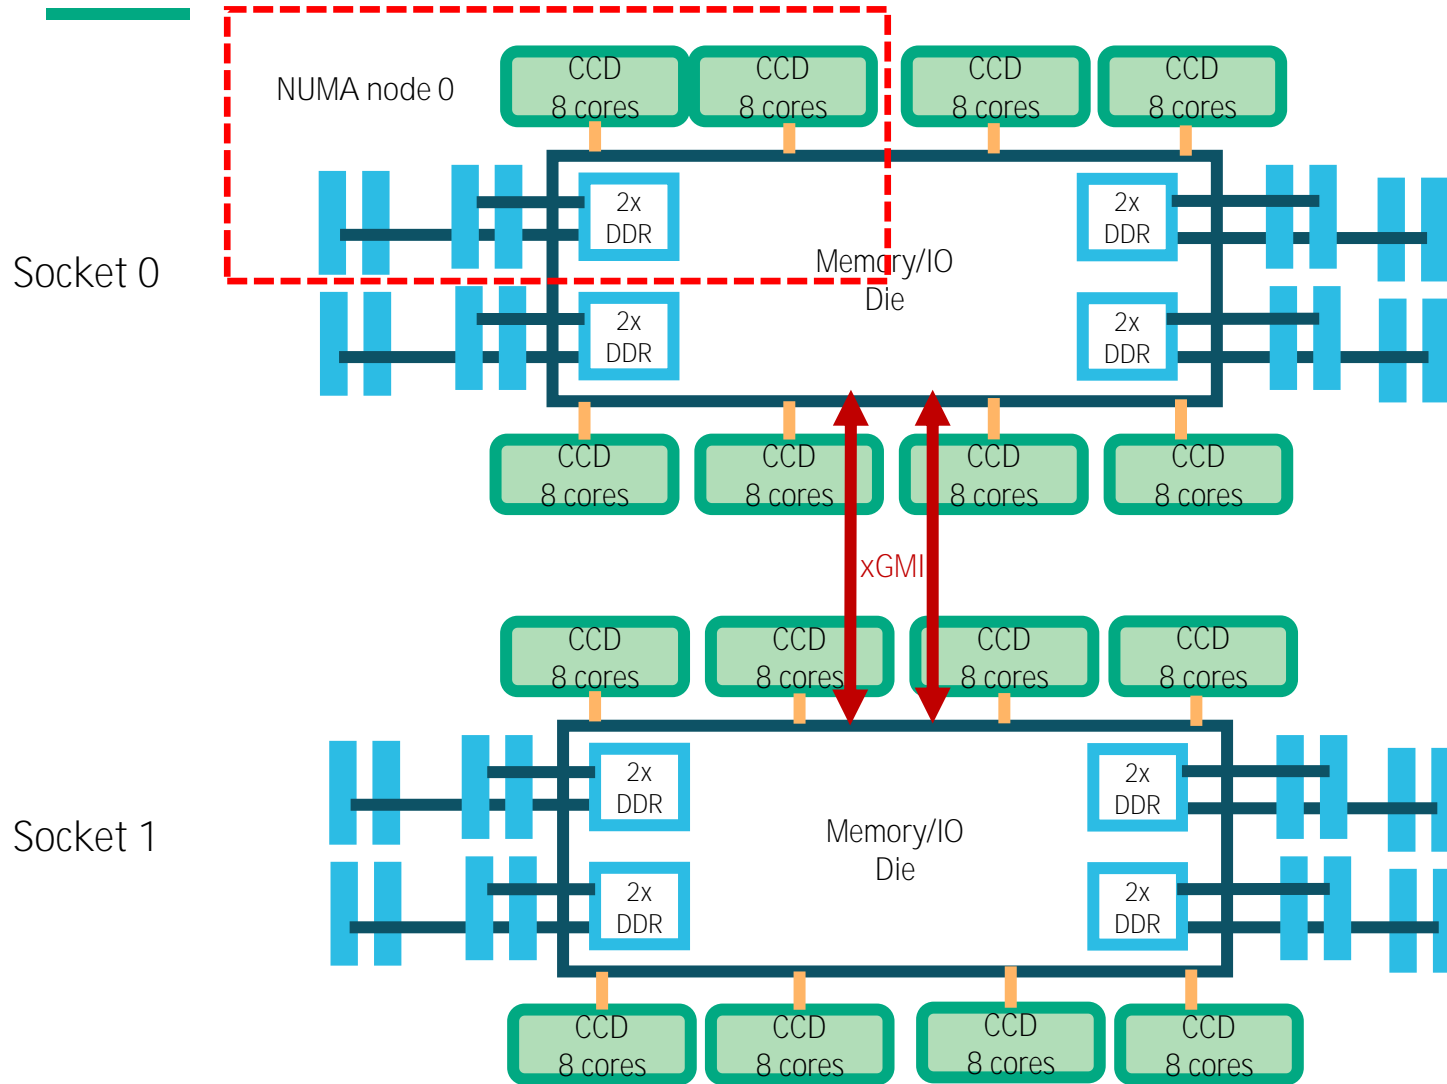
Compute blade



Module CPU - HPE Cray EX425

- Antero Module
 - 4 nodes per module
- Each Antero Node
 - 2 processors AMD Milan 64 cores
 - 1 network interface Slingshot-11 (200Gbps)


INTRA-NODE PLACEMENT – Interconnect With Two CPU Sockets



- CPUs in a compute node can be seen as a small cluster of interconnect and cache coherent compute groups called NUMA (Non-Uniform Memory Access) nodes.
- A NUMA node is composed of a set of cores associated with a memory controller and dedicated affinity to devices such as network interface(s) and/or GPU(s).
- Each LUMI-C node is divided into eight NUMA nodes.

INTRA-NODE PLACEMENT – Potential Performance Issues

Quick reminder: MPI tasks in the same node may exchange data via shared memory communications.



Interconnect

Higher latency:

When fetching data in a different NUMA node (“**NUMA effects**”), especially on a different CPU socket (higher “**NUMA distance**”).

Saturation:

Heavily using the interconnect (additional hops or link saturation) leads to suboptimal utilization of memory / IO throughputs.

Memory

Saturation:

Collocating more tasks on one NUMA node may lead to unbalanced memory usage (performance and capacity).

INTRA-NODE PLACEMENT – What Can Be Done?

Different strategies depending on the workload. You may want to profile the application first.

A few guidelines:

- Task placement limiting NUMA effects: If possible, place ranks exchanging highest volume of data on the same NUMA node or NUMA nodes with lowest NUMA distance (same CPU socket). Reordering MPI tasks might be an option.
- Avoid making each task affinity spread across multiple NUMA nodes (without specific reason): by default, data get assigned to a NUMA node on first access (first touch policy) based on NUMA the CPU core used. Using multiple NUMA nodes with one task may increase NUMA effects.
- Task placement with balanced NUMA node usage: use same number of tasks per NUMA node to avoid saturating one memory controller and slow down the whole application.

Controlling Intra-Node Task Placement

- Manually reordering MPI tasks
- Coarse grain control using the `--distribution` (*srun*) argument (second level) and `-c` to define the quantity of cores/hyperthreads per task.
- Fine grain control using a CPU map (`--cpu-bind=map_cpu`) or a CPU mask (`--cpu-bind=mask_cpu`) to select CPU cores/hyperthreads accessible by each task.

INTRA-NODE PLACEMENT – Getting Node Information, lscpu

```
lscpu | grep -Ei "CPU\(s\)|Thread|Core\(s\)|Socket\(s\)|Numa|Model\ name| MHz|cache"
```

| | |
|----------------------|----------------|
| CPU(s) : | 128 |
| On-line CPU(s) list: | 0-127 |
| Thread(s) per core: | 2 |
| Core(s) per socket: | 64 |
| Socket(s) : | 1 |
| NUMA node(s) : | 4 |
| Model name: | AMD EPYC 7A53 |
| CPU MHz: | 1925.019 |
| CPU max MHz: | 3541.0149 |
| CPU min MHz: | 1500.0000 |
| L1d cache: | 32K |
| L1i cache: | 32K |
| L2 cache: | 512K |
| L3 cache: | 32768K |
| NUMA node0 CPU(s) : | 0-15, 64-79 |
| NUMA node1 CPU(s) : | 16-31, 80-95 |
| NUMA node2 CPU(s) : | 32-47, 96-111 |
| NUMA node3 CPU(s) : | 48-63, 112-127 |

- Check **/proc/cpuinfo** ON the compute nodes
- If **lscpu** is installed on a system, it will list the configuration
- Hyperthreading (aka SMT) is turned ON
 - From a binding point of view, here we have CORES=CPUs

LUMI-G



```
lstopo-no-graphics -.ascii --only pu
```

INTRA-NODE PLACEMENT – Getting Node Information, numactl

```
numactl --hardware | grep -E "nodes|cpus"
```

```
available: 4 nodes (0-3)
```

```
node 0 cpus: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 64 65 66 67 68 69 70 71 72 73 74 75 76 ...
```

```
node 1 cpus: 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 80 81 82 83 84 85 86 87 88 89 ...
```

```
...
```

Each NUMA NODE (4 in total) has

- 16 physical core
- Each core has hardware threads – aka Hyperthreads or Simultaneous MultiThreading (SMT)


```
lscpu | grep NUMA
```

```
NUMA node(s): 4
NUMA node0 CPU(s): 0-15, 64-79
NUMA node1 CPU(s): 16-31, 80-95
NUMA node2 CPU(s): 32-47, 96-111
NUMA node3 CPU(s): 48-63, 112-127
```

INTRA-NODE PLACEMENT – Hyperthreads and Numbering

- The numbering of **Lumi-C**
 - Actual cores from 0-127
 - Hyperthreads from 128-255.
- It is not mandatory to use the hyperthreads (disabled by default)
 - The hyperthreads are still there but not utilized.


Option for salloc and srun
`--hint=[no]multithread`



CPU 0-127 are available,
 CPU 128-255 are ignored

INTRA-NODE PLACEMENT – Low Noise Mode Configuration (1/2)

- The LUMI-G compute nodes have the low-noise mode activated.
 - Helps reduce jitter and variability from run to run
 - This mode reserve 1 core of each CCD to the operating system (8 cores in total)



...

INTRA-NODE PLACEMENT – Low Noise Mode Configuration (2/2)

Jobs requesting > 56 cores (112 SMT threads) per node will never run.

➔ Only 56 cores / 112 SMT threads available for user applications

```
> srun -p small-g --nodes=1 --hint=nomultithread sh -c 'echo $SLURM_JOB_CPUS_PER_NODE'
56
```

```
> srun -p small-g --nodes=1 --hint=multithread sh -c 'echo $SLURM_JOB_CPUS_PER_NODE'
112
```

```
> srun -p small-g --nodes=1 --hint=multithread --cpus-per-task=113 ./myApp
```

```
srun: error: Unable to allocate resources: Requested node configuration is not available
```

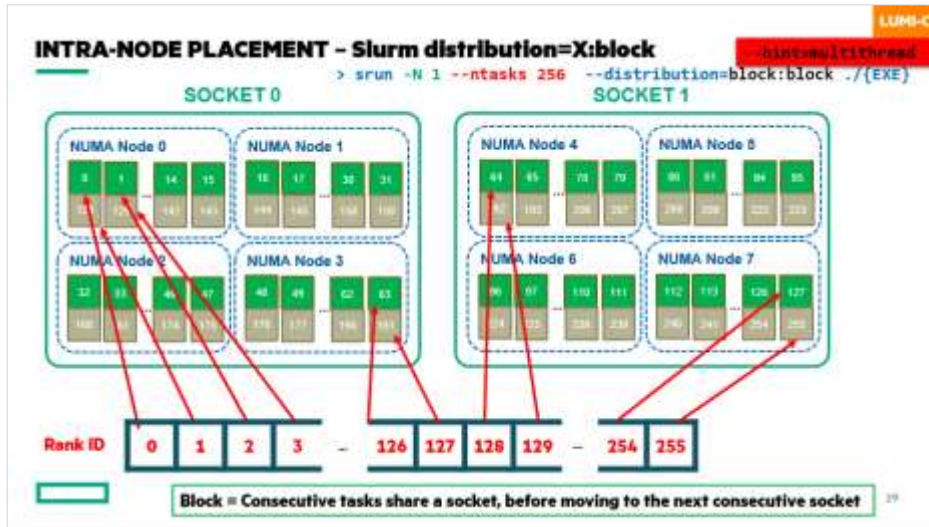
- In general, applications are bandwidth bound (memory or transfers with GPUs)



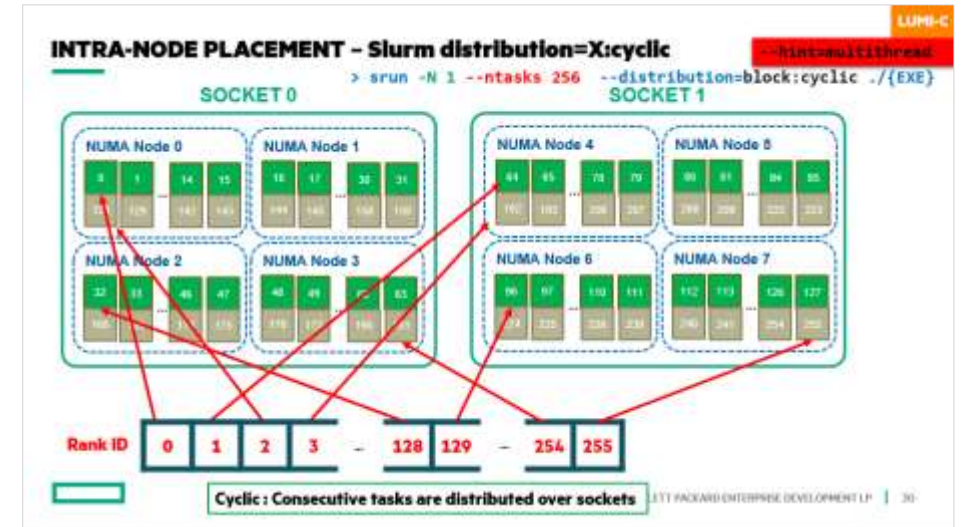
INTRA-NODE PLACEMENT – Slurm task distribution (level 2)

For the second distribution method, the ranks collected in a node in the first distribution step, can be distributed over:

- Sockets on LUMI-C
- NUMA nodes on LUMI-G (NUMA nodes as declared as sockets on LUMI-G)



--distribution=[block|cyclic]:block
consecutive tasks share a CPU socket



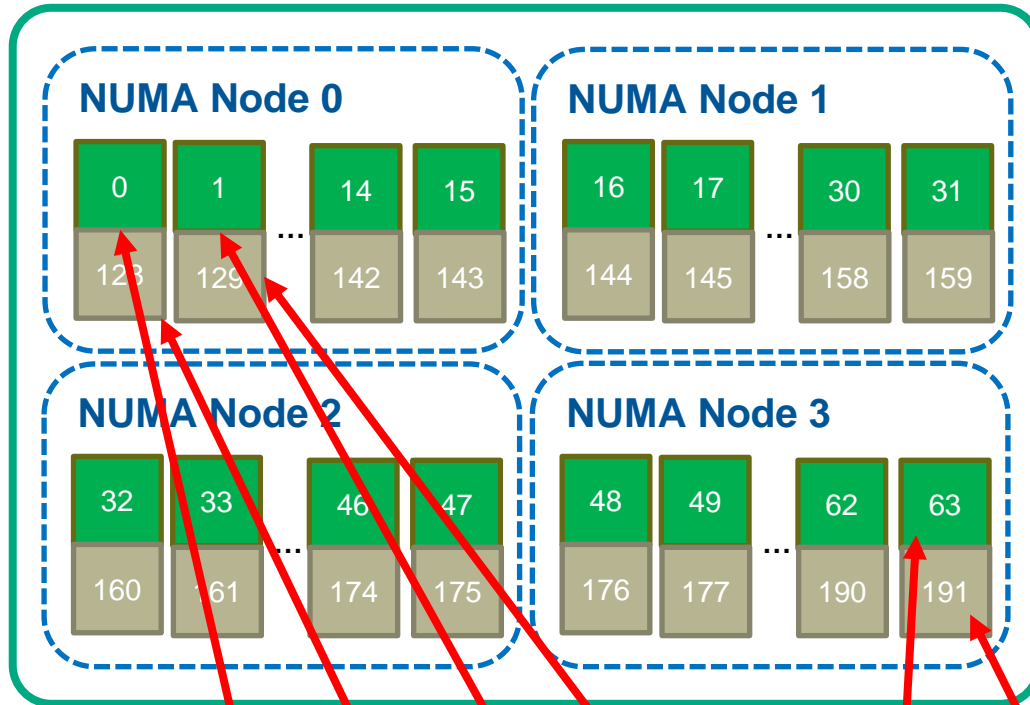
--distribution=[block|cyclic]:cyclic
consecutive tasks are distributed over CPU sockets

INTRA-NODE PLACEMENT – Slurm distribution=X:block

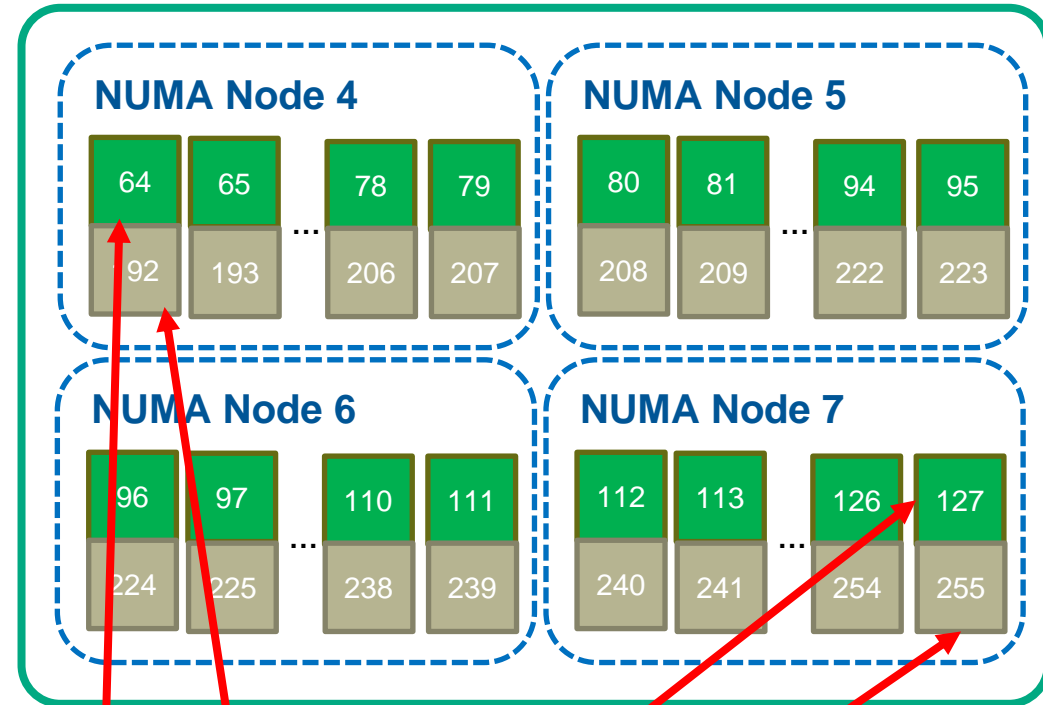
```
--hint=multithread
```

```
> srun -N 1 --ntasks 256 --distribution=block:block ./{EXE}
```

SOCKET 0



SOCKET 1



Rank ID



Block = Consecutive tasks share a socket, before moving to the next consecutive socket

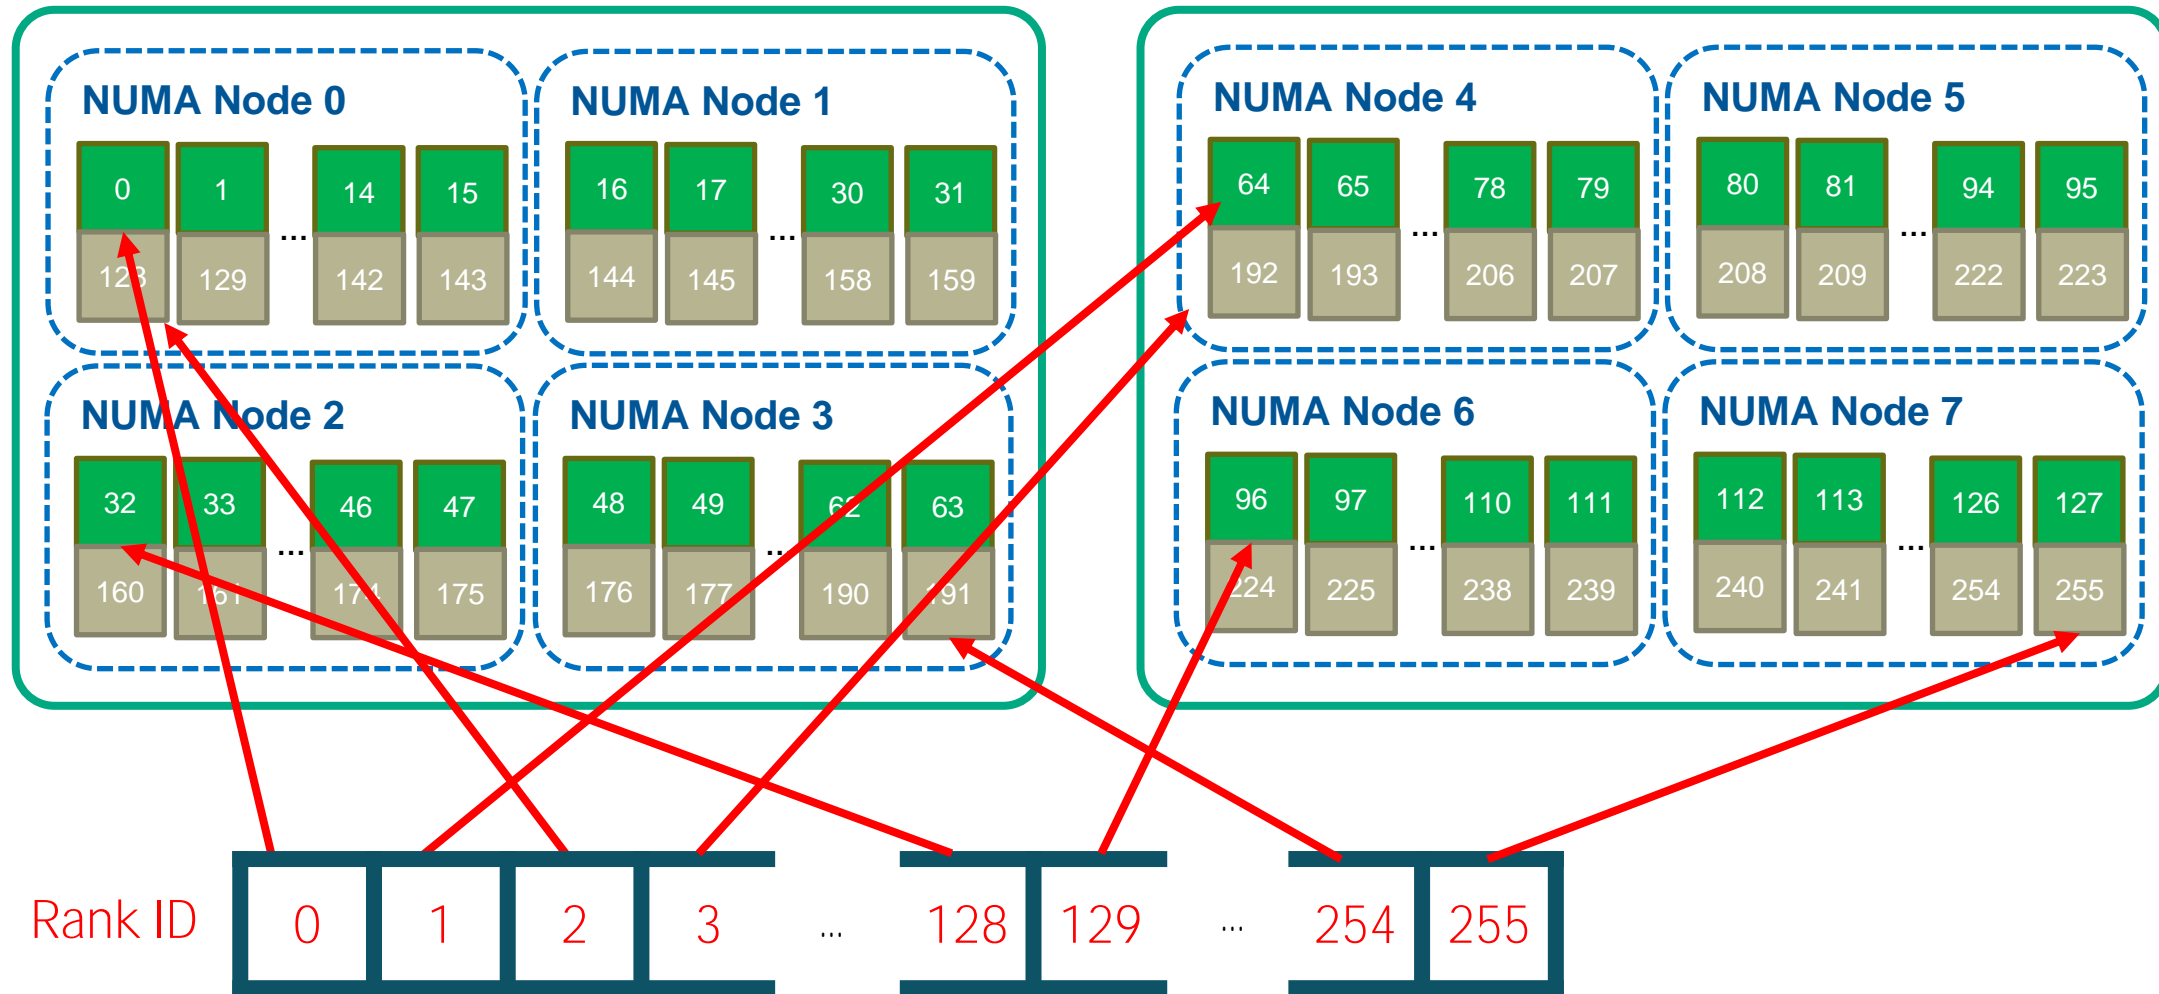
INTRA-NODE PLACEMENT – Slurm distribution=X:cyclic

```
--hint=multithread
```

```
> srun -N 1 --ntasks 256 --distribution=block:cyclic ./{EXE}
```

SOCKET 0

SOCKET 1

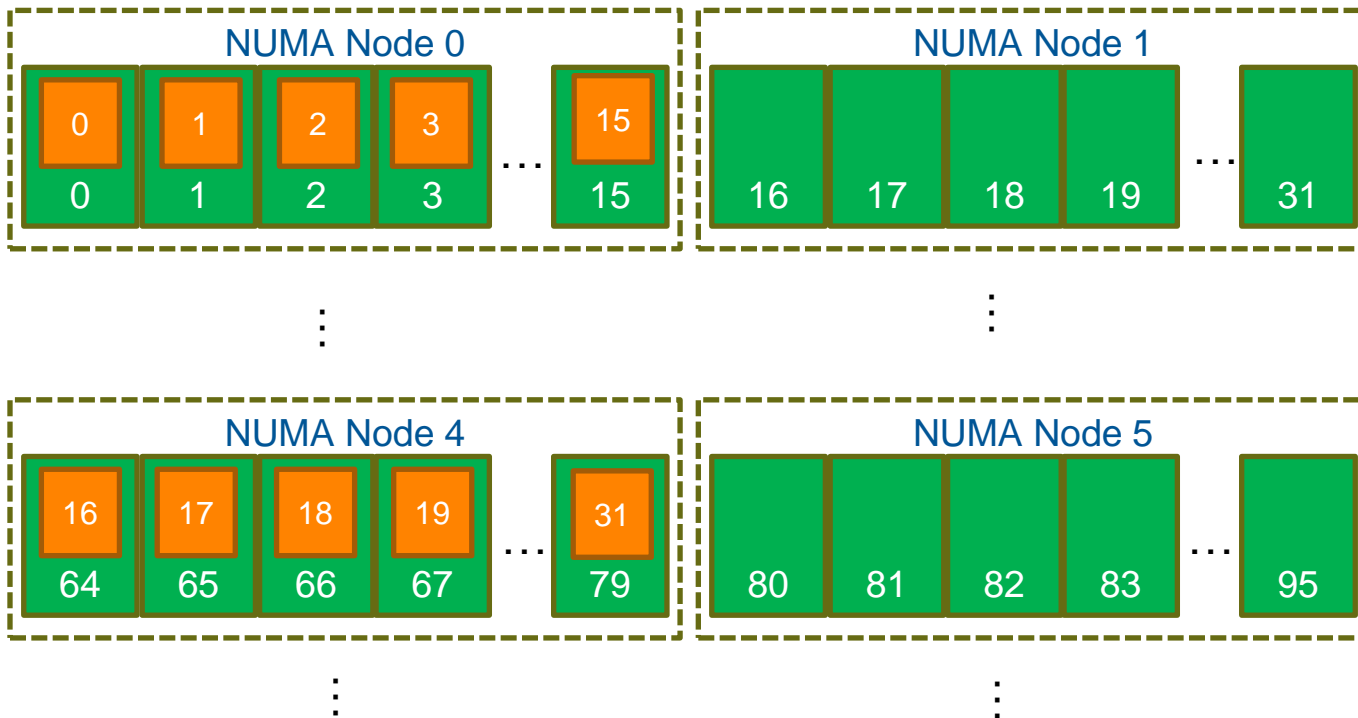


Cyclic : Consecutive tasks are distributed over sockets

INTRA-NODE PLACEMENT – Specifying Quantity of Tasks per Socket

- The number of tasks per socket can be limited

```
srun --nodes 1 -n 32 --ntasks-per-socket=16 --hint=nomultithread ./${EXE}
```




| host | rank | thr | count | mask |
|-----------|------|-----|-------|--------|
| nid002459 | 0 | 0 | 1 | cpu 0 |
| | 1 | 0 | 1 | cpu 1 |
| ... | | | | |
| | 14 | 0 | 1 | cpu 14 |
| | 15 | 0 | 1 | cpu 15 |
| | 16 | 0 | 1 | cpu 64 |
| | 17 | 0 | 1 | cpu 65 |
| ... | | | | |
| | 30 | 0 | 1 | cpu 78 |
| | 31 | 0 | 1 | cpu 79 |

INTRA-NODE PLACEMENT – Custom Mapping, One Core per Task

- The user may explicitly specify a core/hyperthread accessible by each task on the node.
- The mapping is the same for all nodes
- Useful when underpopulating a node to access more memory bandwidth per task

```
export bind=0,2,16,18,32,34,48,50  
srun -N 1 -n 8 --cpu-bind=map_cpu:${bind} ./${EXE}
```

Only makes sense if the
node is reserved in
--exclusive mode



INTRA-NODE PLACEMENT – Custom Mapping, Multiple Cores per Task (1/2)

- The user may explicitly specify a CPU mask to define multiple cores/hyperthreads accessible by each task on the node.
- The mapping is the same for all nodes
- A mask is a bitmap. When converted to binary, the position (from right to left hand side) of the ones defines the core IDs. For instance: 0xFE -> 1111 1110 (core ID 0 would be skipped and cores from ID 1 to ID 7 selected)
- Following aliases (add to your .bashrc) might be useful for conversions:

```
# Convert hexa binary
0x () {
    local val=$(tr '[:a-z:]' '[:A-Z:]' <<< $1)
    echo "binary: `BC_LINE_LENGTH=0 bc <<< \"ibase=16;obase=2;$val\"`"
}

# Convert binary to hexa
0b () {
    local val=$(tr '[:a-z:]' '[:A-Z:]' <<< $1)
    echo "hexa: `BC_LINE_LENGTH=0 bc <<< \"ibase=2;obase=16;$val\"`"
}
```

Example:

```
> 0x ef
```

```
binary: 11101111
```

```
> 0b 11101111
```

```
hexa: EF
```


INTRA-NODE PLACEMENT – Takeaways

- Understand node hierarchy and core to NUMA node mapping to
 - Reduce NUMA effects
 - Avoid unbalances (overloading one specific NUMA node)
- 8 cores are not available on LUMI-G nodes

Cheat sheet

| | |
|---|---|
| <code>lscpu, lstopo, numactl -hardware</code> | Display node information |
| <code>salloc sbatch --hint=[no]multithread</code> | Use or hide hardware threads (hyperthreads) |
| <code>srun --distribution=X:block cyclic</code> | Modify task distribution pattern across CPU sockets |
| <code>srun --ntasks-per-socket=X</code> | Specify how many tasks per socket |
| <code>srun --cpu-bind=map_cpu:<core_list></code> | Provide a CPU core affinity per task |
| <code>srun -c X --cpu-bind=mask_cpu:<mask_list></code> | Provide multiple CPU core affinities per task |
| <code>export MPICH_RANK_REORDER_METHOD=3</code> <code>export MPICH_RANK_REORDER_FILE=<file></code> | Manually defining rank reordering (more details in another presentation) |

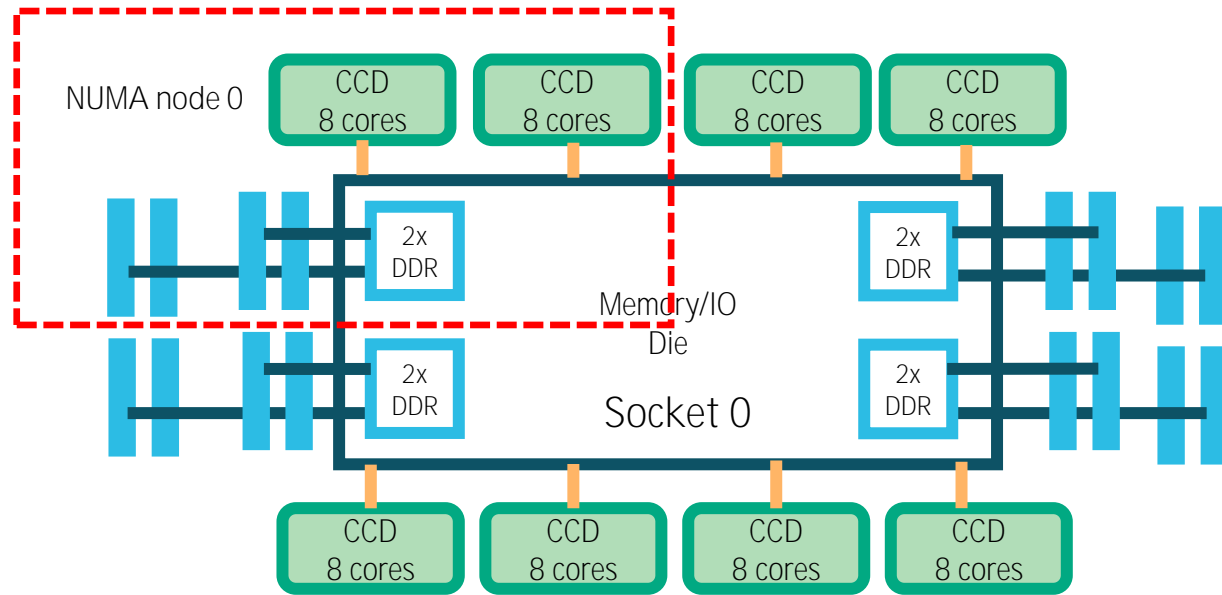





COMPUTE RESOURCES PLACEMENT (CPU CORE AFFINITY)

Hybrid MPI/OpenMP, thread binding to reduce data movements

CORE AFFINITY – AMD CPU, CCD Zoom in



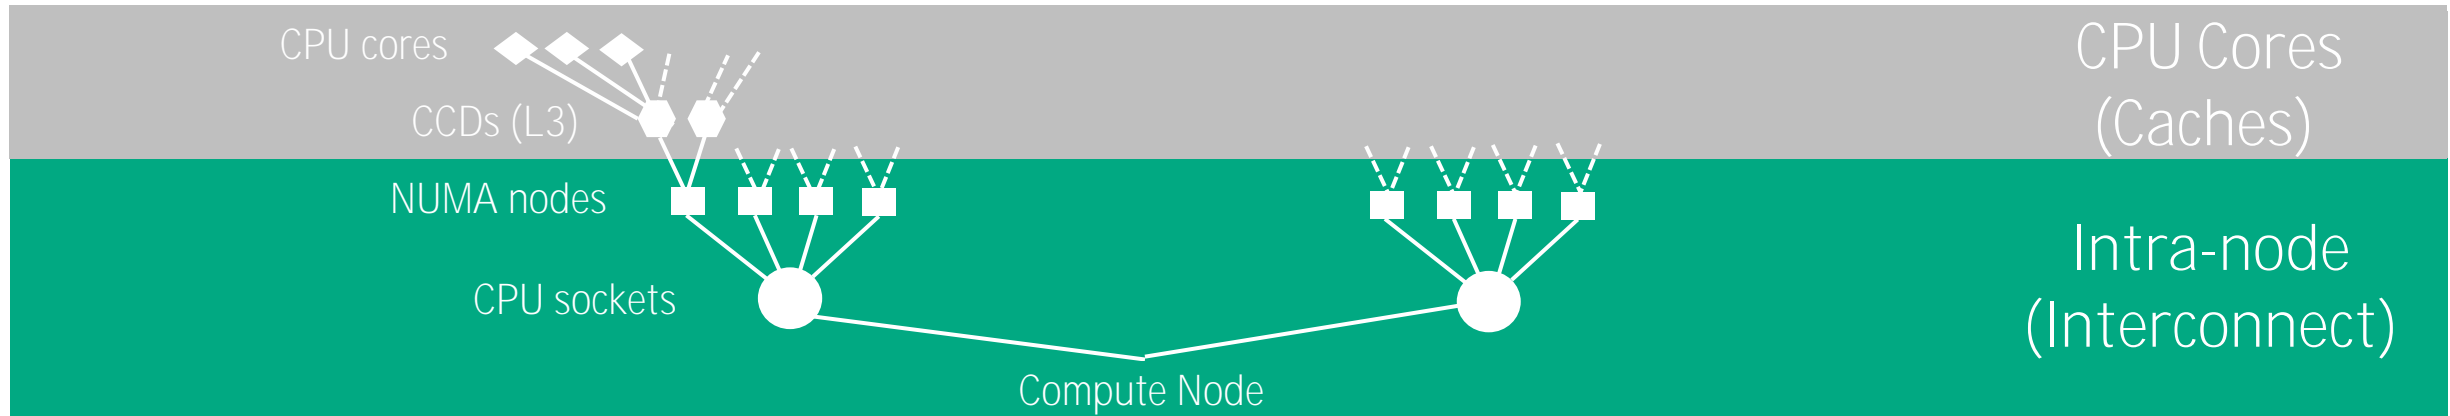
Compute Complex Die (CCD)



Compute Complex Dies (CCDs)

- host cores and L2/L3 cache
 - L1 cache 32kB / core
 - L2 cache 512kB / core
 - L3 cache 32MB / 8-cores

CORE AFFINITY – Binding And Potential Performance Issues



A task may have an affinity to multiple cores (`--cpu-bind=mask_cpu`). The same may apply to threads. We call this binding a process or a thread to CPUs.

- The concept binding is crucial for optimal performance :
 - memory/cache locality (minimize the data movement internally on the processor)
 - Make sure threads runs only on one core/hyperthread
 - To reduce OS costs when moving it (less locality: NUMA effects, different cache hierarchy)
 - To make sure no core/hyperthread is oversubscribed (resource sharing + context switches)
- It can be hard to spot performance problems relating to binding


BINDING only makes sense
if the node is reserved in
`--exclusive mode`

Binding Threads with OpenMP (in hybrid context MPI + OpenMP)



OPENMP – Set the Number of OpenMP Threads

export OMP_NUM_THREADS=2




Outer : create 1 new thread
to make 1 team of 2 threads

Inner : No new threads created
to make 2 teams of 1 threads

outer: thread_num = 0 ... level = 1
outer: thread_num = 1 ... level = 1
inner: thread_num = 0 ... level = 2
inner: thread_num = 0 ... level = 2

export OMP_NUM_THREADS=2,2



Outer : create 1 new thread
to make 1 team of 2 threads

Inner : 4 new threads to
make 2 teams of 3 threads

outer: thread_num = 0 ... level = 1
outer: thread_num = 1 ... level = 1
inner: thread_num = 0 ... level = 2
inner: thread_num = 1 ... level = 2
inner: thread_num = 2 ... level = 2
inner: thread_num = 0 ... level = 2
inner: thread_num = 1 ... level = 2
inner: thread_num = 2 ... level = 2

```
int main() {  
    #pragma omp parallel  
    printf("outer: ...");  
    #pragma omp parallel  
    printf("inner: ... ");  
}
```

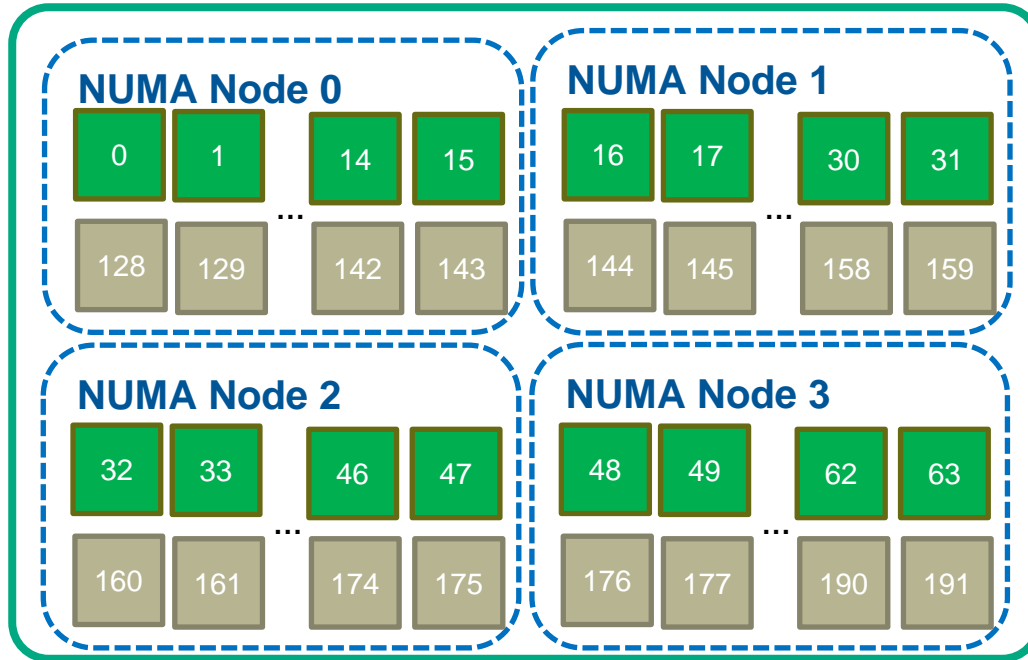
OPENMP BINDING – OMP_PLACES

- A list of places (affinity) that threads can be pinned on.
- Each “Place” defines a location where a **thread can “float”**
- Places do NOT depend on the number of threads (OMP_NUM_THREADS)
- Keep in mind that the MPI task already has affinity which may restrict an OMP places
- **OMP_PLACES=values** where possible values are:
 - **threads**: Each place corresponds to a single hardware thread (hyperthread) on the target machine.
 - **cores**: Each place corresponds to a single core (having one or more hardware threads)
 - **sockets**: Each place corresponds to a single socket (consisting of one or more cores)
 - Or a list with explicit values e.g., “{0:4}:4:4 = {0,1,2,3},{4,5,6,7},{8,9,10,11},{12,13,14,15}”



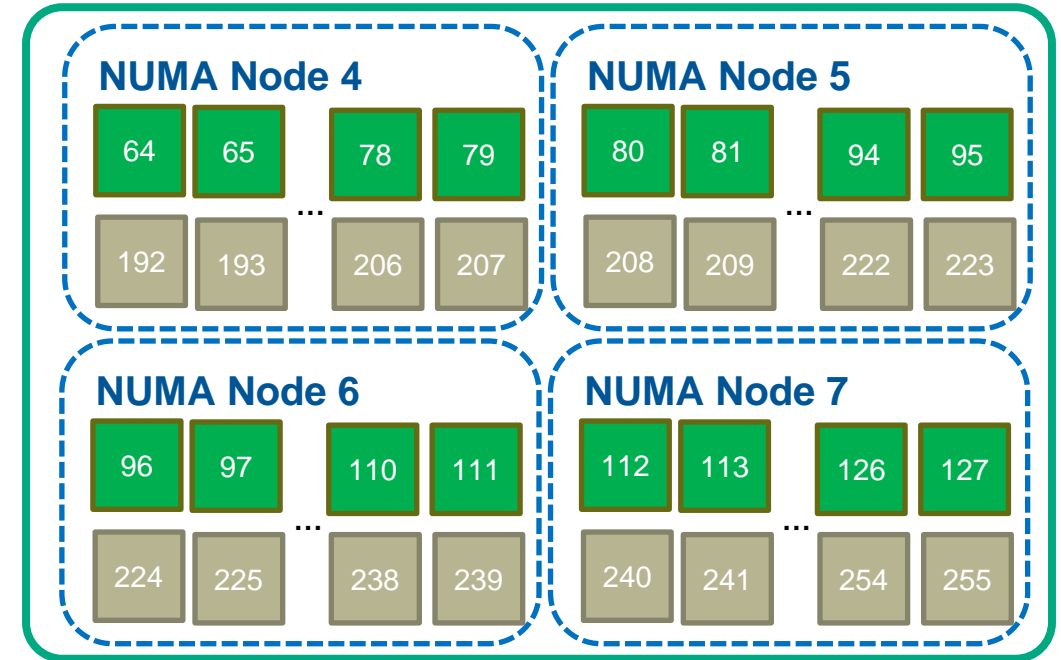
OPENMP BINDING – OMP_PLACES=threads

SOCKET 0



SOCKET 1

--hint=multithread



> **srun -N 1 --ntasks 8 -c 32**

OMP_DISPLAY_ENV=true

OMP_PLACES=threads

1 place = a single
hardware thread

```
OMP_PLACES = '{0},{128},{1},{129},{2},{130},{3} ...
OMP_PLACES = '{16},{144},{17},{145},{18},{146},{19} ...
OMP_PLACES = '{32},{160},{33},{161},{34},{162},{35} ...
OMP_PLACES = '{48},{176},{49},{177},{50},{178},{51} ...
OMP_PLACES = '{64},{192},{65},{193},{66},{194},{67} ...
OMP_PLACES = '{80},{208},{81},{209},{82},{210},{83} ...
OMP_PLACES = '{96},{224},{97},{225},{98},{226},{99} ...
OMP_PLACES = '{112},{240},{113},{241},{114},{242},{115} ...
```


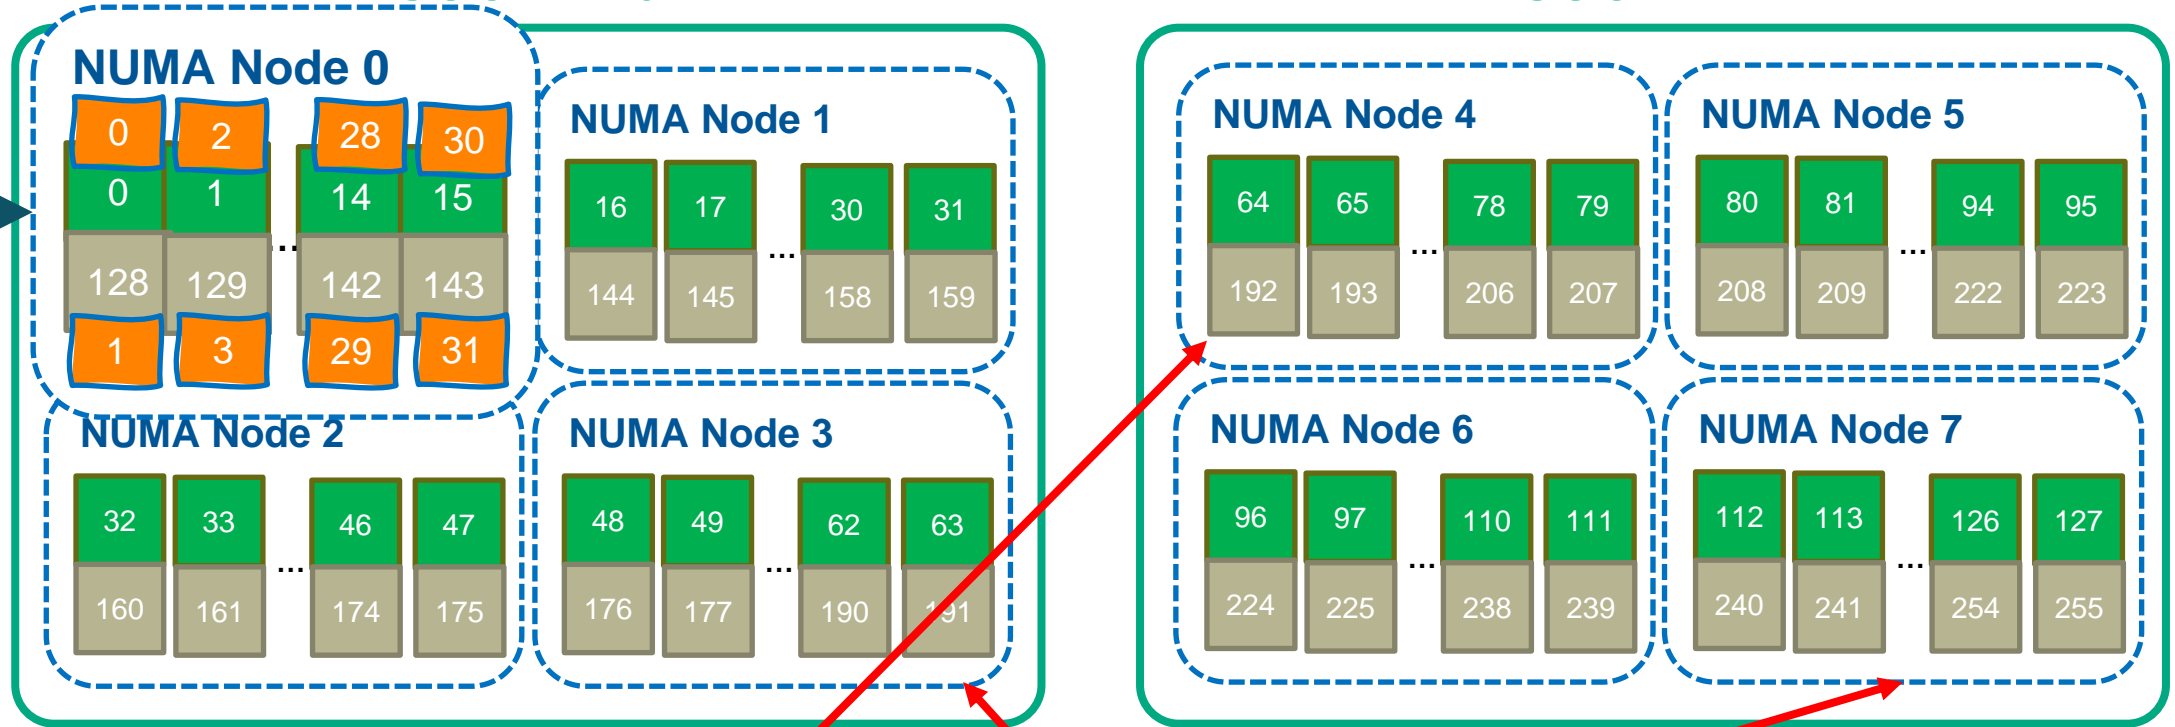
OPENMP BINDING – OMP_PLACES=threads

--hint=multithread

> **srun -N 1 --ntasks 8 -c 32 --distribution=block:cyclic ./\${EXE}**

SOCKET 0

SOCKET 1



32 threads per rank



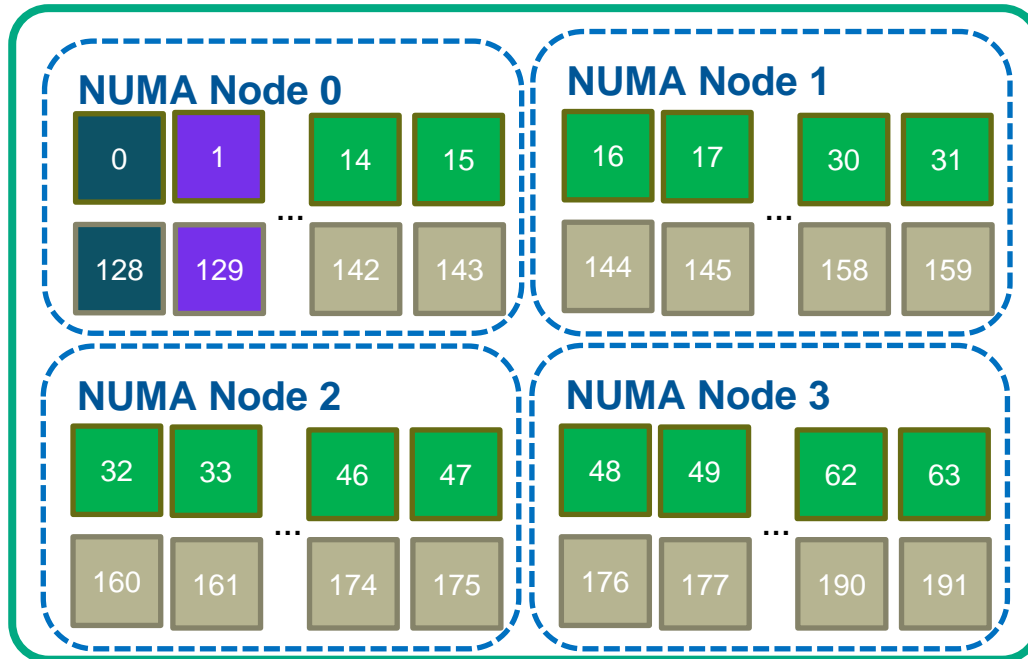
MPI TASKS

export OMP_PLACES=threads
export OMP_NUM_THREADS=32

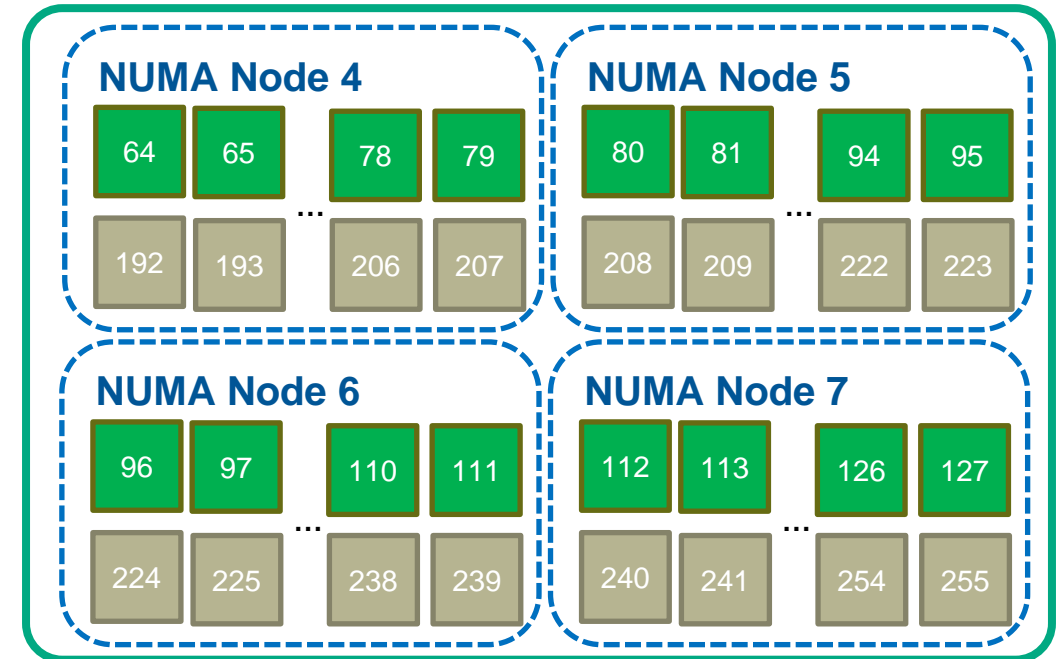
OPENMP BINDING – OMP_PLACES=cores

--hint=multithread

SOCKET 0



SOCKET 1



> srun -N 1 --ntasks 8 -c 32

OMP_DISPLAY_ENV=true

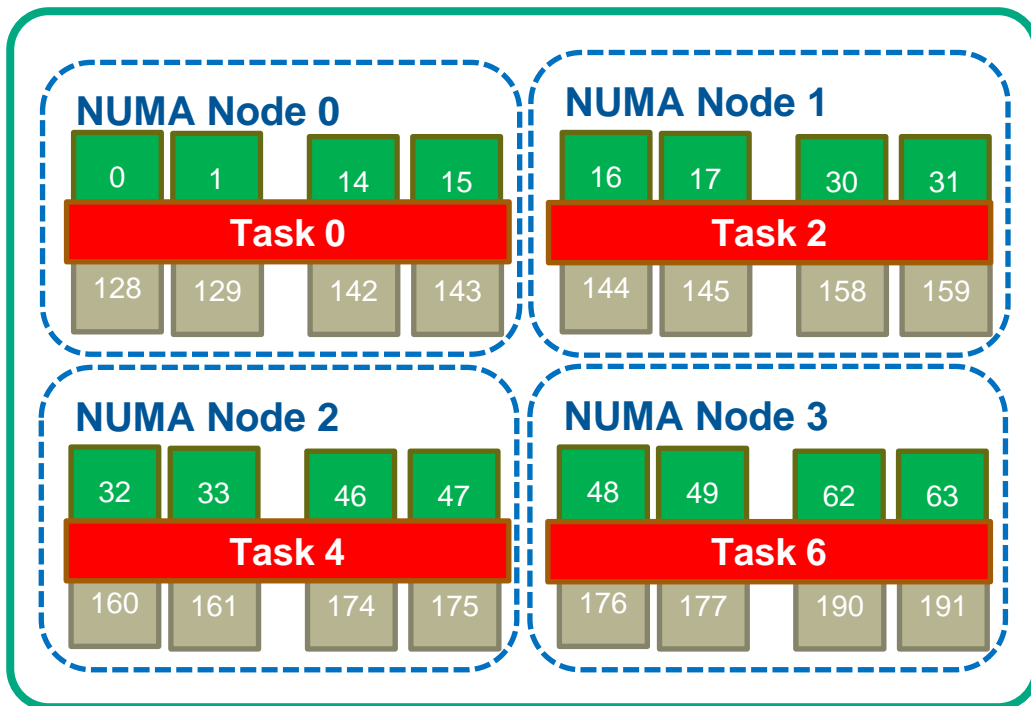
OMP_PLACES=cores
1 place = a single core

```
OMP_PLACES = '{0,128},{1,129},{2,130},{3,131} ...
OMP_PLACES = '{16,144},{17,145},{18,146},{19,147} ...
OMP_PLACES = '{32,160},{33,161},{34,162},{35,163} ...
OMP_PLACES = '{48,176},{49,177},{50,178},{51,179} ...
OMP_PLACES = '{64,192},{65,193},{66,194},{67,195} ...
OMP_PLACES = '{80,208},{81,209},{82,210},{83,211} ...
OMP_PLACES = '{96,224},{97,225},{98,226},{99,227} ...
OMP_PLACES = '{112,240},{113,241},{114,242},{115,243} ...
```

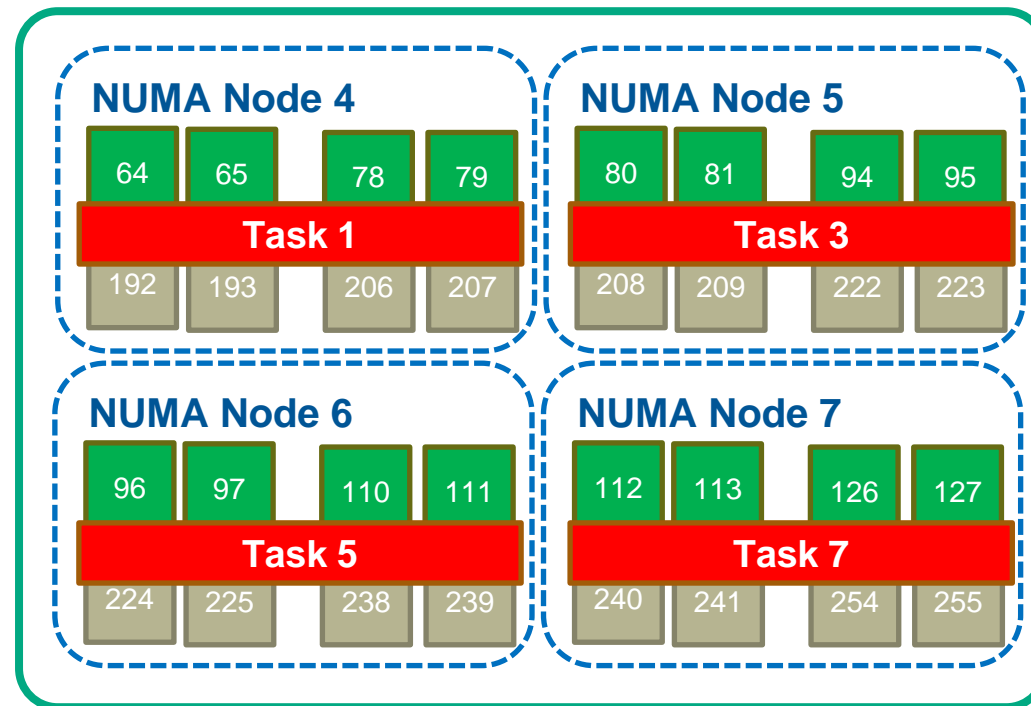

OPENMP BINDING - OMP_PLACES=sockets

```
> srun --nodes 1 -n 8 -c 32 --distribution=block:cyclic --hint=multithread
```

SOCKET 0



SOCKET 1



```
OMP_PLACES = '{0:16,128:16}'
OMP_PLACES = '{16:16,144:16}'
OMP_PLACES = '{32:16,160:16}'
OMP_PLACES = '{48:16,176:16}'
OMP_PLACES = '{64:16,192:16}'
OMP_PLACES = '{80:16,208:16}'
OMP_PLACES = '{96:16,224:16}'
OMP_PLACES = '{112:16,240:16}'
```

| host | rank | thr | count | mask |
|-----------|------|-----|-------|----------------------|
| nid002241 | 0 | 0 | 32 | cpus 0-15 128-143 |
| | 1 | 0 | 32 | cpus 64-79 192-207 |
| | 2 | 0 | 32 | cpus 16-31 144-159 |
| | 3 | 0 | 32 | cpus 80-95 208-223 |
| | 4 | 0 | 32 | cpus 32-47 160-175 |
| | 5 | 0 | 32 | cpus 96-111 224-239 |
| | 6 | 0 | 32 | cpus 48-63 176-191 |
| | 7 | 0 | 32 | cpus 112-127 240-255 |

OMP_PLACES=sockets
OMP_NUM_THREADS=1

threads can "float"

OPENMP BINDING – OMP_PROC_BIND

- Sets the binding of threads to processors.
 - **close**: Bind threads close to the master thread while still distributing for load balancing.
 - **spread**: Bind threads as evenly distributed (spread) as possible.
 - **master**: Bind threads to the same place as the master thread.
 - **false**: turns off OMP binding





OPENMP BINDING – OMP_PROC_BIND=close

--hint=multithread

- Here we specify 8 MPI tasks with 4 OpenMP threads with places = threads or cores


SOCKET 0



```
export OMP_PROC_BIND=close
export OMP_PLACES=threads
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|-------|------|
| nid002241 | 0 | 0 | 1 cpu | 0 |
| | | 1 | 1 cpu | 128 |
| | | 2 | 1 cpu | 1 |
| | | 3 | 1 cpu | 129 |
| | 1 | 0 | 1 cpu | 64 |

...



```
export OMP_PROC_BIND=close
export OMP_PLACES=cores
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|--------|--------|
| nid002241 | 0 | 0 | 2 cpus | 0 128 |
| | | 1 | 2 cpus | 1 129 |
| | | 2 | 2 cpus | 2 130 |
| | | 3 | 2 cpus | 3 131 |
| | 1 | 0 | 2 cpus | 64 192 |

...


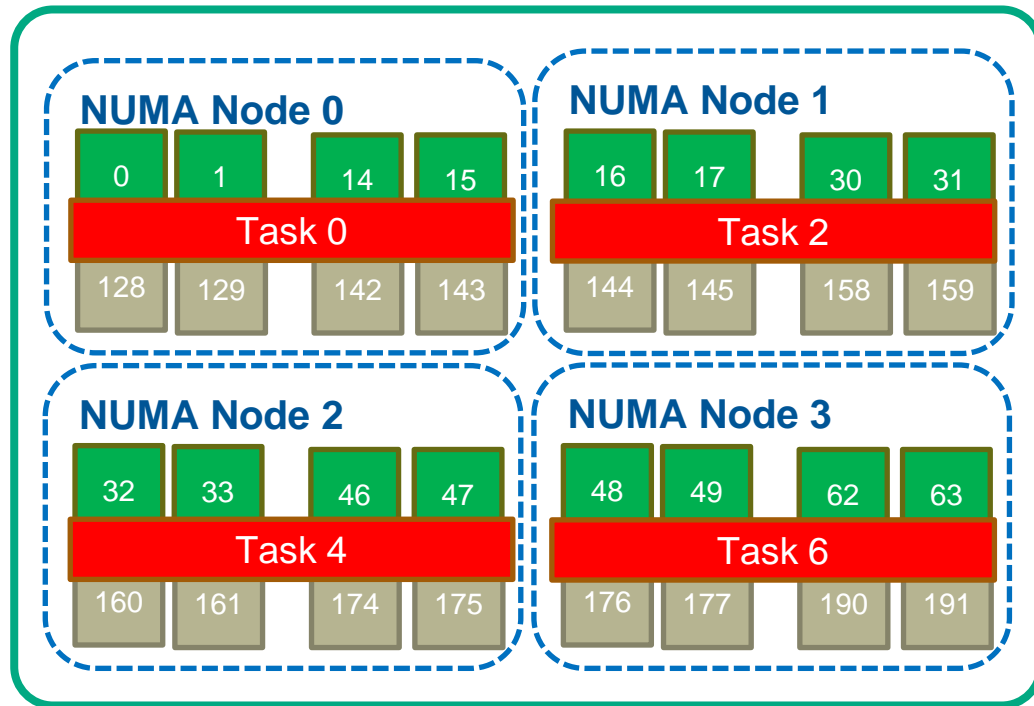
```
> srun -N 1 --ntasks 8 -c 32 --distribution=block:cyclic --hint=multithread ./${EXE}
```

OPENMP BINDING – OMP_PROC_BIND=spread

--hint=multithread

- Here we specify 8 MPI tasks with 4 OpenMP threads with places = threads or cores


SOCKET 0



```
export OMP_PROC_BIND=spread
export OMP_PLACES=threads
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|-------|------|
| nid002241 | 0 | 0 | 1 cpu | 0 |
| | | 1 | 1 cpu | 4 |
| | | 2 | 1 cpu | 8 |
| | | 3 | 1 cpu | 12 |
| | 1 | 0 | 1 cpu | 64 |

...



```
export OMP_PROC_BIND=spread
export OMP_PLACES=cores
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|--------|--------|
| nid002241 | 0 | 0 | 2 cpus | 0 128 |
| | | 1 | 2 cpus | 4 132 |
| | | 2 | 2 cpus | 8 136 |
| | | 3 | 2 cpus | 12 140 |
| | 1 | 0 | 2 cpus | 64 192 |

...

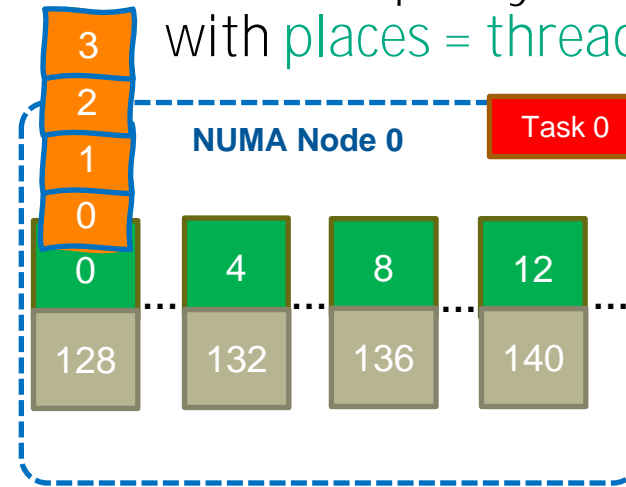
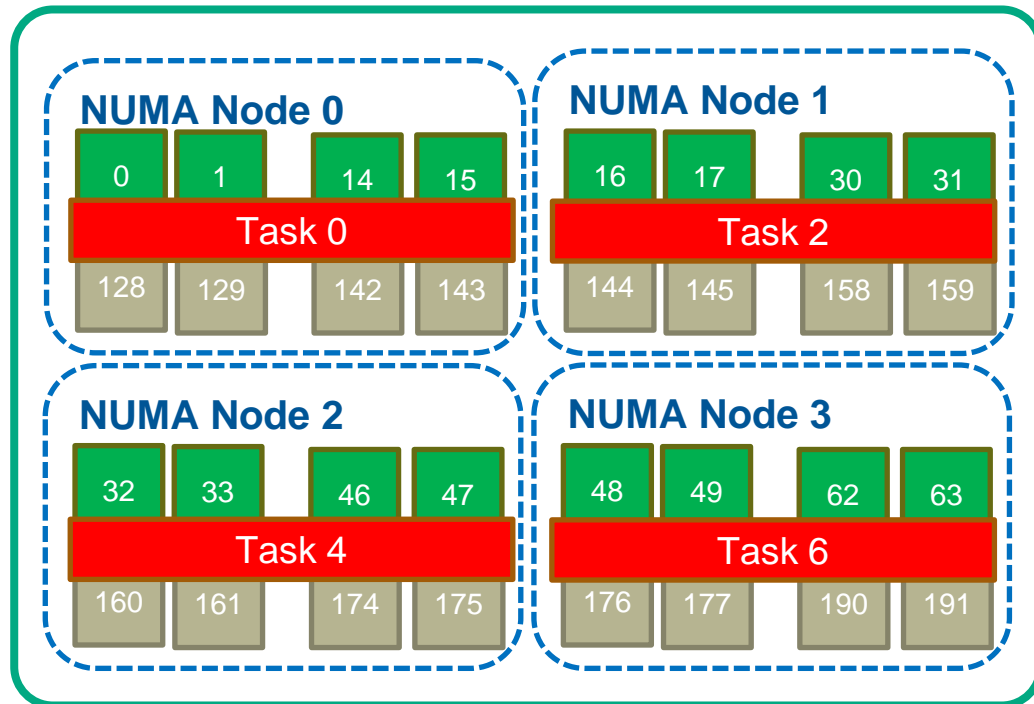
> srun -N 1 --ntasks 8 -c 32 --distribution=block:cyclic --hint=multithread ./{EXE}

OPENMP BINDING – OMP_PROC_BIND=master

--hint=multithread

- Here we specify 8 MPI tasks with 4 OpenMP threads with places = threads or cores

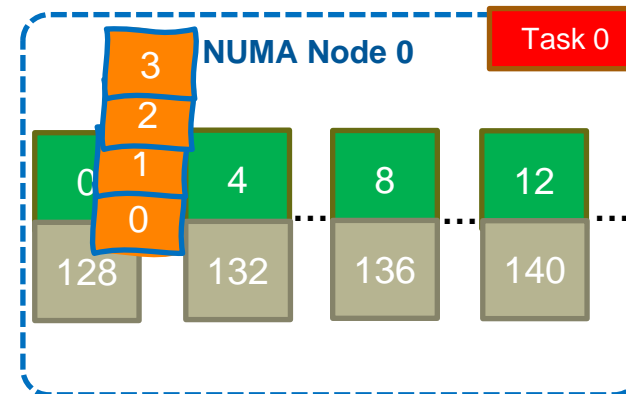
SOCKET 0



```
export OMP_PROC_BIND=master
export OMP_PLACES=threads
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|-------|------|
| nid002241 | 0 | 0 | 1 cpu | 0 |
| | | 1 | 1 cpu | 0 |
| | | 2 | 1 cpu | 0 |
| | | 3 | 1 cpu | 0 |
| | 1 | 0 | 1 cpu | 64 |

...



```
export OMP_PROC_BIND=master
export OMP_PLACES=cores
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|--------|--------|
| nid002241 | 0 | 0 | 2 cpus | 0 128 |
| | | 1 | 2 cpus | 0 128 |
| | | 2 | 2 cpus | 0 128 |
| | | 3 | 2 cpus | 0 128 |
| | 1 | 0 | 2 cpus | 64 192 |

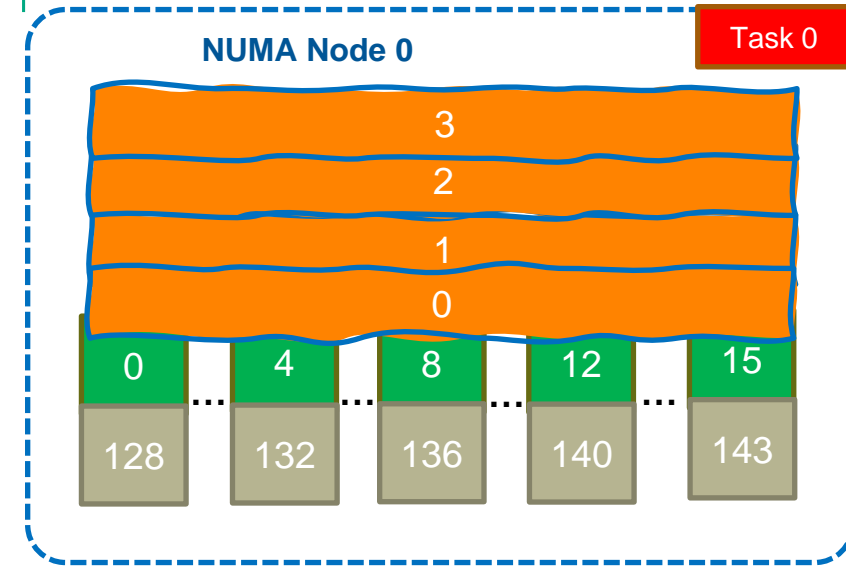
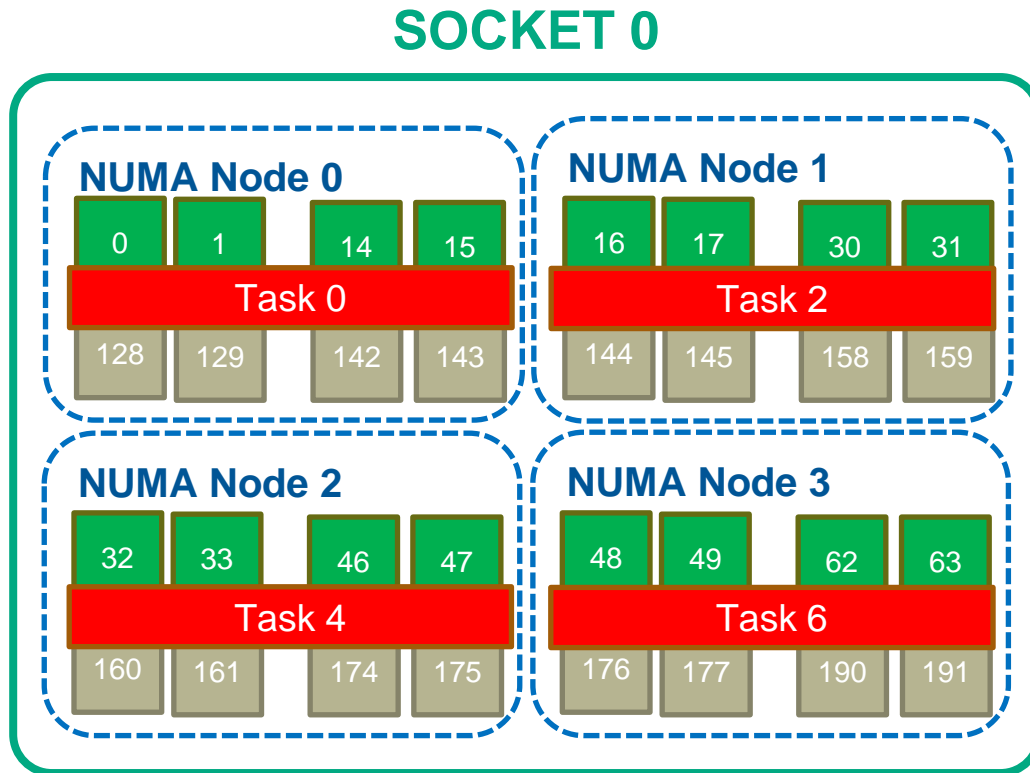
...

> srun -N 1 --ntasks 8 -c 32 --distribution=block:cyclic --hint=multithread ./{EXE}

OPENMP BINDING – OMP_PROC_BIND=false

--hint=multithread

- Here we specify 8 MPI tasks with 4 OpenMP threads with places = threads or cores



```
export OMP_PROC_BIND=false
export OMP_PLACES=threads | cores | sockets #same results
export OMP_NUM_THREADS=4
```

| host | rank | thr | count | mask |
|-----------|------|-----|-------|--------------------|
| nid002241 | 0 | 0 | 32 | cpus 0-15 128-143 |
| | | 1 | 32 | cpus 0-15 128-143 |
| | | 2 | 32 | cpus 0-15 128-143 |
| | | 3 | 32 | cpus 0-15 128-143 |
| | 1 | 0 | 32 | cpus 64-79 192-207 |

> **srun -N 1 --ntasks 8 -c 32--distribution=block:cyclic --hint=multithread ./{EXE}**

CORE AFFINITY – Takeaways

- Multiple levels of setting affinities (ex. MPI Ranks, OpenMP): OpenMP threads affinities are restricted by the parent process (MPI task), finer control can be set using OMP environment variables.
- Ensure cores or hyperhtreads are not oversubscribed (shared compute resource, OS context switches)
- Prefer binding a thread to a single core: Allowing threads to float between cores may impact cache locality and induce NUMA effects

Cheat sheet

| | |
|---|---|
| <code>export OMP_NUM_THREADS=X</code> | Set number of OpenMP threads per task |
| <code>export OMP_PLACES=threads cores sockets <list></code> | The way the threads are distributed |
| <code>export OMP_PROC_BIND=close spread master false</code> | Assinging a core/hyperthread to each thread |





INTRA-NODE PLACEMENT FOR GPUS

Distributing tasks to better harness GPUs

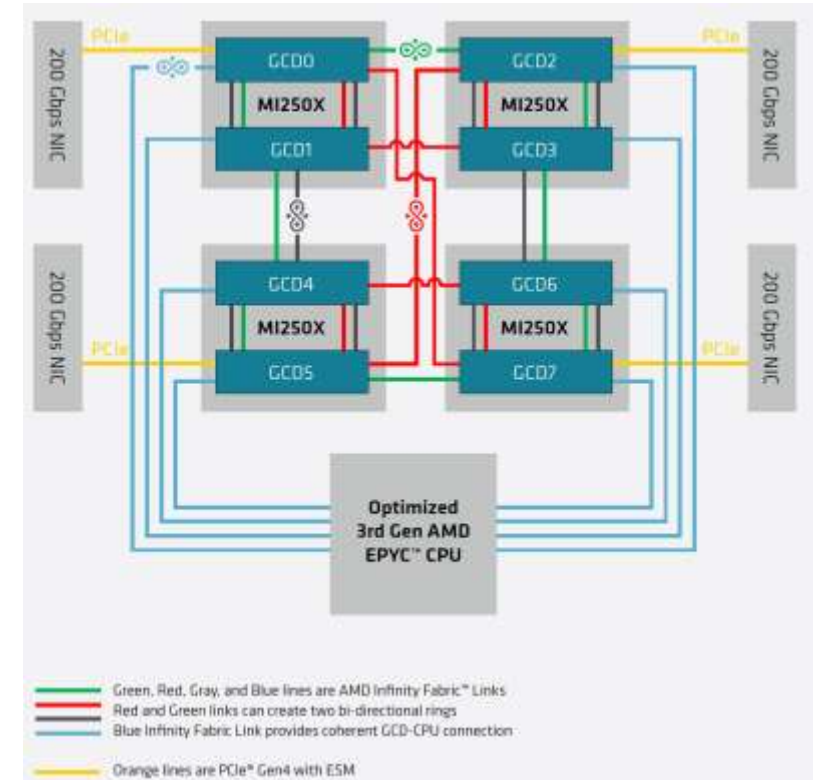
PLACEMENT FOR GPUS – Resources on Lumi-G

On **Lumi-G**:

- 4 MI250X per node = 8 GCDs (Graphics Complex Die) in total
- Will show as 8 separate GPUs according to Slurm terminology

Note :

- Just Keep in mind that 2 GDCs can be on the same physical package = better connectivity

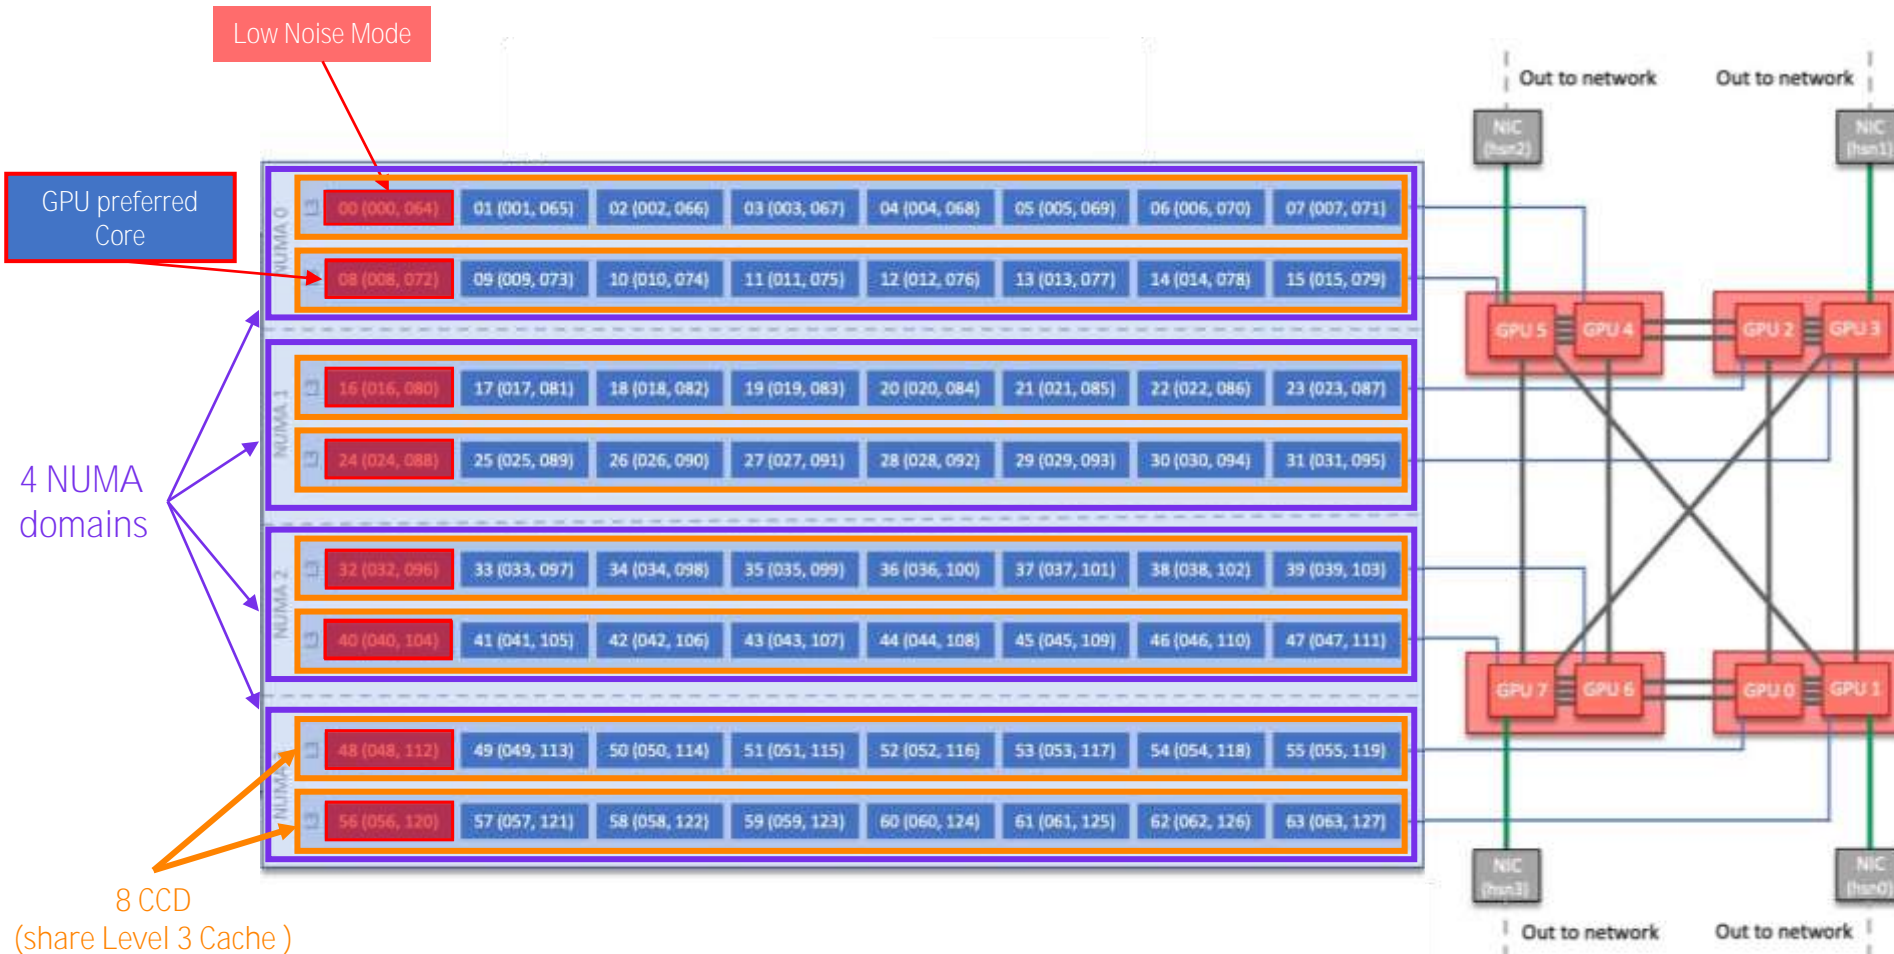


One MI250X GPU is splitted in Two GCDs

```
uan02:~> sinfo -o "%10R  %10c  %10m  %20G  %20F" -p standard,standard-g
```

| PARTITION | CPUS | MEMORY | GRES | NODES(A/I/O/T) |
|------------|------|--------|------------------|-------------------------|
| standard | 256 | 229376 | (null) | 1151 / 317 / 260 / 1728 |
| standard-g | 128 | 491520 | gpu:mi250:8(S:0) | 2346 / 125 / 92 / 2563 |

PLACEMENT FOR GPUS – CPU / GPU Affinity



MPI Ranks ≥ 8

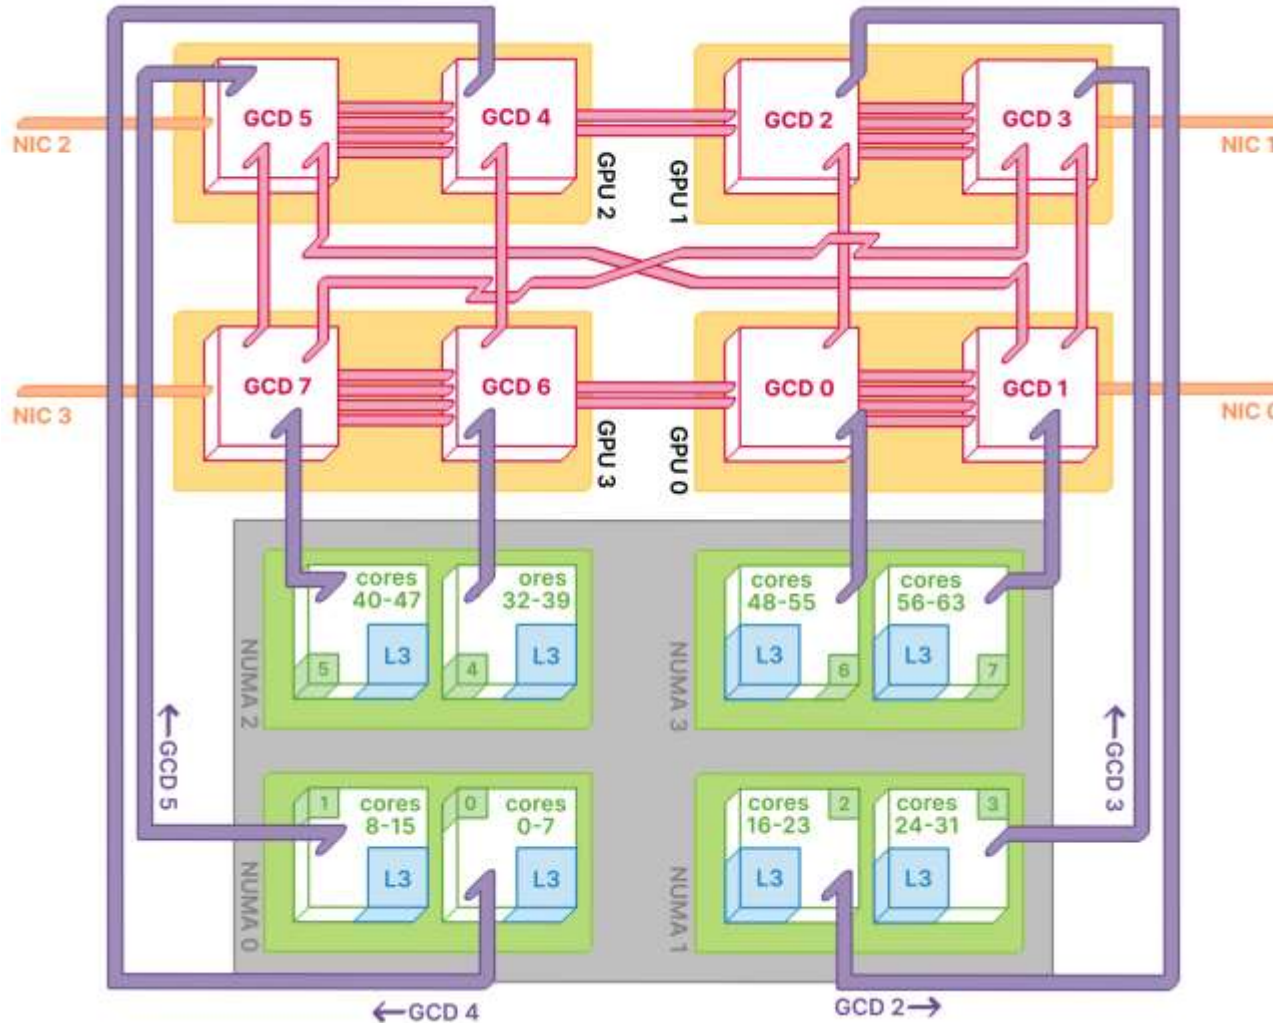
- Pinned to the 8 CCD

OpenMP Threads ≤ 7




- Cores
0,8,16,24,32,40,48,56
not available

The first cores in each CDD are reserved for
- servicing the OS + GPU management

PLACEMENT FOR GPUS –Potential Performance Issues



- NUMA affinity for each GPU: Improper task placement may impact transfers between the GCDs and main memory (aka Host to Device and Device to Host transfers).
- Asymmetric interconnect between GPUs (GCDs): Task ordering matters to ensure faster inter-GCD communications (aka IPC).
- Network interfaces directly attached to the GPUs: Transfers could be slightly slower when initiated by the CPU.

 **Infinity fabric GPU-GPU**
 50+50 GB/s
 **Infinity fabric CPU-GPU**
 36+36 GB/s * (28 + 28 GB/s limited by DMA engines)
 **Cray Slingshot-11 interconnect**
 25+25 GB/s

PLACEMENT FOR GPUS – GPUs to NUMA Domains Mapping

- GPUs are associated to NUMA nodes
- Can use rocm-smi to see the topology

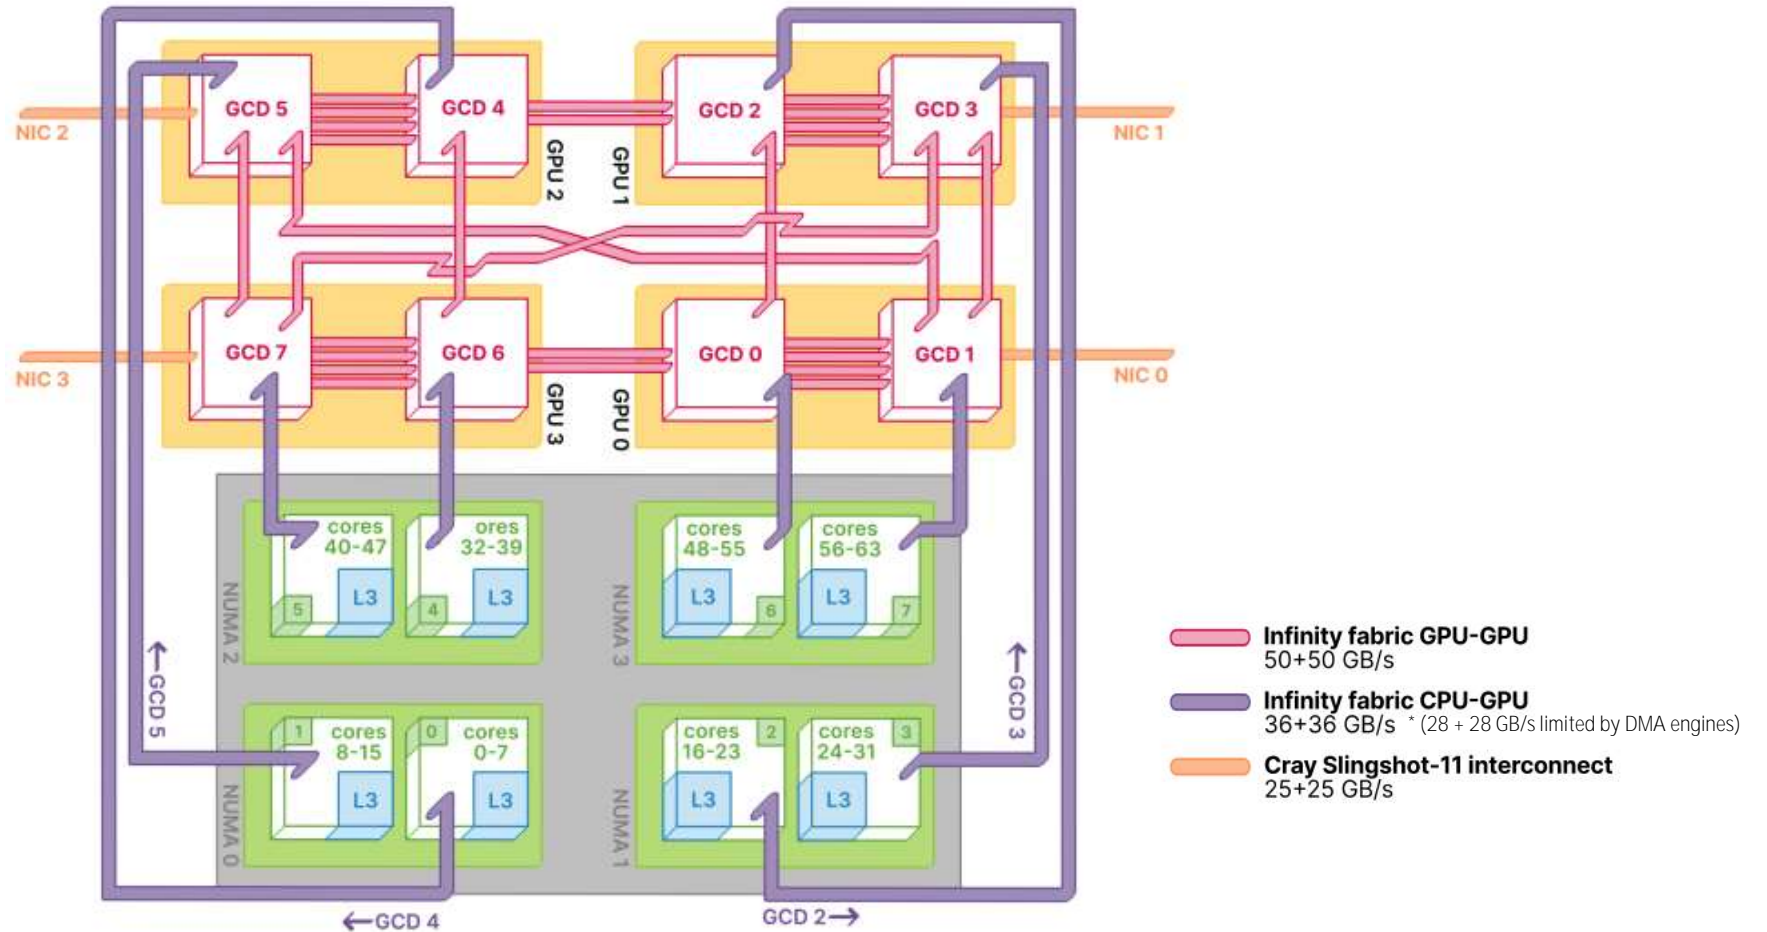
```
> srun --nodes=1 -t "00:10:00" --ntasks=1 --gres=gpu:8 rocm-smi -showtoponuma | grep Affinity
```

```
...
```

```
GPU[0]      : (Topology) Numa Affinity: 3
GPU[1]      : (Topology) Numa Affinity: 3
GPU[2]      : (Topology) Numa Affinity: 1
GPU[3]      : (Topology) Numa Affinity: 1
GPU[4]      : (Topology) Numa Affinity: 0
GPU[5]      : (Topology) Numa Affinity: 0
GPU[6]      : (Topology) Numa Affinity: 2
GPU[7]      : (Topology) Numa Affinity: 2
```



PLACEMENT FOR GPUS – GPUs to NUMA Domains Mapping



| NUMA ID | 0 | 0 | 1 | 1 | 2 | 2 | 3 | 3 |
|----------------|---|---|---|---|---|---|---|---|
| Optimal GPU ID | 4 | 5 | 2 | 3 | 6 | 7 | 0 | 1 |

PLACEMENT FOR GPUS – Assuming one GPU per MPI Rank

- Don't use `--gpus-per-task` or `--gpu-bind` (Slurm issue with cgroups preventing to use intra-node GPU to GPU communications).
- Associate each MPI rank to a given GPU with two environment variables:
 - **ROCR_VISIBLE_DEVICES**
 - Limit the number of GPU devices that are available for a given process
 - **SLURM_LOCALID**
 - Node local task ID for the process within a job
- Remember, you can check GPU bindings with the `gpu_check` tool

```
MPI 000 - OMP 000 - HWT 001 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
MPI 000 - OMP 001 - HWT 002 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
MPI 000 - OMP 002 - HWT 003 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
MPI 000 - OMP 003 - HWT 004 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
...
MPI 007 - OMP 004 - HWT 060 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
MPI 007 - OMP 005 - HWT 061 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
MPI 007 - OMP 006 - HWT 062 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
MPI 007 - OMP 007 - HWT 063 - Node nid005015 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,de
```

PLACEMENT FOR GPUS – Naive Approach (1/2)

- Let's consider the example of the 8 MPI tasks/node shown before
- Use a script to properly set **ROCR_VISIBLE_DEVICES** during the job submission

job.slurm

```
#!/bin/bash
#SBATCH -p <partition>
#SBATCH -A <your_project>
#SBATCH --time=00:02:00
#SBATCH --nodes=2
#SBATCH --gres=gpu:8
#SBATCH --exclusive
#SBATCH --ntasks-per-node=8
#SBATCH --cpus-per-task=7
#SBATCH --hint=nomultithread
...
${ASRUN} ./select_gpu.sh <my_app>
```

select_gpu.sh

```
#!/bin/bash

export ROCR_VISIBLE_DEVICES=$SLURM_LOCALID

exec $*
```

- Naïve approach example, Each MPI rank :
 - Detects the availability of a single GPU device
 - Is essentially associated with a different GPU device to ensure that two MPI ranks do not share the same GPU device



PLACEMENT FOR GPUS – Naive Approach (2/2)

- What we want

| | | | | | | | | |
|----------------|---|---|---|---|---|---|---|---|
| Local Task ID | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| NUMA ID | 0 | 0 | 1 | 1 | 2 | 2 | 3 | 3 |
| Optimal GPU ID | 4 | 5 | 2 | 3 | 6 | 7 | 0 | 1 |

- Naïve approach binding map

| | | | | | | | | |
|---------------|-------|--------|---------|---------|---------|---------|---------|---------|
| Local Task ID | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| Cores ID | 1 - 7 | 9 - 15 | 17 - 23 | 25 - 31 | 33 - 39 | 41 - 47 | 49 - 55 | 57 - 63 |
| NUMA ID | 0 | 0 | 1 | 1 | 2 | 2 | 3 | 3 |
| GPU ID | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |

- This mapping is not optimal



PLACEMENT FOR GPUS – Affinity Script

- Adapt `select_gpu.sh` script for the optimal GPU ID

```
#!/bin/bash
GPUSID="4 5 2 3 6 7 0 1"
GPUSID=(${GPUSID})
if [ ${#GPUSID[@]} -gt 0 -a -n "${SLURM_NTASKS_PER_NODE}" ]; then
    if [ ${#GPUSID[@]} -gt $SLURM_NTASKS_PER_NODE ]; then
        export ROCR_VISIBLE_DEVICES=${GPUSID[$(($SLURM_LOCALID))]}
    else
        export ROCR_VISIBLE_DEVICES=${GPUSID[$(($SLURM_LOCALID / ($SLURM_NTASKS_PER_NODE / ${#GPUSID[@]})))]}
    fi
fi
exec $*
```

MPI 000 - OMP 000 - GPU_ID 4 - Bus_ID d1
 MPI 001 - OMP 000 - GPU_ID 5 - Bus_ID d6
 MPI 002 - OMP 000 - GPU_ID 2 - Bus_ID c9
 MPI 003 - OMP 000 - GPU_ID 3 - Bus_ID ce
 MPI 004 - OMP 000 - GPU_ID 6 - Bus_ID d9
 MPI 005 - OMP 000 - GPU_ID 7 - Bus_ID de
 MPI 006 - OMP 000 - GPU_ID 0 - Bus_ID c1
 MPI 007 - OMP 000 - GPU_ID 1 - Bus_ID c6

| Local Task ID | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 |
|----------------|---|---|---|---|---|---|---|---|
| NUMA ID | 0 | 0 | 1 | 1 | 2 | 2 | 3 | 3 |
| Optimal GPU ID | 4 | 5 | 2 | 3 | 6 | 7 | 0 | 1 |

PLACEMENT FOR GPUS – Combining MPI tasks, OpenMP and GPUs (1/2)

- 8 tasks/node
- 7 threads/task
- 8 GPU/node
- 1 GPU/task

```
#!/bin/bash
#SBATCH --nodes=2
#SBATCH --gres=gpu:8
#SBATCH --ntasks-per-node=8
#SBATCH --cpus-per-task=7
#SBATCH --hint=nomultithread
#SBATCH --exclusive

export OMP_PLACES=cores #replace this by sockets
export OMP_PROC_BIND=close
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}

srun ${cpu_bind} ./select_gpu.sh gpu_check
```

command : `gpu_check | awk '{ gsub("- Node nid[0-9]+ - RT_GPU_ID [0-9]+ - ", ""); print }'`




PLACEMENT FOR GPUS – Combining MPI tasks, OpenMP and GPUs (2/2)

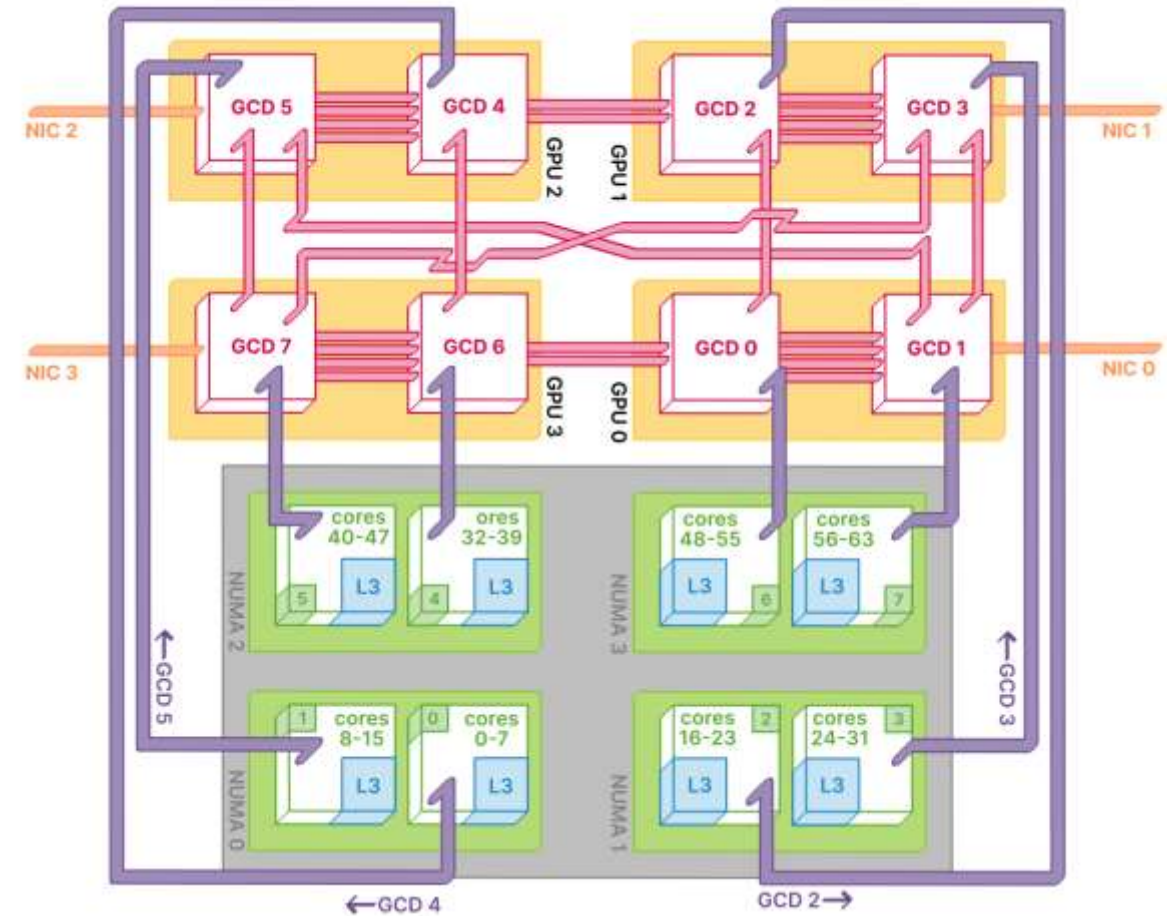
```

MPI 000 - OMP 000 - HWT 001 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 000 - OMP 001 - HWT 002 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 000 - OMP 002 - HWT 003 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 000 - OMP 003 - HWT 004 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 000 - OMP 004 - HWT 005 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 000 - OMP 005 - HWT 006 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 000 - OMP 006 - HWT 007 (CCD0) GPU_ID 4 - Bus_ID d1(GCD4/CCD0)
MPI 001 - OMP 000 - HWT 009 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 001 - OMP 001 - HWT 010 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 001 - OMP 002 - HWT 011 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 001 - OMP 003 - HWT 012 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 001 - OMP 004 - HWT 013 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 001 - OMP 005 - HWT 014 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 001 - OMP 006 - HWT 015 (CCD1) GPU_ID 5 - Bus_ID d6(GCD5/CCD1)
MPI 002 - OMP 000 - HWT 017 (CCD2) GPU_ID 2 - Bus_ID c9(GCD2/CCD2)

```

...

 **Infinity fabric GPU-GPU**
 50+50 GB/s
 **Infinity fabric CPU-GPU**
 36+36 GB/s
 **Cray Slingshot-11 interconnect**
 25+25 GB/s



PLACEMENT FOR GPUS – Mapping Processes to Network Interfaces

- On compute nodes that offer multiple GPU devices and multiple Network Interface Controllers (NIC), Cray MPI offers a flexible way to offer the ideal mapping between a process and the default NIC
 - We have a NIC per each Mi250X, i.e. 4 NICs per node
- For GPU-enabled parallel applications that involve MPI operations that access application arrays resident are on GPU-attached memory regions, users can set **MPICH_OFI_NIC_POLICY** to **“GPU”**
 - In this case, for each MPI process, Cray MPI strives to select a NIC device that is closest to the GPU device being used
- To display information pertaining to NIC selection:
export MPICH_OFI_NIC_VERBOSE=2
- More info in **mpi** man page



PLACEMENT FOR GPUS – Takeaways

- Don't use `--gpus-per-task` or `--gpu-bind` (Slurm issue with cgroups preventing to use intra-node GPU to GPU communications)
- Use instead a binding script to ensure proper NUMA affinity (impact transfers between Host and Device)
- When using GPU to GPU communications (different compute nodes), set `MPICH_GPU_SUPPORT_ENABLED=1` and `MPICH_OFI_NIC_POLICY=GPU`
- Potentially reorder tasks to better harness the asymmetric interconnect between GPUs in the same node

Cheat sheet

| | |
|---|---|
| <code>rocm-smi -showtopo</code> | Display GPU topology |
| <code>export ROCR_VISIBLE_DEVICES=<list></code> | Set GPUs visible by the task (tools and applications) |
| <code>export HIP_VISIBLE_DEVICES=<list></code> | Set GPUs visible by the task (HIP only) |
| <code>export MPICH_GPU_SUPPORT_ENABLED=1</code> | Enable MPI communications between GPUs |
| <code>export MPICH_OFI_NIC_POLICY=GPU</code> | Pick network interconnect closer to GCD |
| <code>export MPICH_RANK_REORDER_METHOD=3</code> <code>export MPICH_RANK_REORDER_FILE=<file></code> | Manually defining rank reordering (more details in another presentation) |



CONCLUSION



LUMI-C : OPTIMAL INITIAL PLACEMENT

2 sockets

x

8 CCD

=

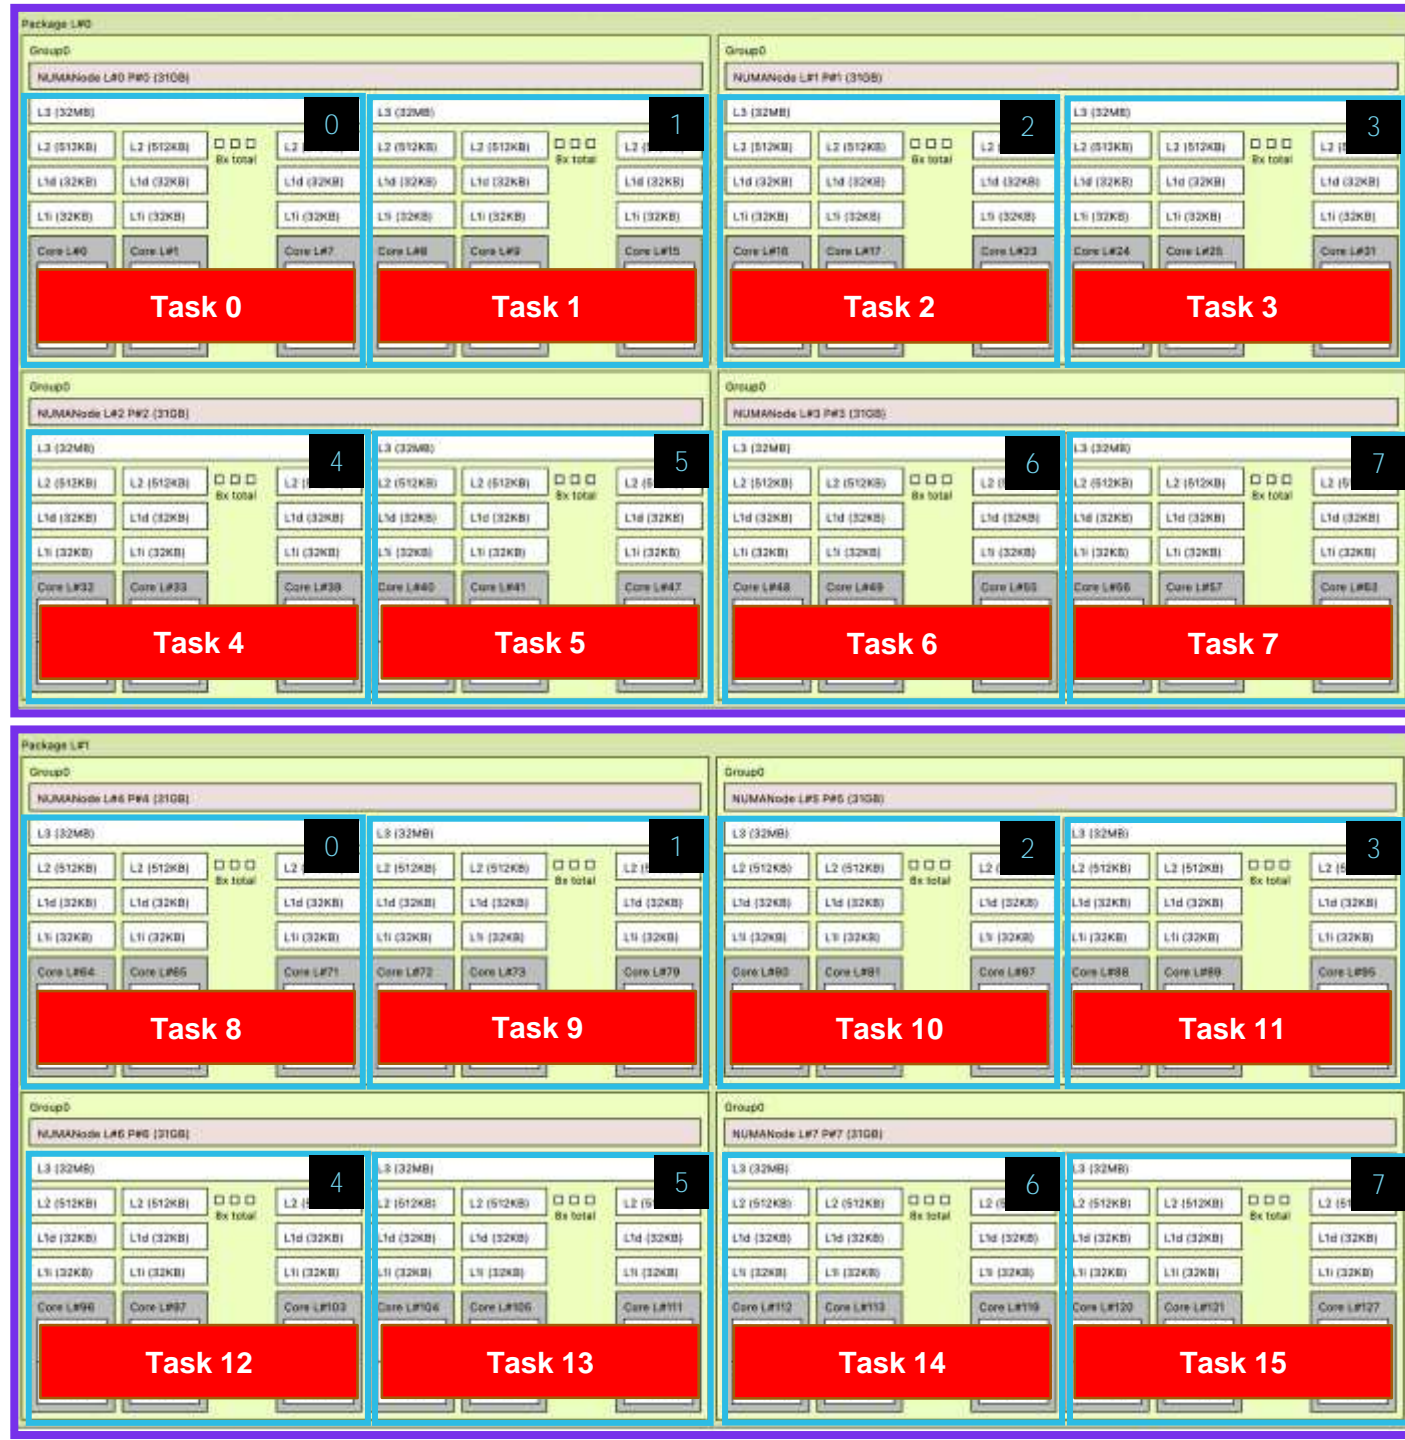
16 MPI TASKS

each with 8 cores

```
#SBATCH --exclusive
#SBATCH --hint=nomultithread
#SBATCH --ntasks-per-node=16
#SBATCH -N X
```

```
export OMP_NUM_THREADS=8
export OMP_PROC_BIND=close
export OMP_PLACES=cores
srun -c 8 <app>
```

(Using all cores in the node, with cores per task matching hardware)



LUMI-G : OPTIMAL INITIAL PLACEMENT

1 socket × 8 CCD = 8 MPI TASKS

each with 7 cores

```
#SBATCH --exclusive
#SBATCH --ntasks-per-node=8
#SBATCH --hint=nomultithread
#SBATCH --gres=gpu:8
#SBATCH -N X
```

```
export OMP_NUM_THREADS=7
export OMP_PROC_BIND=close
export OMP_PLACES=cores
export MPICH_GPU_SUPPORT_ENABLED=1
export MPICH_OFI_NIC_POLICY=GPU
```

```
srun -c 7 ./gpu_bind_script <app>
```



(Assuming one GPU per task. Using all available cores in the node, with cores per task matching hardware)

Documentation

- [Slurm documentation](#)
 - Slurm manpage
man srun
 - [LUMI Documentation](#)
 - [ARCHER2 Documentation](#)
 - [OpenMP documentation](#)
-
- Just be aware that some details are site specific





Questions?

/project/project_465001098/Exercises/HPE/day2/gpu_perf_binding/

```
✓ gpu_perf_binding
  ✓ check
    $ job.slurm
  ✓ himeno
    C himeno.c
    $ job.slurm
    M Makefile
    C param.h
    $ paramset.sh
    $ gpu_env.sh
    ⓘ README.md
    $ select_gpu_naive.sh
    $ select_gpu_opti.sh
```

1) Get the right environment

- (copy the exercises to your home directory)
- `cd ~/<your_path>/gpu_perf_binding`
- `source ../../lumi_g.sh`
- `source gpu_env.sh`

2) Start with example in *check* directory

- `cd check`
- `sbatch job.slurm`
 - test different bindings
 - `cpu_bind`, `gpu_bind` variables...

3) Launch *Himeno* from the root directory

- `cd himeno && make`
- `sbatch job.slurm`
 - test different bindings



check directory

Default run. Binding issues?

| | | | | | | | | | | | | |
|---------|---|---------|---|---------|---|----------------|---|-------------|---|----------|---|-----------|
| MPI 000 | - | OMP 000 | - | HWT 001 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 0 | - | Bus_ID c1 |
| MPI 000 | - | OMP 001 | - | HWT 001 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 0 | - | Bus_ID c1 |
| MPI 001 | - | OMP 000 | - | HWT 009 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 1 | - | Bus_ID c6 |
| MPI 001 | - | OMP 001 | - | HWT 009 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 1 | - | Bus_ID c6 |
| MPI 002 | - | OMP 000 | - | HWT 017 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 2 | - | Bus_ID c9 |
| MPI 002 | - | OMP 001 | - | HWT 017 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 2 | - | Bus_ID c9 |
| MPI 003 | - | OMP 000 | - | HWT 025 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 3 | - | Bus_ID ce |
| MPI 003 | - | OMP 001 | - | HWT 025 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 3 | - | Bus_ID ce |
| MPI 004 | - | OMP 000 | - | HWT 033 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 4 | - | Bus_ID d1 |
| MPI 004 | - | OMP 001 | - | HWT 033 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 4 | - | Bus_ID d1 |
| MPI 005 | - | OMP 000 | - | HWT 041 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 5 | - | Bus_ID d6 |
| MPI 005 | - | OMP 001 | - | HWT 041 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 5 | - | Bus_ID d6 |
| MPI 006 | - | OMP 000 | - | HWT 049 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 6 | - | Bus_ID d9 |
| MPI 006 | - | OMP 001 | - | HWT 049 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 6 | - | Bus_ID d9 |
| MPI 007 | - | OMP 000 | - | HWT 057 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 7 | - | Bus_ID dc |
| MPI 007 | - | OMP 001 | - | HWT 057 | - | Node nid005177 | - | RT_GPU_ID 0 | - | GPU_ID 7 | - | Bus_ID dc |



check directory

Enabling CPU masks for OpenMP threads and optimal GPU bindings:

```
MPI 000 - OMP 000 - HWT 001 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 4 - Bus_ID d1
MPI 000 - OMP 001 - HWT 002 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 4 - Bus_ID d1
MPI 001 - OMP 000 - HWT 009 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 5 - Bus_ID d6
MPI 001 - OMP 001 - HWT 010 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 5 - Bus_ID d6
MPI 002 - OMP 000 - HWT 017 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 2 - Bus_ID c9
MPI 002 - OMP 001 - HWT 018 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 2 - Bus_ID c9
MPI 003 - OMP 000 - HWT 025 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 3 - Bus_ID ce
MPI 003 - OMP 001 - HWT 026 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 3 - Bus_ID ce
MPI 004 - OMP 000 - HWT 033 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 6 - Bus_ID d9
MPI 004 - OMP 001 - HWT 034 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 6 - Bus_ID d9
MPI 005 - OMP 000 - HWT 041 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 7 - Bus_ID dc
MPI 005 - OMP 001 - HWT 042 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 7 - Bus_ID dc
MPI 006 - OMP 000 - HWT 049 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 0 - Bus_ID c1
MPI 006 - OMP 001 - HWT 050 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 0 - Bus_ID c1
MPI 007 - OMP 000 - HWT 057 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 1 - Bus_ID c6
MPI 007 - OMP 001 - HWT 058 - Node nid005174 - RT_GPU_ID 0 - GPU_ID 1 - Bus_ID c6
```



himeno directory

Default run:

```
MPI 000 - OMP 000 - HWT 001 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 001 - OMP 000 - HWT 002 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 002 - OMP 000 - HWT 003 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 003 - OMP 000 - HWT 004 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 004 - OMP 000 - HWT 005 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 005 - OMP 000 - HWT 006 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 006 - OMP 000 - HWT 007 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
MPI 007 - OMP 000 - HWT 009 - Node nid005175 - RT_GPU_ID 0,1,2,3,4,5,6,7 - GPU_ID 0,1,2,3,4,5,6,7 - Bus_ID c1,c6,c9,ce,d1,d6,d9,dc
```

cpu : 50.143480 sec.

Loop executed for 50 times

Gosa : 8.030980e-05

MFLOPS measured : 9012.093693

Score based on Pentium III 600MHz : 108.789156



himeno directory

Adding GPU binding:

```
MPI 000 - OMP 000 - HWT 001 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 4 - Bus_ID d1
MPI 001 - OMP 000 - HWT 002 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 5 - Bus_ID d6
MPI 002 - OMP 000 - HWT 003 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 2 - Bus_ID c9
MPI 003 - OMP 000 - HWT 004 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 3 - Bus_ID ce
MPI 004 - OMP 000 - HWT 005 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 6 - Bus_ID d9
MPI 005 - OMP 000 - HWT 006 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 7 - Bus_ID dc
MPI 006 - OMP 000 - HWT 007 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 0 - Bus_ID c1
MPI 007 - OMP 000 - HWT 009 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 1 - Bus_ID c6
```

cpu : 5.642638 sec.

Loop executed for 50 times

Gosa : 7.531330e-04

MFLOPS measured : 80086.260270

Score based on Pentium III 600MHz : 966.758333



himeno directory

Adding CPU binding and GPU binding:

```
MPI 000 - OMP 000 - HWT 001 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 4 - Bus_ID d1
MPI 001 - OMP 000 - HWT 009 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 5 - Bus_ID d6
MPI 002 - OMP 000 - HWT 017 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 2 - Bus_ID c9
MPI 003 - OMP 000 - HWT 025 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 3 - Bus_ID ce
MPI 004 - OMP 000 - HWT 033 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 6 - Bus_ID d9
MPI 005 - OMP 000 - HWT 041 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 7 - Bus_ID dc
MPI 006 - OMP 000 - HWT 049 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 0 - Bus_ID c1
MPI 007 - OMP 000 - HWT 057 - Node nid005173 - RT_GPU_ID 0 - GPU_ID 1 - Bus_ID c6
```

cpu : 5.302780 sec.

Loop executed for 50 times

Gosa : 9.317707e-04

MFLOPS measured : 85219.030186

Score based on Pentium III 600MHz : 1028.718375

