



**Hewlett Packard
Enterprise**

Advanced Performance Analysis

Comprehensive General LUMI Course

April 23–26, 2024

Agenda

- In-depth performance analysis with **perftools**
 - Sampling, tracing, and loop work estimates
 - Automatic Performance Analysis
- Reveal
 - Compiler Feedback and Variable Scoping
- OpenMP profiling
- Perftools API
 - Customized performance analysis
- Hardware performance counters
- Load imbalance analysis
- pat_run
 - Profile existing dynamically linked executables
- Concluding remarks

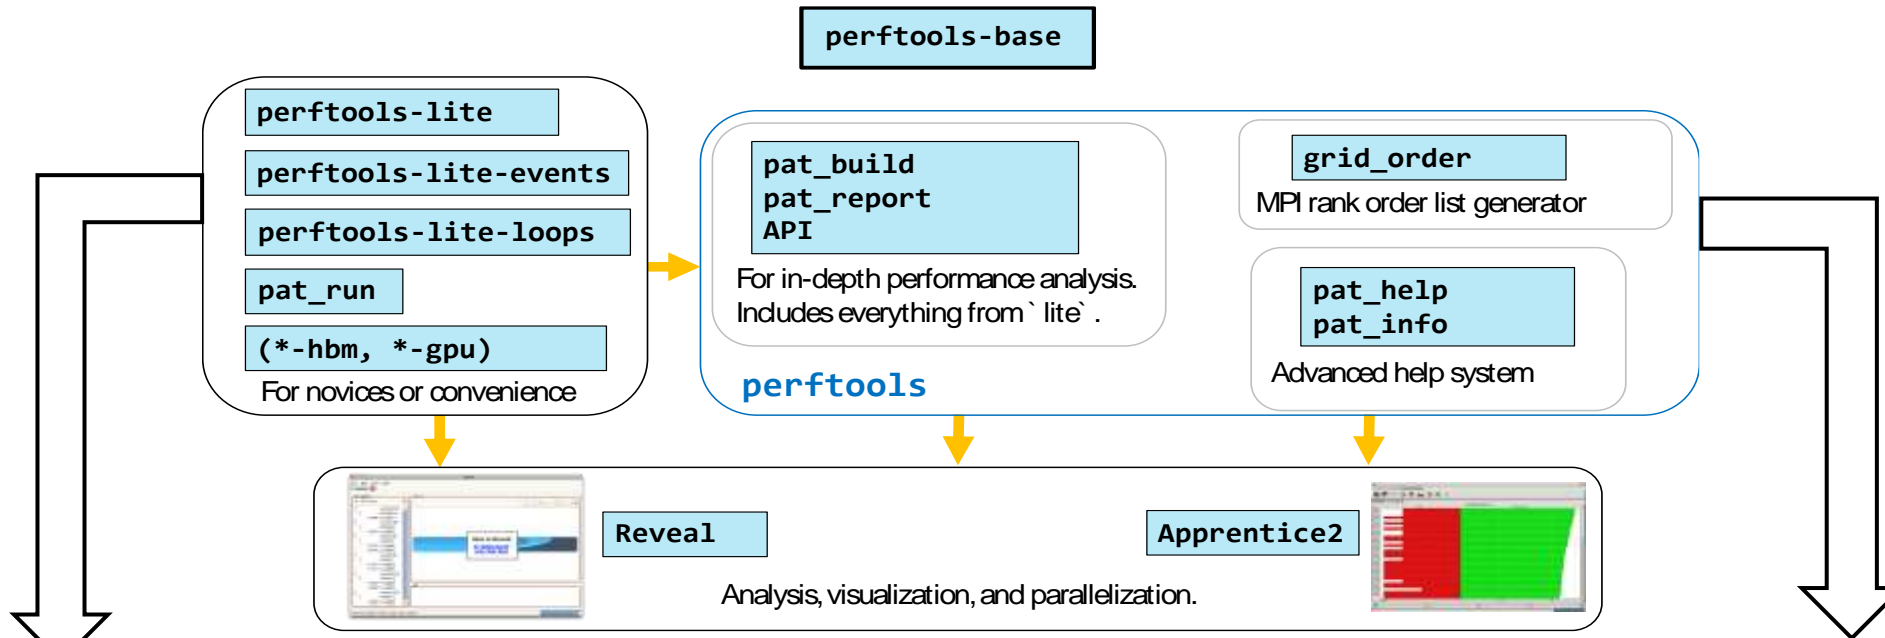


PerfTools

In-depth Performance Analysis



Overview



Previous talk:

1. Load the **perftools-lite*** module wrt desired profiling
2. Generate the binary
3. Run the binary
4. Look at the **automatically** generated report

This talk:

1. Load the **perftools** module
2. Generate the binary
3. **Manually** generate an **instrumented binary** wrt desired profiling
4. Run the **instrumented binary**
5. **Manually** generate the report
6. Look at the **manually** generated report

Sampling with Perftools

```
> module load perftools-base  
> module load perftools
```

```
> make clean; make  
> pat_build -S app.exe
```

- This generates a new executable **app.exe+pat** and preserves **app.exe**.
- Object files must be present during this stage.

```
> srun -n 8 ./app.exe+pat  
> pat_report -o myrep.trace.rpt app.exe+pat+*/
```

- Running the **app.exe+pat** creates an **app.exe+pat+*/** directory.
- **pat_report** reads that directory and prints a lot of human-readable performance data.
- If worthwhile, a MPI rank reordering file will be produced.

Event Tracing using Perftools

```
> module load perftools-base  
> module load perftools
```

```
> make clean; make  
> pat_build -u -g mpi app.exe
```

- This generates a new executable **app.exe+pat** and preserves **app.exe**.
- If object files and user libraries have already been compiled with perftools enabled, just relink the application with **rm app.exe; make** instead.
- Traces MPI functions calls and functions defined in the program source files.

```
> srun -n 8 ./app.exe+pat  
> pat_report -o myrep.trace.rpt app.exe+pat+*/
```

- Running the **app.exe+pat** creates a **app.exe+pat+*/** directory.
- **pat_report** reads that directory and prints human-readable performance data.

Page 10 of 10

```

Number of PES (MPI ranks): 8
Numbers of PEs per Node: 8
Numbers of Threads per PE: 1
Number of Cores per Socket: 16
Execution start time: Tue Mar 7 21:28:21 2017
System name and speed: nid00036 2301 MHz (approx)
Intel haswell CPU Family: 6 Model: 63 Stepping: 2

```

Table 1: Profile by Function Group and Function

Time%	Time	Imb.	Imb.	Calls	Group
	Time	Time%		Function	
PE=HIDE					
100.0%	36.703706	--	--	3,036.0	Total
97.2%	35.666406	--	--	457.0	USER
93.8%	34.432069	0.683771	1.22%	2.0	llacobi
3.4%	1.232095	0.012880	1.2%	1.0	initmt
0.0%	0.000941	0.001077	61.0%	1.0	main
0.0%	0.000486	0.000232	36.9%	150.0	sendp
0.0%	0.000356	0.000445	63.5%	150.0	sendp3
0.0%	0.000266	0.000024	9.3%	1.0	initcomm
0.0%	0.000172	0.000192	60.3%	150.0	sendp2
0.0%	0.000020	0.000008	31.8%	1.0	exit
0.0%	0.000002	0.000001	29.9%	1.0	initmaxclone_12853_1

All user defined routines traced and shown (-T option).

```

||-----||
|| 2.8% | 1.019139 | -- | -- | 2,423.0 | MPI
||-----||
|| 2.3% | 0.833848 | 0.688832 | 51.7% | 450.0 | MPI_Waitall
|| 0.5% | 0.169196 | 0.025124 | 14.8% | 900.0 | MPI_Isend
|| 0.0% | 0.012631 | 0.016372 | 64.5% | 900.0 | MPI_Irecv
|| 0.0% | 0.001162 | 0.000014 | 1.3% | 152.0 | MPI_Allreduce
|| 0.0% | 0.001077 | 0.000043 | 4.4% | 1.0 | MPI_Cart_create
|| 0.0% | 0.000703 | 0.000038 | 5.8% | 3.0 | MPI_Type_commit
||-----||
||-----||
|| 0.0% | 0.018161 | -- | -- | 156.0 | MPI_SYNC
||-----||
|| 0.0% | 0.014154 | 0.014143 | 99.9% | 2.0 | MPI_Barrier(sync)
|| 0.0% | 0.003969 | 0.003266 | 82.3% | 152.0 | MPI_Allreduce(sync)
|| 0.0% | 0.000029 | 0.000027 | 92.5% | 1.0 | MPI_Init(sync)
|| 0.0% | 0.000009 | 0.000009 | 99.6% | 1.0 | MPI_Finalize(sync)
||-----||

```

Table 2: Load Balance with MPI Message Stats (limited entries shown)

Time%	Time Msg Count	MPI	MPI Msg Bytes Msg Size	Avg MPI Group PE=[mmm]
100.0%	36.703700	602.0	197,839,216.0	328,636.57 Total
97.2%	35.666406	0.0	0.0	- USER
99.1%	36.361896	0.0	0.0	- pe.3
96.9%	35.581600	0.0	0.0	- pe.7
95.3%	34.961048	0.0	0.0	- pe.5
2.8%	1.019132	602.0	197,839,216.0	328,636.57 MPI
4.7%	1.722054	602.0	197,839,216.0	328,636.57 pe.5
2.4%	0.881184	602.0	197,839,216.0	328,636.57 pe.1
0.9%	0.332689	602.0	197,839,216.0	328,636.57 pe.3
0.0%	0.018161	0.0	0.0	- MPI_SYNC
0.1%	0.029523	0.0	0.0	- pe.6
0.1%	0.019500	0.0	0.0	- pe.1
0.0%	0.001622	0.0	0.0	- pe.0

MPI calls traced

Load balance with MPI message statistics

Table 3: MPI Message Stats by Caller (limited entries shown)

	MPI Msg	MPI Msg	Bytes Msg	MPI Count	MPI Count	MsgSz <16	MsgSz <1MiB	64KiB<= Caller	Function
	Bytes%							PE=[mmm]	
	100.0%	197,839,216.0	602.0	152.0	450.0	Total			
	100.0%	197,838,600.0	450.0	0.0	450.0	MPI_Isend			
3	40.0%	79,104,600.0	150.0	0.0	150.0	sendp2			
4						sendp			
5						jacobi			
						main			
6	40.0%	79,104,600.0	150.0	0.0	150.0	pe.0			
6	40.0%	79,104,600.0	150.0	0.0	150.0	pe.4			
6	40.0%	79,104,600.0	150.0	0.0	150.0	pe.7			
	40.0%	79,104,600.0	150.0	0.0	150.0	sendp1			
3						jacobi			
4						main			
	40.0%	79,104,600.0	150.0	0.0	150.0	pe.0			
5	40.0%	79,104,600.0	150.0	0.0	150.0	pe.4			
5	40.0%	79,104,600.0	150.0	0.0	150.0	pe.7			
	20.0%	39,629,400.0	150.0	0.0	150.0	sendp3			

Load balance with MPI message statistics

Remarks for Tracing

- More information is given in the [pat_build](#) man page
- Only true function calls can be traced. Functions that are inlined by the compiler or that have local scope in a compilation unit cannot be traced
- [-g tracegroup](#) Instrument the program to trace all function references belonging to the trace function group tracegroup. Only those functions actually executed by the program at runtime are traced. A selection of tracegroup values is:
 - [blas](#) Basic Linear Algebra subprograms
 - [netcdf](#) Network Common Data Form
 - [hdf5](#) HDF5 I/O library
 - [heap](#) dynamic heap
 - [io](#) includes stdio and sysio groups
 - [lapack](#) Linear Algebra Package
 - [mpi](#) MPI
 - [omp](#) OpenMP API
 - [sysio](#) I/O system calls
 - [syscall](#) system calls
 - [hip](#) hip calls
- More information on the various tracegroup values is given in [\\$CRAYPAT_ROOT/share/traces](#) after loading the [perftools-base](#) module



Event Tracing using Perftools for GPU (HIP)

```
> module load PrgEnv-cray  
> module load craype-x86-trento craype-accel-amd-gfx90a rocm  
> module load perftools-base  
> module load perftools
```

```
> make clean; make  
> pat_build -u -g hip app.exe
```

- This generates a new executable **app.exe+pat** and preserves **app.exe**.
- If object files and user libraries have already been compiled with perftools enabled, just relink the application with **rm app.exe; make** instead.
- Traces MPI functions calls, openMP offload directives and functions defined in the program source files.

```
> srun -n 8 ./app.exe+pat  
> pat_report -T -o myrep.trace.rpt app.exe+pat+*/
```

- Running the **app.exe+pat** creates a **app.exe+pat+*/** directory.
- **pat_report** reads that directory and prints human-readable performance data.

Output: Event Tracing with Perftools for GPU (HIP)

General job information

CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22 02:03:49
Number of PEs (MPI ranks): 1
Numbers of PEs per Node: 1
Numbers of Threads per PE: 2
Number of Cores per Socket: 64
Execution start time: Sun Feb 12 10:34:24 2024
System name and speed: nid007295 2.010 GHz (nominal)
AMD Trento CPU Family: 25 Model: 48 Stepping: 1

Table 1: Profile by Function Group and Function

Time%	Time	Imb.	Imb.	Calls	Group
	Time	Time%			Function
					Thread=HIDE
100.0%	1.055954	--	--	1,255.0	Total

99.9%	1.054614	--	--	808.0	HIP

88.1%	0.930584	--	--	110.0	hipDeviceSynchronize
11.7%	0.123213	--	--	644.0	hipMalloc
0.1%	0.000684	--	--	18.0	hipLaunchKernel
0.0%	0.000057	--	--	10.0	hipMemcpy
0.0%	0.000039	--	--	10.0	hipMemsetAsync
0.0%	0.000012	--	--	4.0	hipFree
0.0%	0.000011	--	--	1.0	hipKernel.count_kernel
0.0%	0.000008	--	--	3.0	hipGetDeviceProperties
0.0%	0.000004	--	--	4.0	__hipPushCallConfiguration
0.0%	0.000003	--	--	4.0	__hipPopCallConfiguration
=====					

Hip routines

```
| 0.1%|0.001300| --| --| 446.0|ETC
|-----|
|| 0.1%|0.001225| --| --| 3.0|__hip_module_dtor
|| 0.0%|0.000044| --| --| 41.0|hiprtcCreateProgram
|| 0.0%|0.000025| --| --| 400.0|call_init.part.0
|| 0.0%|0.000004| --| --| 1.0|__hip_module_ctor
|| 0.0%|0.000002| --| --| 1.0|__hip_register_globals
|=====|
| 0.0%|0.000040| --| --| 1.0|USER
|-----|
|| 0.0%|0.000040| --| --| 1.0|main
|=====|
|=====|
```

Table 2: Profile of maximum function times

Time%	Time	Imb.	Imb.	Function
	Time	Time%		Thread=HIDE

100.0%	0.930584	--	--	hipDeviceSynchronize
13.2%	0.123213	--	--	hipMalloc
0.1%	0.001225	--	--	__hip_module_dtor
0.1%	0.000684	--	--	hipLaunchKernel
0.0%	0.000057	--	--	hipMemcpy
0.0%	0.000044	--	--	hiprtcCreateProgram
0.0%	0.000040	--	--	main
0.0%	0.000039	--	--	hipMemsetAsync
0.0%	0.000025	--	--	call_init.part.0
0.0%	0.000012	--	--	hipFree
0.0%	0.000011	--	--	hipKernel.count_kernel
0.0%	0.000008	--	--	hipGetDeviceProperties
0.0%	0.000004	--	--	__hip_module_ctor
0.0%	0.000004	--	--	__hipPushCallConfiguration
0.0%	0.000003	--	--	__hipPopCallConfiguration
0.0%	0.000002	--	--	__hip_register_globals
=====				

Event Tracing using Perftools for GPU (MPI+OpenMP offload)

```
> module load PrgEnv-cray  
> module load craype-x86-trento craype-accel-amd-gfx90a rocm  
> module load perftools-base  
> module load perftools
```

```
> make clean; make  
> pat_build -u -g mpi,omp app.exe
```

- This generates a new executable **app.exe+pat** and preserves **app.exe**.
- If object files and user libraries have already been compiled with perftools enabled, just relink the application with **rm app.exe; make** instead.
- Traces MPI functions calls, openMP offload directives and functions defined in the program source files.

```
> srun -n 8 ./app.exe+pat  
> pat_report -T -o myrep.trace.rpt app.exe+pat+*/
```

- Running the **app.exe+pat** creates a **app.exe+pat+*/** directory.
- **pat_report** reads that directory and prints human-readable performance data.

Output: Event Tracing with Perftools for GPU (MPI+openMP offload)

CrayPat/X: Version 22.06.0 Revision 4b5ab6256 05/21/22 02:03:49
Number of PEs (MPI ranks): 8
Numbers of PEs per Node: 4 PEs on each of 2 Nodes
Numbers of Threads per PE: 1
Number of Cores per Socket: 64
Execution start time: Sat Feb 11 10:37:06 2024
System name and speed: nid007370 2.009 GHz (nominal)
AMD Trento CPU Family: 25 Model: 48 Stepping: 1

Table 1: Profile by Function Group and Function

Time%	Time	Imb.	Imb.	Calls	Group	Function
						PE=HIDE
100.0%	24.551284	--	--	3,449.0	Total	

90.3%	22.160000	--	--	1,733.0	OACC	

83.2%	20.418715	0.373699	2.1%	1,094.0	jacobi.ACC_COPY@li.242	
4.7%	1.164765	0.021092	2.0%	90.0	jacobi.ACC_COPY@li.236	
2.3%	0.576451	0.038592	7.2%	429.0	jacobi.ACC_COPY@li.268	
0.0%	0.000038	0.000013	28.6%	60.0	jacobi.ACC_KERNEL@li.242	
0.0%	0.000030	0.000020	45.6%	60.0	jacobi.ACC_KERNEL@li.268	
=====						

6.2%	1.515424	--	--	130.0	USER

6.1%	1.505413	0.002802	0.2%	1.0	initmt
0.0%	0.009833	0.000202	2.3%	4.0	jacobi
0.0%	0.000090	0.000026	25.7%	1.0	main
0.0%	0.000049	0.000018	30.9%	60.0	jacobi.ACC_REGION@li.242
0.0%	0.000032	0.000020	44.9%	60.0	jacobi.ACC_REGION@li.268
0.0%	0.000008	0.000001	15.6%	4.0	jacobi.ACC_DATA_REGION@li.236
=====					



General job information

openMP offload regions

MPI routines

User traced routines

HIP kernels calls during openMP offload

3.5%	0.862093	--	--	983.0	MPI

3.4%	0.825582	0.313861	31.5%	180.0	MPI_Waitall
0.1%	0.028702	0.002903	10.5%	360.0	MPI_Isend
0.0%	0.006636	0.016262	81.2%	360.0	MPI_Irecv
0.0%	0.001061	0.000024	2.5%	62.0	MPI_Allreduce
0.0%	0.000075	0.000010	12.9%	1.0	MPI_Cart_create
0.0%	0.000021	0.000004	19.6%	2.0	MPI_Barrier
0.0%	0.000007	0.000002	21.2%	4.0	MPI_Wtime
0.0%	0.000003	0.000003	54.7%	3.0	MPI_Type_commit
0.0%	0.000002	0.000002	54.4%	3.0	MPI_Type_vector
0.0%	0.000002	0.000000	24.1%	1.0	MPI_Cart_get
0.0%	0.000001	0.000000	20.9%	3.0	MPI_Cart_shift
0.0%	0.000000	0.000000	21.7%	1.0	MPI_Finalize
0.0%	0.000000	0.000000	20.0%	1.0	MPI_Init
0.0%	0.000000	0.000000	14.0%	1.0	MPI_Comm_size
0.0%	0.000000	0.000000	25.1%	1.0	MPI_Comm_rank
=====					
0.0%	0.005111	--	--	66.0	MPI_SYNC

0.0%	0.002822	0.002792	98.9%	2.0	MPI_Barrier(sync)
0.0%	0.002211	0.000429	19.4%	62.0	MPI_Allreduce(sync)
0.0%	0.000066	0.000062	92.8%	1.0	MPI_Init(sync)
0.0%	0.000012	0.000010	80.8%	1.0	MPI_Finalize(sync)
=====					

0.0%	0.000606	--	--	121.0	HIP

0.0%	0.000357	0.000092	23.3%	60.0	hipKernel__omp_offloading_43b2fce4_1f016d09_jacobi_I242
0.0%	0.000204	0.000084	33.1%	60.0	hipKernel__omp_offloading_43b2fce4_1f016d09_jacobi_I268
0.0%	0.000044	0.000000	0.9%	1.0	hipGraphicsUnmapResources
=====					
0.0%	0.000041	0.000008	17.9%	400.0	ETC

0.0%	0.000041	0.000008	17.9%	400.0	__libc_csu_init
=====					

Loop Work Estimates using Perftools

```
> module load perftools-base  
> module load perftools
```

- Use the compiler flag **-h profile_generate** for Fortran or **-finstrument-loops** for C for your build (**Cray Compiler only**). This flag turns off OpenMP and significant compiler loop restructuring optimizations except for vectorization.

```
> make clean; make  
> pat_build -w app.exe
```

- This generates a new executable **app.exe+pat** and preserves **app.exe**.
- Loop profiling requires a rebuild of object files and user libraries.
- Automatically traces functions defined in the program source files despite **-w**

```
> srun -n 8 ./app.exe+pat  
> pat_report -T -o myrep.trace.rpt app.exe+pat+*/
```

- Running the **app.exe+pat** creates a directory **app.exe+pat+*/** experiment directory.
- **pat_report** reads that directory and prints a lot of human-readable performance data.

Automatic Profiling Analysis (1/2)

```
> module load perftools-base  
> module load perftools
```

```
> make clean; make  
> pat_build app.exe
```

- The APA (**-O apa**) is the default experiment. No option needed.
- The **pat_build** generates a binary instrumented for sampling (different from the pure sampling shown before with **-S** option)

```
> srun -n 8 app.exe+pat  
> pat_report -o myrep.txt app+pat+*/
```

- Running the “**+pat**” binary creates an experiment directory.
- Applying **pat_report** to the **app+pat+*/** directory generates an ***.apa** file therein.

Automatic Profiling Analysis (2/2)

```
> vi app.exe+pat+*/*.apa
```

- The ***.apa** file contains instructions for the next step, i.e. tracing. Modify it according to your needs.

```
> pat_build -O app.exe+pat+*/*.apa
```

- Generates an instrumented binary **app.exe+apa** for tracing according to the instructions in the **app.exe+pat+*/*.apa** file.

```
> srun -n 8 ./app.exe+apa  
> pat_report -o myrep.txt app.exe+apa+*/
```

- Running the **app.exe+apa** binary creates a new data file or directory.
- Applying **pat_report** to the **app+apa+*/** directory generates a new report.

Output: APA with perftools

Sampling profiling

```
CrayPat/X: Version 6.4.1 Revision 6a6694f 06/27/16 17:24:11
Number of PEs (MPI ranks): 8
Numbers of PEs per Node: 8
Numbers of Threads per PE: 1
Number of Cores per Socket: 16

Execution start time: Wed Mar 8 14:24:42 2017
System name and speed: nid00036 2301 MHz (approx)
Intel haswell CPU Family: 6 Model: 63 Stepping: 2
...

Samp% | Samp | Imb. | Imb. | Group
      | Samp | Samp | Samp | Function
      |     |     |     | PE=HIDE
-----|-----|-----|-----|-----
100.0% | 3,685.0 | -- | -- | Total
-----|-----|-----|-----|-----
84.2% | 3,104.5 | -- | -- | USER
-----|-----|-----|-----|-----
81.9% | 3,018.1 | 57.9 | 2.2% | jacobi
2.3% | 86.1 | 2.9 | 3.7% | initmt
-----|-----|-----|-----|-----
13.3% | 489.9 | -- | -- | ETC
-----|-----|-----|-----|-----
12.2% | 451.2 | 6.8 | 1.7% | _cray_scopy_HSW
1.0% | 38.6 | 2.4 | 6.6% | _cray_sset_HSW
-----|-----|-----|-----|-----
2.5% | 90.6 | -- | -- | MPI
-----|-----|-----|-----|-----
1.9% | 69.6 | 46.4 | 45.7% | MPI_Waitall
-----|-----|-----|-----|-----

Total
-----|-----|-----|-----|-----
CPU_CLK_THREAD_UNHALTED:THREAD_P 109,408,212,031
CPU_CLK_THREAD_UNHALTED:REF_XCLK 3,588,344,079
DTLB_LOAD_MISSES:MISS_CAUSES_A_WALK 53,188,056
PM_ENERGY:NODE 30.382 /sec 1,112 J
CPU_CLK 3.05GHz
TLB utilization 268.83 refs/miss 0.53 avg uses
D1 cache hit,miss ratios 70.8% hits 29.2% misses
D1 cache utilization (misses) 3.43 refs/miss 0.43 avg hits
D2 cache hit,miss ratio 67.8% hits 32.2% misses
D1+D2 cache hit,miss ratio 90.6% hits 9.4% misses
D1+D2 cache utilization 10.65 refs/miss 1.33 avg hits
D2 to D1 bandwidth 6,149.405MiB/sec 236,055,684,984 bytes
```

APA file

```
# -----
# Collect the default PERFCTR group.

-Drtenv=PAT_RT_PERFCTR=default

...

# Libraries to trace.

-g mpi

# -----
# User-defined functions to trace, sorted by % of samples.
...

-w # Enable tracing of user-defined functions.
# Note: -u should NOT be specified as an additional option.

# 81.90% 3735 bytes
-T jacobi

# 2.34% 2450 bytes
-T initmt

# -----
-o himeno.exe+apa # New instrumented program.
...
```

Suggestion to collect Performance counters

Augment list if needed,
i.e. -g mpi,io

Create the binary
for tracing

Add or remove
functions as needed.

Event Tracing tuned profiling

```
CrayPat/X: Version 6.4.1 Revision 6a6694f 06/27/16 17:24:11

Number of PEs (MPI ranks): 8
Numbers of PEs per Node: 8
Numbers of Threads per PE: 1
Number of Cores per Socket: 16
...
Table 1: Profile by Function Group and Function

Time% | Time | Imb. | Imb. | Calls | Group
      | Time | Time% | PE=HIDE | Function
-----|-----|-----|-----|-----|-----
100.0% | 37.748701 | -- | -- | 2,584.0 | Total
-----|-----|-----|-----|-----|-----
95.9% | 36.184768 | -- | -- | 5.0 | USER
-----|-----|-----|-----|-----|-----
92.5% | 34.935081 | 1.178364 | 3.7% | 2.0 | jacobi
3.3% | 1.248950 | 0.016487 | 1.5% | 1.0 | initmt
=====|=====|=====|=====|=====|=====
4.1% | 1.537996 | -- | -- | 2,423.0 | MPI
-----|-----|-----|-----|-----|-----
3.6% | 1.344349 | 1.096798 | 51.3% | 450.0 | MPI_Waitall
=====|=====|=====|=====|=====|=====

USER / jacobi
-----|-----|-----|-----|-----|-----
Time% 92.5%
Time 34.935081 secs
Imb. Time 1.178364 secs
Imb. Time% 3.7%
Calls 0.057 /sec 2.0 calls
CPU_CLK_THREAD_UNHALTED:THREAD_P 104,712,083,433

CPU_CLK_THREAD_UNHALTED:REF_XCLK 3,465,232,627
DTLB_LOAD_MISSES:MISS_CAUSES_A_WALK 44,033,261
DTLB_STORE_MISSES:MISS_CAUSES_A_WALK 26,116,487

L1D:REPLACEMENT 6,955,645,277
L2_RQSTS:ALL_DEMAND_DATA_RD 3,675,034,522
L2_RQSTS:DEMAND_DATA_RD_HIT 1,408,925,969
MEM_UOPS_RETIRED:ALL_LOADS 22,406,881,264
CPU_CLK 3.02GHz
TLB utilization 319.41 refs/miss 0.62 avg uses
D1 cache hit,miss ratios 69.0% hits 31.0% misses
D1 cache utilization (misses) 3.22 refs/miss 0.40 avg hits
```


Using **pat_report**

- Always need to run **pat_report** at least once to perform data conversion
 - Experiment directory contains *.xf files in the xf-files/ subdirectory
 - Perftools combines information from **xf** output and executable to produce **ap2** stored in the ap2-files/ subdirectory : **instrumented binary must still exist when data is converted!**
 - Resulting **ap2** files are the input for subsequent **pat_report** calls and Reveal or Apprentice²
 - **Always use the entire experiment directory as input and not single ap2 files**
 - **xf** files and instrumented binary files could be removed once **ap2** file is generated.
- Generates a text report of performance results
 - Data laid out in tables
 - Many options for sorting, slicing or dicing data in the tables.
 - > **pat_report -O <table option> app.exe+pat+*/**
 - > **pat_report -O help (list of available profiles)**
 - Volume and type of information depends upon sampling vs tracing



Using `pat_report`

- The performance numbers reported are in general an average over all tasks (also explains non-integer values)
- Not always meaningful
 - controller/worker schemes
 - MPMD
- To solve this you can filter the *.ap2 file

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function PE=HIDE
100.0%	20.643909	--	--	1149.0	Total
98.8%	20.395989	--	--	219.0	USER
91.1%	18.797060	0.115535	0.7%	2.0	jacobi
7.7%	1.597866	0.006647	0.5%	1.0	initmt
0.0%	0.000402	0.000167	33.5%	53.0	sendp3

> `pat_report -sfilter_input='condition' ...`

- The 'condition' should be an expression involving 'pe' such as 'pe<1024' or 'pe%2==0'.
- This option is also useful when the size of the full data file makes a report incorporating data from all PEs take too long or exceed the available memory
- More help:
 - `pat_report -h` => usage
 - `pat_report -O -h` => available report tables
 - `pat_report -s -h` => options for content and format
 - `pat_report -d -h` => options for data columns



General Remarks

- Always check that the **instrumenting binary** has not affected the run time significantly compared to the original executable
- Collecting event traces on **large numbers of frequently called functions**, or setting the **sampling interval very low** can introduce a **lot of overhead** (check **trace-text-size** option to **pat_build**)
- **Highly recommended to run on Lustre !**
- The runtime analysis can be modified using environment variables of the form **PAT_RT_***
 - Check the **PAT_LD_OBJECT_TMPDIR** variable if you cannot preserve the original build tree
- **pat_build** may recognize for instance MPI in your application and trace **MPI_Init** and **MPI_Finalize** when adding **-g mpi** trace **all MPI calls**
- Processing the **app+*/** directory from **perftools-lite** yields all the options of **pat_build** to reproduce the experiment with regular **perftools**.



Reveal

Performance analysis and code restructuring assistant



Motivation: Compiler Listings (CCE)

- Much information produced by compiler.
 - Can be listed together with source code for better clarity.
 - Use **-fsave-loopmark** for C and **-hlist=m** (or **a**) for Fortran.
 - It generates an ***.lst** file for every source code file.
 - A + sign indicates that more information can be found after the routine definition.
 - Can also be inspected in Reveal with a corresponding program library.

==== Loopmark Legend ====

Primary Loop Type	Modifiers
A - Pattern matched	a - atomic memory operation
	b - blocked
C - Collapsed	c - conditional and/or computed
D - Deleted	
E - Cloned	
F - Flat - No calls	f - fused
G - Accelerated	g - partitioned
I - Inlined	i - interchanged
M - Multithreaded	m - partitioned
	n - non-blocking remote transfer
	p - partial
R - Rerolling	r - unrolled
	s - shortloop
V - Vectorized	w - unwound

+ - More messages listed at end of listing

```
210. +-----< for(i=0 ; i<imax ; ++i)
211. + 1 2-----< for(j=0 ; j<jmax ; ++j)
212. 1 2 Vr2-----< for(k=0 ; k<kmax ; ++k){
213. 1 2 Vr2      a[0][i][j][k]=1.0;
214. 1 2 Vr2      a[1][i][j][k]=1.0;
215. 1 2 Vr2      a[2][i][j][k]=1.0;
216. 1 2 Vr2      a[3][i][j][k]=1.0/6.0;
217. 1 2 Vr2      b[0][i][j][k]=0.0;
218. 1 2 Vr2      b[1][i][j][k]=0.0;
219. 1 2 Vr2      b[2][i][j][k]=0.0;
220. 1 2 Vr2      c[0][i][j][k]=1.0;
221. 1 2 Vr2      c[1][i][j][k]=1.0;
222. 1 2 Vr2      c[2][i][j][k]=1.0;
223. 1 2 Vr2 A-----< p[i][j][k]=(float)((i+it)*(i+it))/(float)((mx-1)*(mx-1));
224. 1 2 Vr2 A-----< wrk1[i][j][k]=0.0;
225. 1 2 Vr2 A-----< wrk2[i][j][k]=0.0;
226. 1 2 Vr2 A-----< bnd[i][j][k]=1.0;
227. 1 2 Vr2----->>> }
228. }
```

CC-6294 CC: VECTOR initmt, File = himeno.c, Line = 210
A loop was not vectorized because a better candidate was found at line 212.

CC-6294 CC: VECTOR initmt, File = himeno.c, Line = 211
A loop was not vectorized because a better candidate was found at line 212.

CC-6005 CC: SCALAR initmt, File = himeno.c, Line = 212
A loop was unrolled 2 times.

CC-6204 CC: VECTOR initmt, File = himeno.c, Line = 212
A loop was vectorized.

Input to Reveal: Program Library

- Mandatory

```
> ftn -O3 -hpl=my_program.pl -c my_program_file1.f90  
> cc -O3 -fcray-program-library-path=my_program.pl -c my_program_file2.c  
> reveal my_program.pl &
```

- Recompile only sources to generate program library `my_program.pl`

- Optional but highly recommended

```
> ftn -O3 -hpl=my_program.pl -c my_program_file1.f90  
> cc -O3 -fcray-program-library-path=my_program.pl -c my_program_file2.c  
> reveal my_program.pl my_program_exe_*/ &
```

- **Collect loop work estimates in a separate experiment** and load it too with `my_program.pl`
- Note that `-hprofile_generate` option disables OpenMP and significant compiler loop restructuring optimizations except for vectorization



Reveal

- View source, performance, and optimization information at the same time.

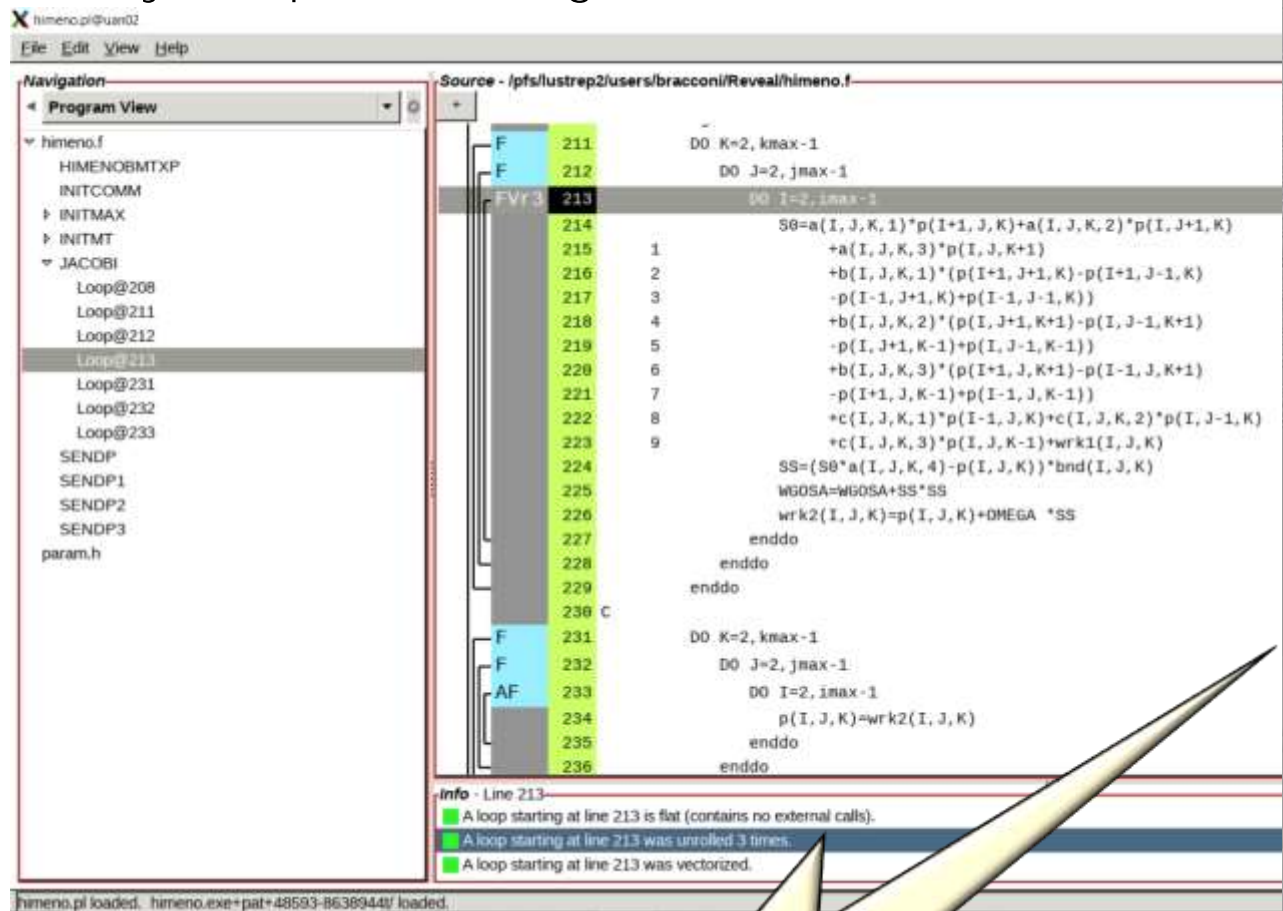
The screenshot displays the Reveal IDE interface. On the left, a 'Navigation' pane shows a tree view of the code structure, including 'Loop Perform' and 'Instance #2'. The main window shows the source code of a file named 'himeno.f'. The code is color-coded and includes comments. A red box highlights the 'Loop Perform' section in the navigation pane. On the right, a 'Loopmark Legend' pane lists various optimization markers and their meanings. The legend includes:

- A Pattern Matched
- C Collapsed
- D Deleted
- E Cloned
- F Flat
- G Accelerated
- I Inlined
- II Not Inlined
- L Loop
- M Multithreaded
- R Region
- S Scoping Analysis
- V Vectorized
- a Atomic Memory Operation
- b Blocked
- c Conditional and/or Computed
- f Fused
- g Partitioned
- i Interchanged
- n Non-blocking Remote Transfer
- p Partial
- r Unrolled
- s Shortloop
- w Unwound

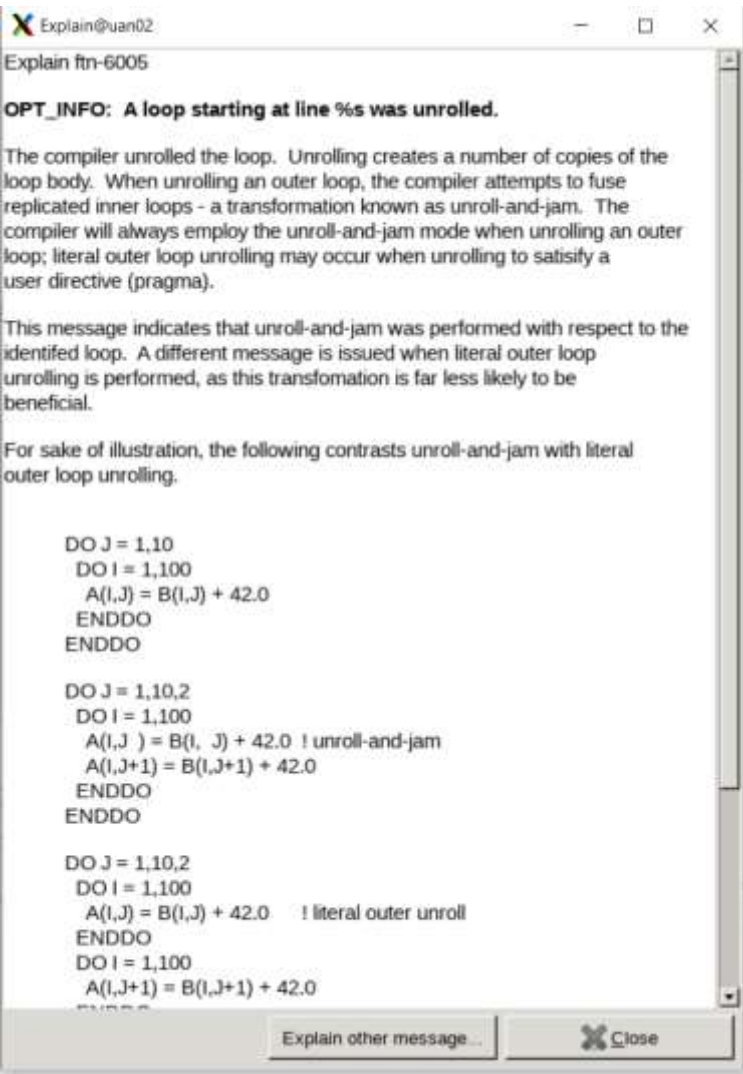
At the bottom, an 'Info' pane provides details about the code, such as 'Line 211' and 'Line 212'.

Reveal

- Access Cray compiler message information



Double click on optimization message for more detailed information



Reveal

- Scope selected loop(s)

The screenshot shows the Reveal OpenMP Scoping tool interface. A red box labeled '1' points to the 'View' menu. A red box labeled '2- loop selection' points to the 'Loop in function JACOBI' entry at line 211 in the 'List of Loops to be Scoped' table. A red box labeled '3- select processor' points to the 'Scope For CPU' and 'Scope For GPU' buttons. The interface includes a 'Navigation' pane on the left, a 'Source' pane in the center showing Fortran code, and an 'Info' pane at the bottom. The 'Scope Loops' dialog box is open, displaying a table of loops to be scoped.

Scope?	Line #	File or Source Line
<input checked="" type="checkbox"/>		/pfs/lustrep2/users/bracconi/Reveal/himeno.f
<input type="checkbox"/>	154	Loop in function INITMT
<input type="checkbox"/>	155	Loop in function INITMT
<input type="checkbox"/>	156	Loop in function INITMT
<input type="checkbox"/>	175	Loop in function INITMT
<input type="checkbox"/>	176	Loop in function INITMT
<input type="checkbox"/>	177	Loop in function INITMT
<input type="checkbox"/>	208	Loop in function JACOBI
<input checked="" type="checkbox"/>	211	Loop in function JACOBI
<input type="checkbox"/>	212	Loop in function JACOBI
<input type="checkbox"/>	213	Loop in function JACOBI
<input type="checkbox"/>	231	Loop in function JACOBI
<input type="checkbox"/>	232	Loop in function JACOBI

Apply Filter Time: 0.000 Trips: 2 Threads: 4 Speedup: 0.010

Scope For CPU Scope For GPU Cancel 1 Loop selected Close

Info - Line 211

- A loop starting at line 211 is flat (contains no external calls).
- A loop starting at line 211 was not vectorized because...
- A loop starting at line 212 is flat (contains no external calls).

Reveal

- Scope selected loop(s) for CPU

Reveal OpenMP Scoping@uan02

Scope Loops

Scoping Results

himeno.f: Loop@211

Name	Type	Scope			
a	Array	P	S	C	U
b	Array	P	S	C	U
bnd	Array	P	S	C	U
c	Array	P	S	C	U
i	Scalar	P	S	C	U
imax	Scalar	P	S	C	U
j	Scalar	P	S	C	U
jmax	Scalar	P	S	C	U
k	Scalar	P	S	C	U
kmax	Scalar	P	S	C	U

First/Last Private

☐ Enable FirstPrivate

☐ Enable LastPrivate

Reduction

None

Find Name:

Insert Directive

Show Directive

Close

OpenMP Directive@uan02

! Directive inserted by Cray Reveal. May be incomplete.

!\$OMP parallel do default(none)

!\$OMP& shared(a, b, bnd, c, imax, jmax, kmax, omega, p, wgosa, wrk1,


!\$OMP& wrk2)

!\$OMP& private(i, j, k, s0, ss)

!\$OMP& reduction(+: wgosa)

Copy Directive

Close



Reveal

- Scope selected loop(s) for GPU

Reveal OpenMP Scoping@uan02

Scope Loops

Scoping Results

himenof: Loop@211 (gpu)

Name	Type	Map	Extents
a	Array	P T F TF A C U	(:::)
b	Array	P T F TF A C U	(:::)
bnd	Array	P T F TF A C U	(:::)
c	Array	P T F TF A C U	(:::)
i	Scalar	P T F TF A C U	
imax	Scalar	P T F TF A C U	
j	Scalar	P T F TF A C U	
jmax	Scalar	P T F TF A C U	
k	Scalar	P T F TF A C U	
kmax	Scalar	P T F TF A C U	
omega	Scalar	P T F TF A C U	
p	Array	P T F TF A C U	(:::)
s0	Scalar	P T F TF A C U	
ss	Scalar	P T F TF A C U	
wgosa	Scalar	P T F TF A C U	

First/Last Private

Map Always

Metadirective

Collapse

Reduction

☐ Enable FirstPrivate

☐ Always

☐ Metadirective

None

☐ Enable LastPrivate

Find Name:

Insert Directive

Show Directive

Close

OpenMP Directive@uan02

Loop Directive:

! Directive inserted by Cray Reveal. May be incomplete.

!\$OMP target teams distribute

!\$OMP& private(i, j, k, s0, ss)

!\$OMP& firstprivate(imax, jmax, kmax, omega)

!\$OMP& map(tofrom: wgosa)

!\$OMP& reduction(+: wgosa)

!\$OMP& map(always, to: a(:::), b(:::), bnd(:::), c(:::),


!\$OMP& p(:::), wrk1(:::))

!\$OMP& map(always, tofrom: wrk2(:::))

There Are No Declare Target Directives For This Loop.

Copy Directive

Close



Reveal

- Generate OpenMP directives.

```
! Directive inserted by Cray Reveal.  May be incomplete.
!$OMP parallel do default(none)
!$OMP& unresolved (dvol,dx,dx0,e,f,flat,p,para,q,r,radius,svel,u,v,w,
!$OMP& xa,xa0)
!$OMP& private (i,j,k,m,n,$$_n,delp2,delp1,shock,temp2,old_flat,
!$OMP& onemfl,hdt,sinxf0,gamfac1,gamfac2,dtheta,deltx,fractn,
!$OMP& ekin)
!$OMP& shared (gamm,isy,js,ks,mypey,ndim,ngeomy,nlefty,npey,nrighty,
!$OMP& recv1,send2,zdy,zxc,zya)
do k = 1, ks
do i = 1, isy
radius = zxc(i+mypey*isy)

! Put state variables into 1D arrays, padding with 6 ghost zones
do m = 1, npey
do j = 1, js
n = j + js*(m-1) + 6
r(n) = recv1(1,k,j,i,m)
p(n) = recv1(2,k,j,i,m)
u(n) = recv1(4,k,j,i,m)
v(n) = recv1(5,k,j,i,m)
w(n) = recv1(3,k,j,i,m)
f(n) = recv1(6,k,j,i,m)
enddo
enddo

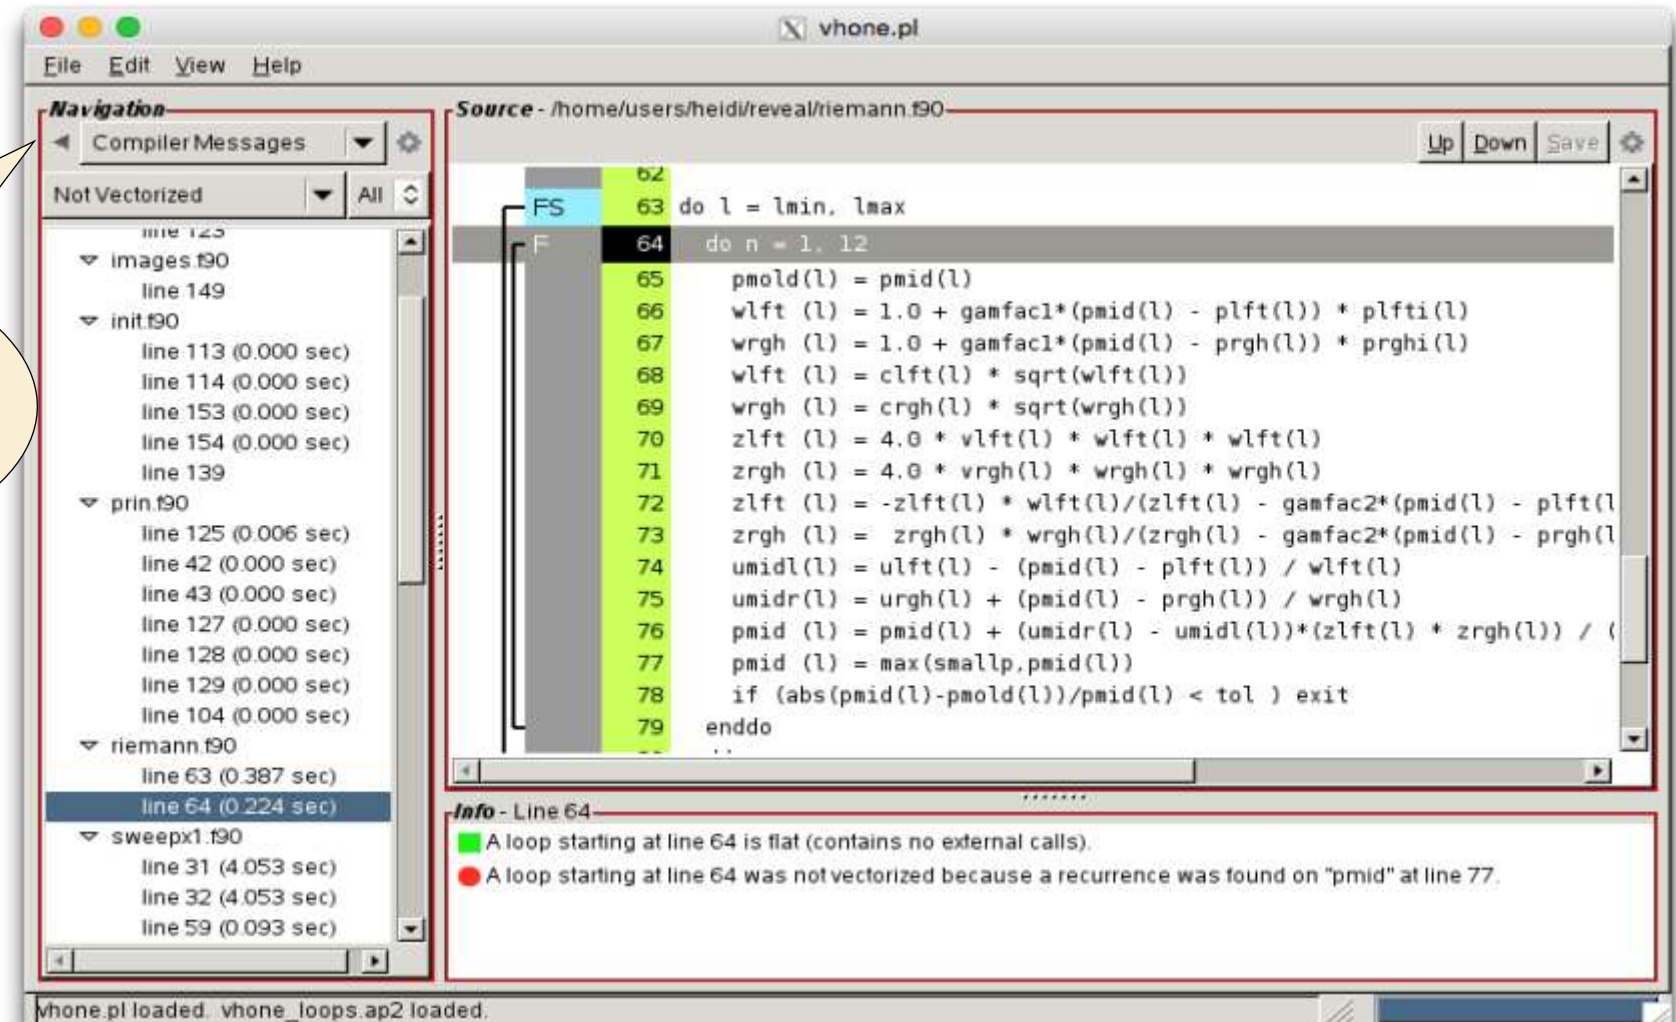
do j = 1, jmax
n = j + 6
```

Reveal generates OpenMP directive with illegal clause marking variables that need addressing

Reveal

- Look for vectorization opportunities.

Choose “Compiler Messages” view to access message filtering, then select desired type of message



OpenMP profiling



OpenMP Data Collection and Reporting

- For OpenMP programs
 - Measure the overhead incurred by entering and leaving parallel regions and work-sharing constructs within parallel regions
 - Show per-thread timings in addition to other data.
 - Calculate the load balance across threads for such constructs.
- For programs that use both MPI and OpenMP
 - Profiles by default show the load balance over PEs of the average time in the threads for each PE
 - But you can also see load balances for each programming model separately.
- Options for **pat_report**
 - **profile_pe.th**
 - Imbalance based on the set of all threads in the program
 - **profile_pe.th** (default view)
 - Highlights imbalance across MPI ranks
 - Uses max for thread aggregation to avoid showing under-performers
 - Aggregated thread data merged into MPI rank data
 - **profile_th_pe**
 - For each thread, show imbalance over MPI ranks
 - Example: Load imbalance shown where thread 4 in each MPI rank didn't get much work



OpenMP Data Collection and Reporting

- OpenMP support needs to be enabled during compilation
- OpenMP tracing calls inserted by default when perftools is loaded

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function PE=HIDE Thread=HIDE
100.0%	2.452453	--	--	1426.8	Total
96.9%	2.377154	--	--	309.8	USER
82.1%	2.013394	0.027282	1.8%	100.0	work.LOOP@li.533
10.6%	0.259470	0.000282	0.1%	1.0	exit
2.4%	0.057711	0.000562	1.3%	1.0	initializeMatrix
1.0%	0.024130	0.000313	1.7%	1.0	setPEsParams.SINGLE@li.355
1.6%	0.039963	--	--	909.0	MPI
1.6%	0.039247	0.079519	89.3%	301.5	MPI_wait
1.2%	0.029108	--	--	101.0	OMP
1.2%	0.029058	0.012000	39.0%	100.0	work.REGION@li.492(ovhd)

Work sharing construct

Region

Overhead

Table 2: Load Imbalance by Thread

Max. Time	Imb. Time	Imb. Time%	Thread PE=HIDE
2.452470	0.316486	17.2%	Total
2.453287	0.000817	0.0%	thread.0
2.078727	0.036293	2.3%	thread.2
2.074969	0.048712	3.1%	thread.1
2.066243	0.043468	2.8%	thread.3

Perftools API

Customized performance analysis.



API for Adding User Instrumentation

The Perftools API calls enable you to insert functions into your source code that write special tracing records into the experiment data file at runtime. Useful for large routines

- API calls are supported in both Fortran and C
- API works for sampling, tracing, and loop profiling as well as with the **perftools-lite*** modules
- When **perftools** module is loaded
 - `-I$CRAYPAT_ROOT/include` is added to compiling flag (C header **pat_api.h**, Fortran header **pat_apif.h** and Fortran 77 header files, **pat_apif.h** and **pat_apif77.h** are available)
 - Macro **CRAYPAT** and `-DCRAYPAT` is added to compiling flag
- **int PAT_region_begin (int id, char *label)**
 - **id** is a unique identifier for the region
 - **label** is the description that will appear in profiling output
- **int PAT_region_end (int id)**
 - **id** must match begin call
- Fortran subroutines with extra final integer argument for return value similar to MPI.
- Data collection through API can be controlled with **PAT_RT_TRACE_API**
- For more information, see **pat_build** man page in the section APPLICATION PROGRAM INTERFACE



PAT Regions Example in C and Fortran

C:

```
#ifdef CRAYPAT
#include "pat_api.h"
#endif
...
#ifdef CRAYPAT
PAT_region_begin( 1, "jacobi_part1");
#endif
// the execution of this code segment will
// appear in CrayPAT output as "jacobi_part1"
...
#ifdef CRAYPAT
PAT_region_end( 1);
#endif
...
#ifdef CRAYPAT
PAT_region_begin( 2, "jacobi_part2");
#endif
// the execution of this code segment will
// appear in CrayPAT output as "jacobi_part2"
...
#ifdef CRAYPAT
PAT_region_end( 2);
#endif
...
```

Fortran:

```
#ifdef CRAYPAT
#include "pat_apif.h"
#endif
...
#ifdef CRAYPAT
Call PAT_region_begin( 1, "jacobi_part1", istat)
#endif
// the execution of this code segment will
// appear in CrayPAT output as "jacobi_part1"
...
#ifdef CRAYPAT
Call PAT_region_end( 1, istat)
#endif
...
#ifdef CRAYPAT
Call PAT_region_begin( 2, "jacobi_part2", istat)
#endif
// the execution of this code segment will
// appear in CrayPAT output as "jacobi_part2"
...
#ifdef CRAYPAT
Call PAT_region_end( 2, istat)
#endif
```

Event Tracing w/ and w/o Perftools API (C)

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function PE=HIDE
100.0%	45.429778	--	--	3,336.0	Total
82.6%	37.513233	--	--	757.0	USER
69.9%	31.777804	7.855317	22.7%	150.0	#1.jacobi_part1
9.9%	4.480465	0.484320	11.1%	150.0	#2.jacobi_part2
2.8%	1.252451	0.066733	5.8%	1.0	initmt
0.0%	0.001015	0.002053	76.5%	1.0	main
0.0%	0.000479	0.000493	58.0%	150.0	sendp
0.0%	0.000298	0.000075	22.9%	2.0	jacobi
0.0%	0.000297	0.000026	9.2%	1.0	initcomm
0.0%	0.000271	0.000211	50.0%	150.0	sendp3
0.0%	0.000132	0.000255	75.2%	150.0	sendp2
0.0%	0.000019	0.000000	2.8%	1.0	exit
17.3%	7.843884	--	--	2,423.0	MPI
16.8%	7.652373	2.858017	31.1%	450.0	MPI_Waitall
0.4%	0.170341	0.033417	18.7%	900.0	MPI_Isend
0.0%	0.016676	0.018946	60.8%	900.0	MPI_Irecv
0.0%	0.001239	0.000039	3.4%	1.0	MPI_Cart_create
0.0%	0.001214	0.000015	1.4%	152.0	MPI_Allreduce
0.0%	0.000933	0.000022	2.6%	3.0	MPI_Type_commit
0.0%	0.000644	0.000031	5.2%	3.0	MPI_Cart_shift
0.0%	0.000459	0.000027	6.5%	3.0	MPI_Type_vector
0.0%	0.000004	0.000002	34.5%	2.0	MPI_Barrier

With API

Regions

Time%	Time	Imb. Time	Imb. Time%	Calls	Group Function PE=HIDE
100.0%	45.649773	--	--	3,036.0	Total
83.4%	38.065287	--	--	457.0	USER
80.6%	36.771025	7.132552	18.6%	2.0	jacobi
2.8%	1.291914	0.077825	6.5%	1.0	initmt
0.0%	0.001031	0.000933	54.3%	1.0	main
0.0%	0.000520	0.000342	45.3%	150.0	sendp
0.0%	0.000337	0.000028	8.7%	1.0	initcomm
0.0%	0.000320	0.000264	51.6%	150.0	sendp3
0.0%	0.000121	0.000322	83.1%	150.0	sendp2
0.0%	0.000017	0.000001	3.9%	1.0	exit
16.4%	7.500291	--	--	2,423.0	MPI
16.0%	7.298155	4.192142	41.7%	450.0	MPI_Waitall
0.4%	0.183874	0.048775	24.0%	900.0	MPI_Isend
0.0%	0.014285	0.020122	66.8%	900.0	MPI_Irecv
0.0%	0.001146	0.000013	1.3%	1.0	MPI_Cart_create
0.0%	0.001003	0.000012	1.4%	152.0	MPI_Allreduce
0.0%	0.000756	0.000012	1.7%	3.0	MPI_Type_commit
0.0%	0.000566	0.000009	1.8%	3.0	MPI_Type_vector
0.0%	0.000501	0.000015	3.3%	3.0	MPI_Cart_shift
0.0%	0.000005	0.000002	34.1%	2.0	MPI_Barrier
0.0%	0.000000	0.000000	15.4%	4.0	MPI_Wtime
0.0%	0.000000	0.000000	18.7%	1.0	MPI_Cart_get

Without API



Load Imbalance Analysis



Load Imbalance

Common cause for performance bottlenecks when running applications at scale

- Look for high imbalance time and percentage
 - User functions
 - $\text{Imbalance time} = \text{Maximum time} - \text{Average time}$
 - Synchronization (Collective communication and barriers)
 - $\text{Imbalance time} = \text{Average time} - \text{Minimum time}$
 - Imbalance percentage of time that the rest of the team is not engaged in useful work on the given function
- Also look for MPI (SYNC) time.
 - Measure for time spent waiting in collectives
 - Only available with event tracing experiments
- Guidance on rank reordering for better load balance might appear in report (MPI Utilization: ...)

Table 1: Profile by Function Group and Function

Time%	Time	Imb. Time	Imb. Time%	Calls	Group	Function
						PE=HIDE
100.0%	1.957703	--	--	42,970.8	Total	
60.0%	1.174021	--	--	3,602.0	USER	
30.8%	0.603850	0.176924	23.0%	1,198.0	function3_	
19.2%	0.375117	0.128748	26.0%	1,200.0	function2_	
9.1%	0.178111	0.081880	32.0%	1,200.0	function1_	
36.0%	0.704928	--	--	9,613.0	MPI_SYNC	
25.8%	0.505174	0.385130	76.2%	9,596.0	mpi_barrier_(sync)	
10.2%	0.199537	0.199518	100.0%	1.0	mpi_init_(sync)	
4.0%	0.078736	--	--	29,754.8	MPI	
2.3%	0.045351	0.003531	7.3%	9,596.0	MPI_BARRIER	
1.1%	0.021520	0.051295	71.6%	8,756.9	MPI_ISEND	

Communication Bottlenecks

- Sort messages by caller
 - Available in default report for tracing (lite-events) experiments
- Analyze message sizes
 - Useful when tuning MPI env vars (eager, rendez-vous, ...)
 - Found in default report
- Put barrier in front of collectives to filter sync times
- Use rank reordering for max. on-node communication
 - Look for guidance and instructions in report
- Visualizing data
 - `grid_order` utility
 - Apprentice2 (`app2 app.exe+*/`) imbalance or activity report
 - `pat_report -s pe=ALL` shows data by MPI rank

MPI	MPI Msg	MPI Msg	MsgSz	4KiB<=	Function
Msg	Bytes	Count	<16	MsgSz	Caller
Bytes%			Count	<64KiB	PE=[nm]
			Count		
100.0%	34,940,767.4	8,771.9	258.6	8,513.3	Total

100.0%	34,940,647.4	8,756.9	243.6	8,513.3	IMPI_ISEND

56.2%	19,622,700.0	4,837.5	56.2	4,781.2	calc2_31

56.4%	19,718,400.0	7,200.0	2,400.0	4,800.0	pe.0
56.4%	19,699,200.0	4,800.0	0.0	4,800.0	pe.32
42.3%	14,784,000.0	4,800.0	1,200.0	3,600.0	pe.63

Total	
MPI Msg Bytes%	100.0%
MPI Msg Bytes	4,465,684,125.8
MPI Msg Count	13,057.0 msgs
MsgSz <16 Count	719.0 msgs
16<= MsgSz <256 Count	28.0 msgs
256<= MsgSz <4KiB Count	0.7 msgs
4KiB<= MsgSz <64KiB Count	279.8 msgs
64KiB<= MsgSz <1MiB Count	12,029.6 msgs

MPI_Send	
MPI Msg Bytes%	100.0%
MPI Msg Bytes	4,465,680,353.8
MPI Msg Count	12,318.0 msgs
MsgSz <16 Count	0.0 msgs
16<= MsgSz <256 Count	0.0 msgs
256<= MsgSz <4KiB Count	0.7 msgs
4KiB<= MsgSz <64KiB Count	279.8 msgs
64KiB<= MsgSz <1MiB Count	12,029.6 msgs

MPI_Send / LAMMPS_HO::Comm::recv_comm	
MPI Msg Bytes%	49.6%
MPI Msg Bytes	2,171,466,150.3
MPI Msg Count	6,006.0 msgs
MsgSz <16 Count	0.0 msgs
16<= MsgSz <256 Count	0.0 msgs
256<= MsgSz <4KiB Count	0.0 msgs
4KiB<= MsgSz <64KiB Count	0.0 msgs
64KiB<= MsgSz <1MiB Count	6,006.0 msgs

MPI_Send / LAMMPS_HO::Comm::recv_comm / LAMMPS_HO::World::zgm	
MPI Msg Bytes%	49.6%
MPI Msg Bytes	2,169,218,110.3
MPI Msg Count	6,000.0 msgs
MsgSz <16 Count	0.0 msgs
16<= MsgSz <256 Count	0.0 msgs
256<= MsgSz <4KiB Count	0.0 msgs
4KiB<= MsgSz <64KiB Count	0.0 msgs
64KiB<= MsgSz <1MiB Count	6,000.0 msgs

Hardware Performance Counters (HWPC)



Hardware Performance Counters

- Perftools supports the use of hardware counters to collect hardware events
 - All counters accessed through the PAPI interface.
 - Predefined sets of hardware counters are specified that can be instrumented for performance analysis experiment.
 - Number of simultaneous counters limited by hardware.
- Perftools provides information at the function call level on hardware features like caches, vectorization and memory bandwidth.
 - Very useful feature for deep understanding of application performance bottlenecks.
 - Impact of compiler options and code optimization.
- HWPC collection can slow down the execution notably.
 - Should be used within a tracing experiment only for a small set of functions or ideally through an automatic performance analysis.
- Only LUMI, only a limited number of hardware counters are available:
 - Use **papi_avail** on a compute node to get a list



When To Collect Hardware Performance Counters

- Use to understand the “why” of a bottleneck.
- Default set of CPU counters are already collected for whole program
 - Used to present memory and vector summary metrics
- To collect performance counters
 - Set `PAT_RT_PERFCTR` environment variable to list of events or group prior to execution.
(Or use `-Drtenv=PAT_RT_PERFCTR=<event list> | <group>` for `pat_build`.
Environment variable has priority.)
- Run the following utility on a compute node to get list of events for a processor:
 - `papi_native_avail`
 - `papi_avail`
- Use `pat_help` to see counter groups and derived metrics
 - `$> pat_help counters processor_type deriv`
 - Example for LUMI-C: `$> pat_help counters trento deriv`

```
=====
Total
-----
Time%                                100.0%
Time                                70.507153 secs
Imb. Time                           -- secs
Imb. Time%                           --
Calls                               5.744 /sec          405.0 calls
PAPI_BR_TKN                        0.064G/sec      4,506,524,243.125 branch
PAPI_TOT_INS                       12.154G/sec     856,922,640,377.375 instr
PAPI_BR_INS                        0.124G/sec      8,718,726,704.750 branch
PAPI_TOT_CYC                       244,793,361,310.500 cycles
Instr per cycle                     3.50 inst/cycle
MIPS                               97,229.58M/sec
Average Time per Call                0.174092 secs
CrayPat Overhead : Time  0.0%
=====
```



pat_run

Launch a dynamically-linked program instrumented for performance analysis



Profile Existing Dynamically Linked Binaries

```
$> srun -n 16 pat_run ./app.exe  
$> pat_report app.exe+*/ > my_report
```

- Insert **pat_run** before executable. **No instrumentation needed**
- Useful if source code is not accessible but program is dynamically linked

```
$> export PAT_RT_PERFCTR=1  
$> srun -n 8 pat_run ./app.exe
```

- Use existing perftools capability
- Optionally collect a different group of performance counters

```
$> pat_report -P -O callers+src app.exe+*/ > my_callers_report
```

- Create additional views of the data with **pat_report** options
- If at least object files and libraries are available
 - load the **perftools-preload** module before relinking
 - more information like line number hot spots are given

Final notes and Recap on PerfTools



Perftools-lite modules vs Perftools module

- Perftools-lite modules : "novice" approach
 - Little/no knowledge of the application and do not want to dive deeper into the source code or profiling tools
 - Load the `-lite` module, make, run, analysis of automatically generated report
 - Sampling `perftools-lite` , tracing `perftools-lite-events` , loop work estimate `perftools-lite-loops`, GPU `perftools-lite-gpu`
 - Everything is automatic
 - Compilation and run commands are the same : **instrumented binary has the same name as the original one**
 - Report automatically generated
- Perftools module : "advanced" approach
 - Good knowledge of the application and the used libraries (MPI, openMP, gpu, io and more)
 - module load `perftools`, make, **`pat_build`** `opt`, run, **`pat_report`**, analysis of generated report
 - Sampling `opt=-S`; tracing `opt=-u`, loop work estimate `opt=-w`, GPU `opt=-g omp/cuda/hip`
 - AND MUCH MORE `opt=-g mpi,io,blas,lapack,netcdf...`
 - Everything is generated by the user
 - The original binary must be instrumented with **`pat_build`** and the options corresponding to the analysis to be performed
 - Run command must use **the name of the new instrumented binary generated by the `pat_build` command**
 - Report must be generated with **`pat_report`**



Summary

	Sampling	Event Tracing	Loop Work Estimates	APA	Step
perftools-lite	Is the default experiment.	Available through perftools-lite-events	Available through perftools-lite-loops	N/A	1
	Rebuild for all except pat_run . Rerun for all.				2
perftools			Rebuild with -h profile_generate for Fortran or -finstrument-loops for C. Only for Cray Compiler	Is the default experiment	0
	pat_build ...	pat_build [-w -u] [-g]	pat_build -w ...	pat_build ...	1
	For all: Run app.exe+pat				2
	For all: pat_report app+pat+*/				3
				pat_build -O *.apa	4
				Run app.exe+apa	5
				pat_report <apa>.xf	6

Summary

	Sampling	Event Tracing	Loop Work Estimates	APA
Reveal	N/A	N/A	Need a program library obtained with the <code>-hpl=<app>.pl</code> compiler option for Fortran or <code>-fcray-program-library-path=<app.pl></code> for C and optionally an <code><app>+*/</code> directory from a loop profiling experiment (Cray compiler only)	N/A
Hardware performance counters	Can be enabled via <code>-Drtenv=PAT_RT_PERFCTR=<event list> <group></code> option for <code>pat_build</code> or the environment variable <code>PAT_RT_PERFCTR</code> . For <code>perftools-lite*</code> only the environment variable is applicable. Note that for <code>perftools-lite</code> some counters are already collected.			Is typically specified in the <code><app>+*/*.apa</code> file. Note that during the sampling phase, some counters are already collected.
Apprentice2	Can be applied to all <code><app>+*/</code> directories			
pat_run	When you can't compile (only dynamically linked binary or objects and libraries are available).			
pat_info	Can be applied to all <code><app>+*/</code> directories			



What does perftools support?

Cray Performance Tools Components

- Reveal
 - compiler optimization presentation
 - OpenMP
 - memory activity tracking
- Apprentice2 (app2)
- CrayPat runtime
- perftools / perftools-lite experiments
 - pre-defined vs custom instrumentation
 - sampling
 - tracing with runtime summarization
 - predefined trace groups
 - full trace
- pat_build
- pat_report
- pat_run
- pat_region API
- grid_order
- pat_view
- pat_help
- Performance counters
 - PAPI
 - CPU, GPU, network, energy
 - default, predefined groups, individual events

Programming Models

- MPI
- OpenMP
- CUDA
- OpenACC / OpenMP 4.5
- PGAS (upc, Forttran coarrays)
- SHMEM
- OpenCL
- HIP

Distros

- CLE
- Mac
- Windows
- RH/Centos
- SLES

Languages

- Fortran
- C
- C++
- Chapel

Platforms

- Cray XC, CS, Shasta
- HPE Apollo 2000 Gen10Plus, Apollo 80*
- Mac (app2, reveal), Windows (app2)

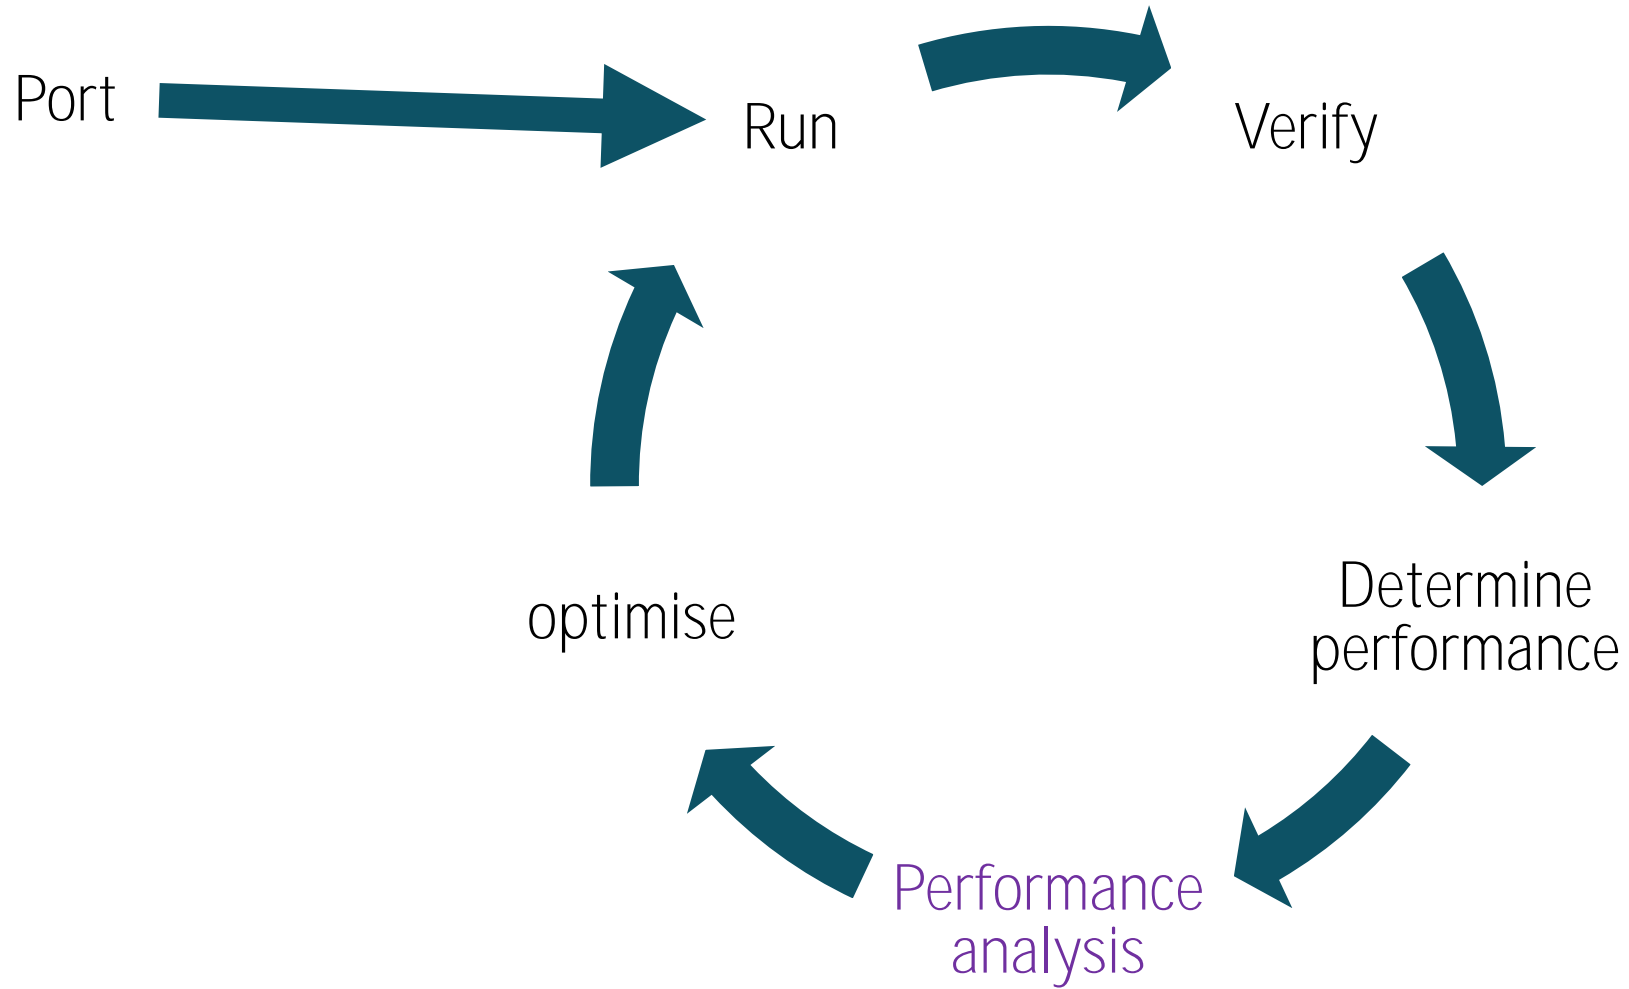
Compilers

- CCE
- GCC
- Intel
- PGI
- AMD (aocc, hipclang)

Processors

- Intel (SNB, IVB, HSW, KNL, BDW, SKL)
- ARM (Cavium TX2)
- AMD (Rome, Milan, MI60/MI100, MI200)
- NVIDIA (Pascal, Volta)

Optimisation Cycle





Questions?